

Kubernetes (V1.18) Study

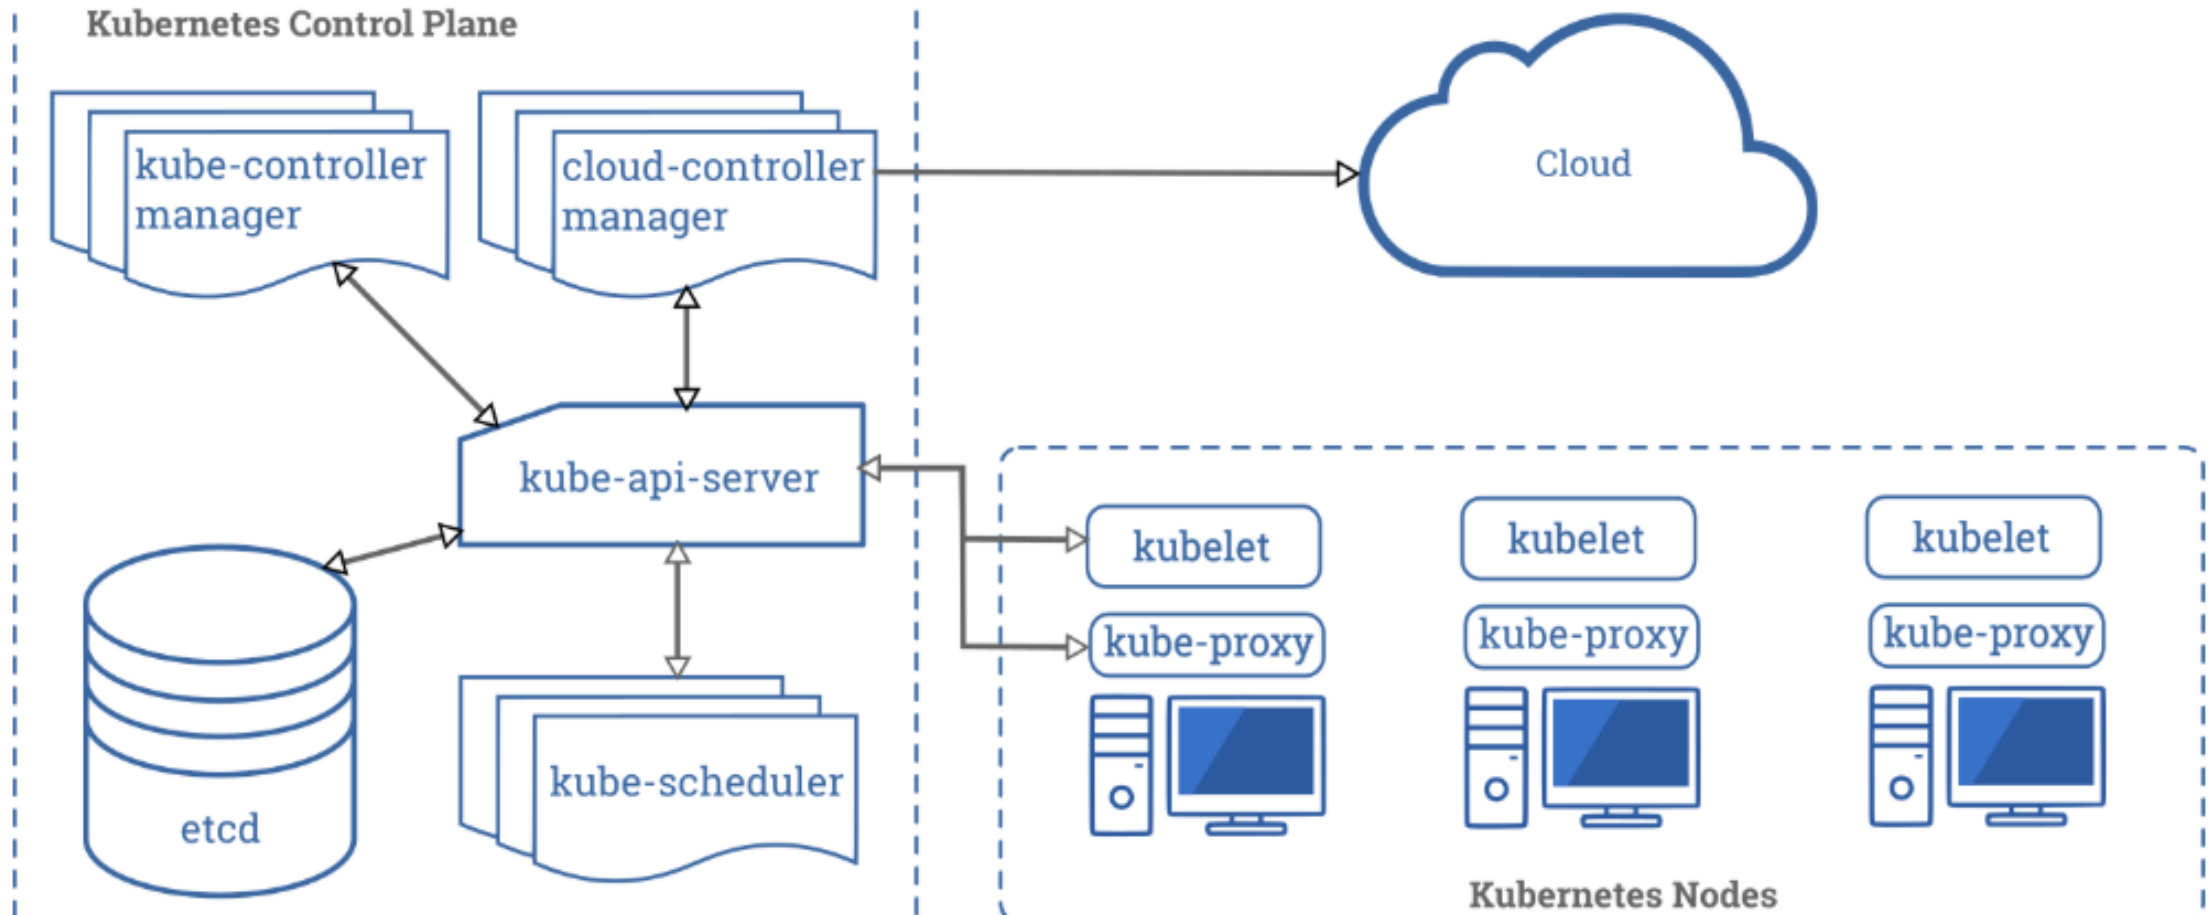
Shenjie Yao

x/xx/2020

Components

- Control Plane
 - api-server
 - kube-controller-manager
 - cloud-controller-manager
 - scheduler
 - etcd
- Node
 - kubelet
 - kube-proxy
 - Container-Runtime

Components

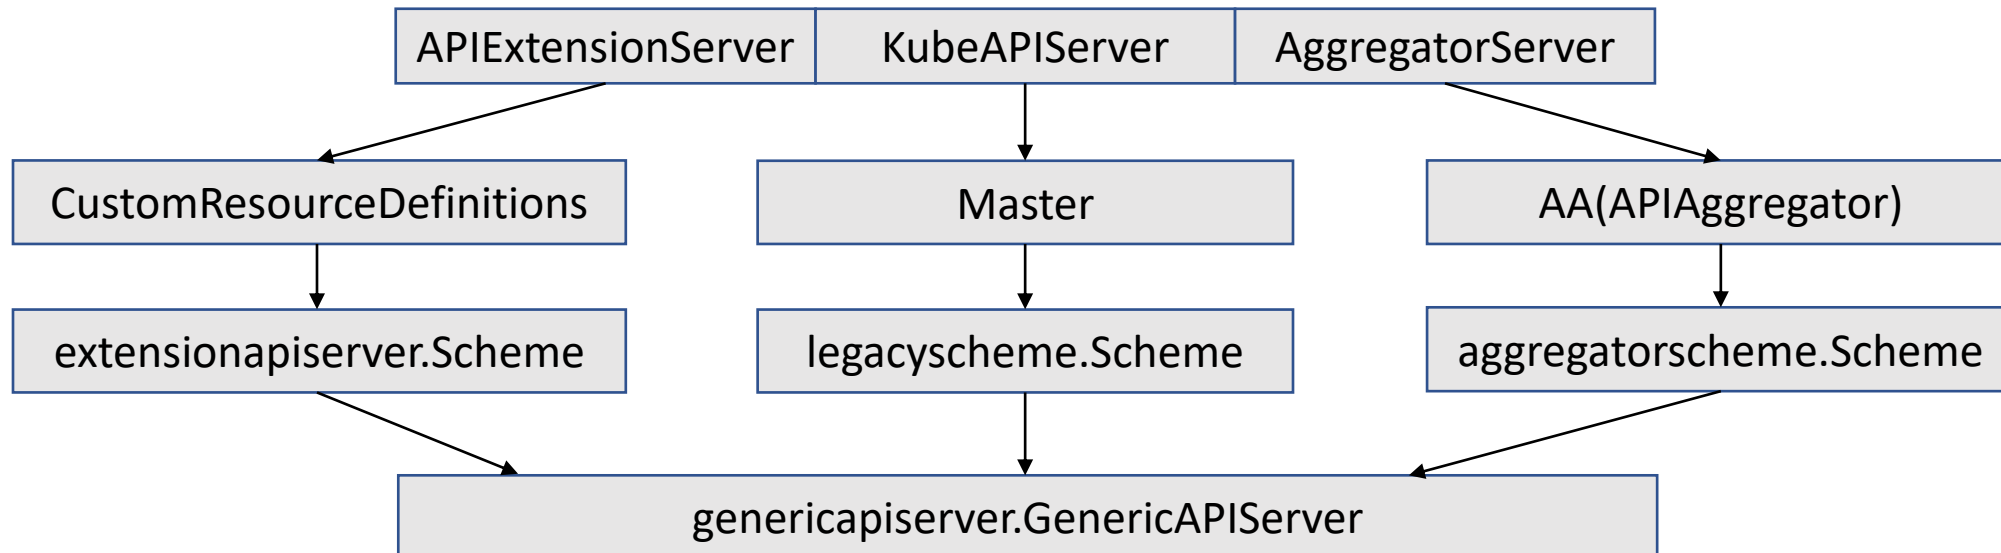


What is kube-apiserver?

- The API server is a component of the Kubernetes [control plane](#) that exposes the Kubernetes API. The API server is the front end for the Kubernetes control plane.
- The main implementation of a Kubernetes API server is [kube-apiserver](#). kube-apiserver is designed to scale horizontally—that is, it scales by deploying more instances. You can run several instances of kube-apiserver and balance traffic between those instances.

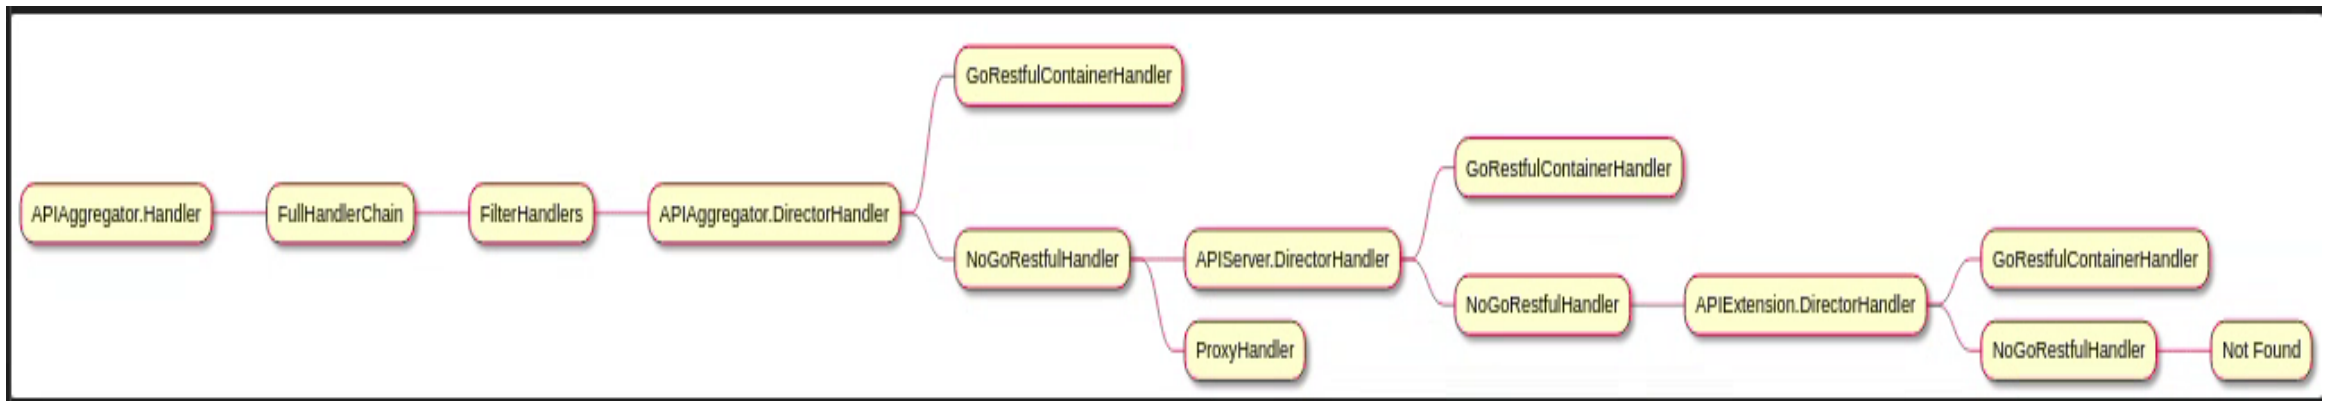
Kube-apiserver

- Kube-apiserver supplied three HTTP Server services for enriching ecosystem. These three HTTP Servers respectively is APIExtensionServer, KubeAPIServer and AggregatorServer. whatever which one Server is, it will depend on GenericAPIServer.
- Kube-apiserver's architecture design as following diagram:



Kube-apiserver

- This is an invoking chain for these three HTTP Servers:
 - AggregatorServer -> APIServer -> APIExtensionServer
- Following diagram shows that when received a HTTP request will invoke which one Server's handler to handle it.



APIAggregator

- [API Aggregation](#) requires programming, but allows more control over API behaviors like how data is stored and conversion between API versions.
- The aggregation layer allows Kubernetes to be extended with additional APIs, beyond what is offered by the core Kubernetes APIs. The additional APIs can either be ready-made solutions such as [service-catalog](#), or APIs that you develop yourself.
- The aggregation layer is different from [Custom Resources](#), which are a way to make the [kube-apiserver](#) recognise new kinds of object.
- The [aggregation layer](#) allows you to provide specialized implementations for your custom resources by writing and deploying your own standalone API server. The main API server delegates requests to you for the custom resources that you handle, making them available to all of its clients.

APIExtension

- The [CustomResourceDefinition](#) API resource allows you to define custom resources. Defining a CRD object creates a new custom resource with a name and schema that you specify. The Kubernetes API serves and handles the storage of your custom resource. The name of a CRD object must be a valid [DNS subdomain name](#).
- This frees you from writing your own API server to handle the custom resource, but the generic nature of the implementation means you have less flexibility than with [API server aggregation](#).
- Refer to the [custom controller example](#) for an example of how to register a new custom resource, work with instances of your new resource type, and use a controller to handle events.

APIAggregator VS APIExtension

- CRDs are easier to use. Aggregated APIs are more flexible. Choose the method that best meets your needs.
- Typically, CRDs are a good fit if:
 - You have a handful of fields
 - You are using the resource within your company, or as part of a small open-source project (as opposed to a commercial product)

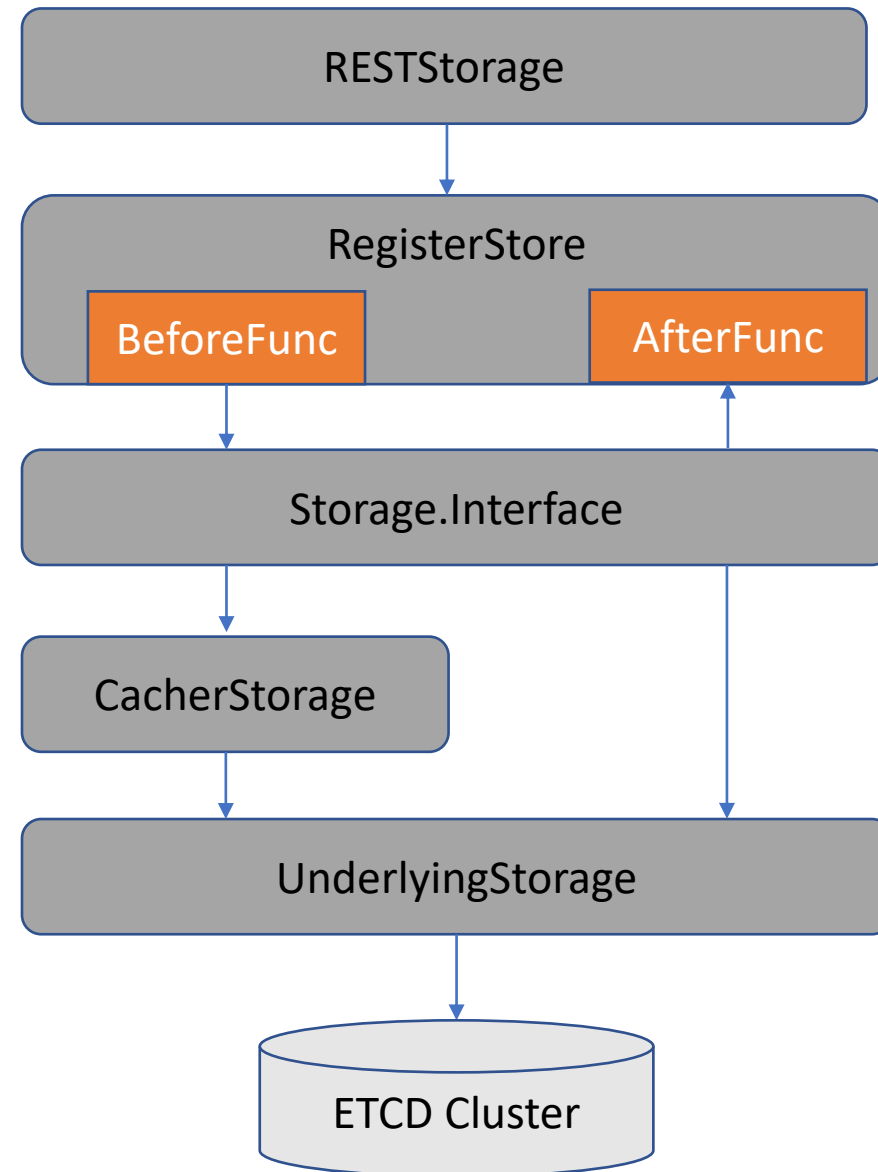
CRDs	Aggregated API
Do not require programming. Users can choose any language for a CRD controller.	Requires programming in Go and building binary and image.
No additional service to run; CRDs are handled by API server.	An additional service to create and that could fail.
No ongoing support once the CRD is created. Any bug fixes are picked up as part of normal Kubernetes Master upgrades.	May need to periodically pickup bug fixes from upstream and rebuild and update the Aggregated API server.
No need to handle multiple versions of your API; for example, when you control the client for this resource, you can upgrade it in sync with the API.	You need to handle multiple versions of your API; for example, when developing an extension to share with the world.

APIServer

- APIServer is a core service. It supplied Kubernetes in-built core resources services, and it doesn't allow developer to change relevant resources.
- API core services is managed by Master object. And through `legacyscheme.Scheme` resource register manage Master relevant resources.

ETCD Storage

- Kubernetes makes a plenty of encapsulation for ETCD storage. And it split this encapsulation to multiple layers, and each layer supports high extensibility.
- The right diagram shows relationship of among these layers

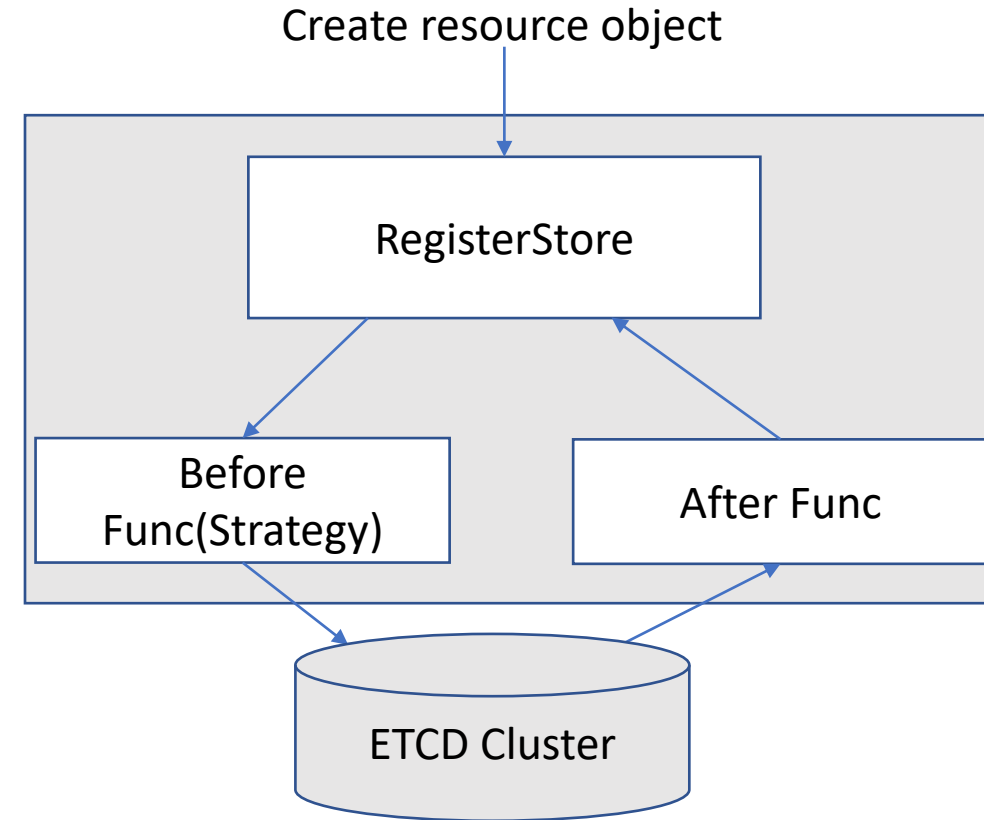


RESTStorage

- RESTStorage is a storage service general interface.
- So Every resources of the Kubernetes supplied RESTful API interface to external resources(RESTStorage interface). And all of resources which expose to external through RESTful API must implement this RESTStorage interface.
- Each RESTStorage interface encapsulated RegisterStore.

RegisterStore

- RegisterStore implemented general operation of the resource storage. Such as, execute some function before or after store resource object.
- Before Func(aka Strategy preprocess): it aims to do some preprocess works before store resource object.
- After Func: it aims to do some post-process works after store resource object.
- Last, it encapsulated Storage.Interface.



RegisterStore Strategy

- **CreateStrategy** : is a preprocess operation as creating a resource object. It defined a BeforeCreate function.
- **UpdateStrategy**: is a preprocess operation as updating a resource object. It defined a BeforeUpdate function.
- **DeleteStrategy**: is a preprocess operation as deleting a resource object. It defined a BeforeDelete function.
- **ExportStrategy**: is a preprocess operation as exporting a resource object. It defined an Export method. only some part of resources implemented this method.

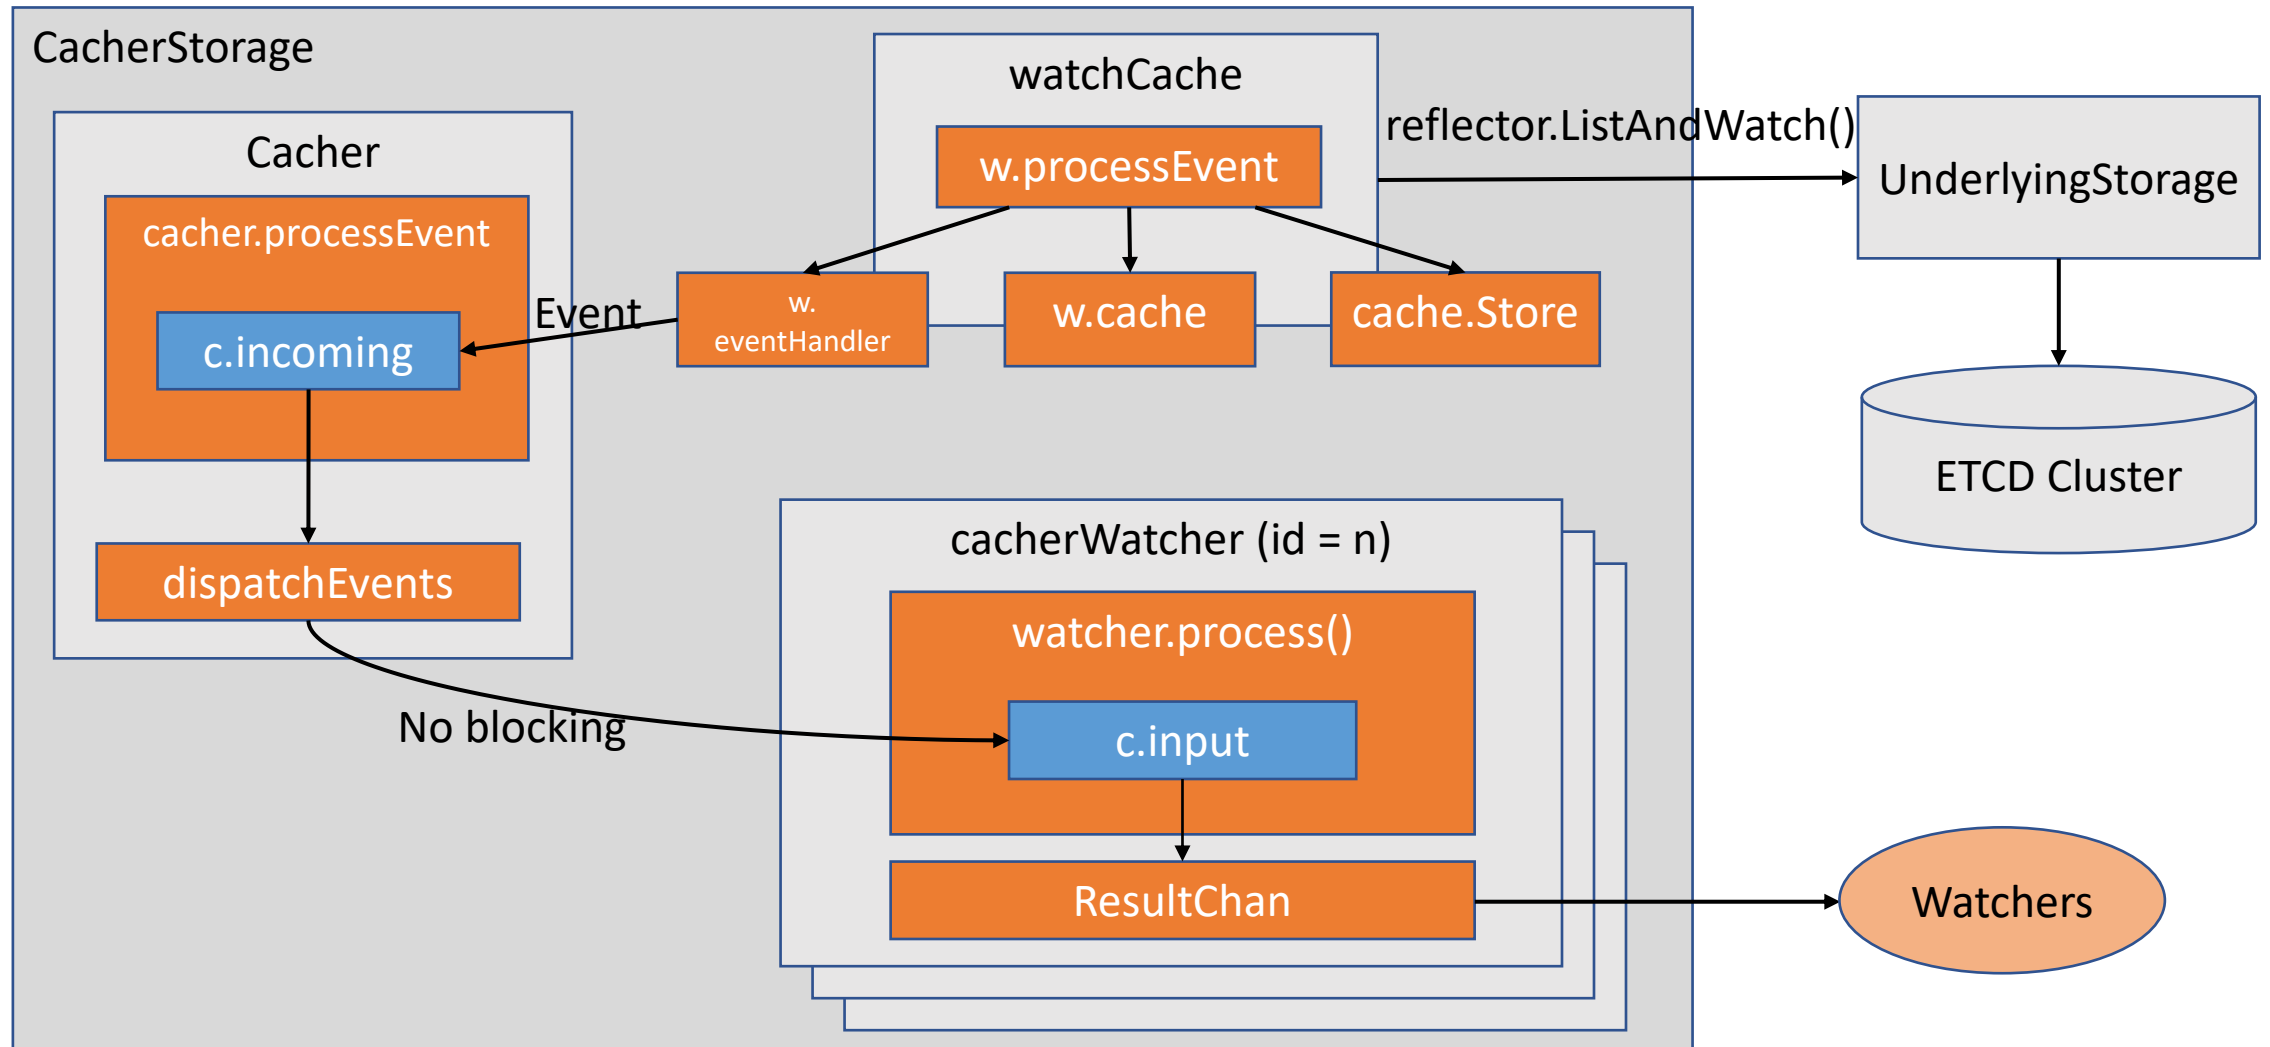
Storage.Interface

- Storage.Interface is a general storage interface and it defined operation method of the resources.
- CacherStorage and UnderlyingStorage implemented all of this interfaces.

CacherStorage

- CacherStorage supplied cache service of the DB data layer. It aims to quickly response and reduce pressure of the DB layer.
- The advantage of this design is that quickly response and return required data. It can reduce connection amount of the ETCD cluster and keep the data consistent between returned data and ETCD cluster.
- CacherStorage is not cache data for all of operations. For some operations ,in order to keep data consistent , so we don't necessary add this cacher layer. such as Create, Delete and Count operation, we can use UnderlyingStorage sending require to ETCD cluster directly.

CacherStorage



CacherStorage

- cacheWatcher: Watcher observer management.
- watchCache: interacting with UnderlyingStorage layer through Reflector framework, let UnderlyingStorage to interact directly with ETCD cluster, and respectively put callback event into w.eventHandler, w.cache and cache.Store.
- Cacher: dispatching events to all of observers who has connected, and It implemented by non-blocking mechanism in dispatching process.

UnderlyingStorage

- UnderlyingStorage also know as BackendStorage in the code. It is a real resource storage object interacting with ETCD cluster.
- UnderlyingStorage encapsulated official ETCD client library. And it also is an implementation of Storage.Interface general storage interface.

References

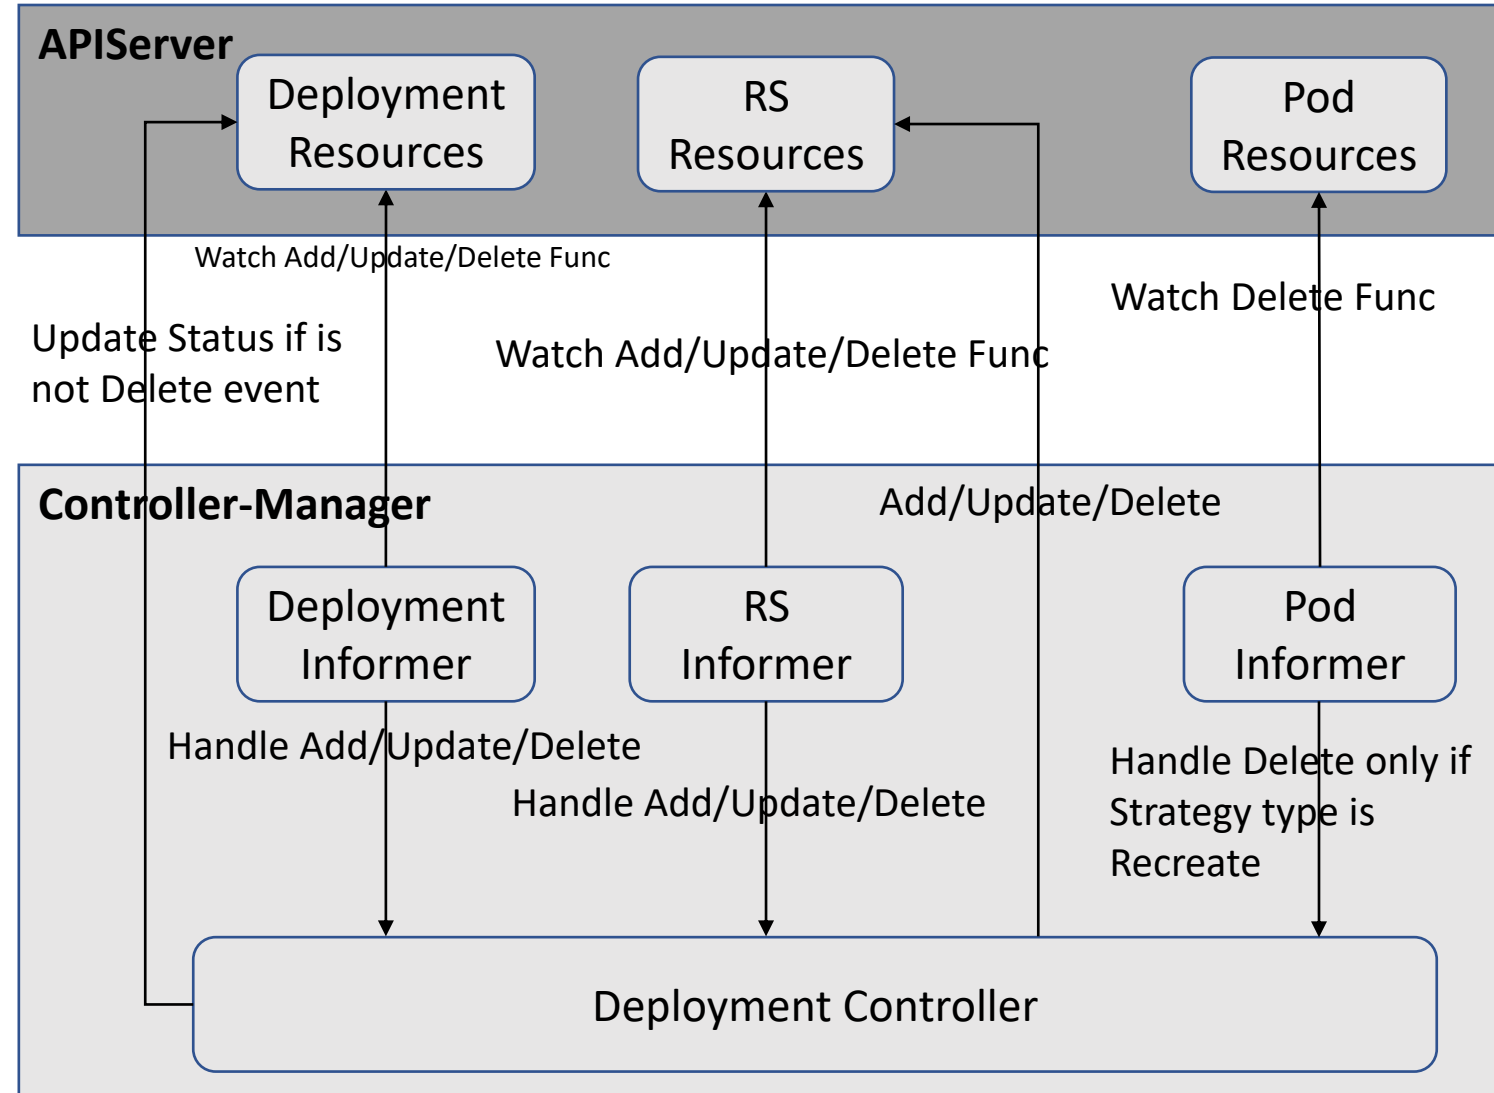
- <https://kubernetes.io/docs/reference/access-authn-authz/controlling-access/>
- <https://kubernetes.io/docs/concepts/extend-kubernetes/api-extension/apiserver-aggregation/>
- <https://kubernetes.io/docs/tasks/extend-kubernetes/custom-resources/custom-resource-definitions/>
- <https://kubernetes.io/docs/concepts/extend-kubernetes/api-extension/custom-resources/>
- <https://kubernetes.io/docs/concepts/extend-kubernetes/>

Controller-manager

- Kube-controller-manager
 - Control Plane component that runs [controller](#) (**a control loop that watches the shared state of the cluster through the api-server and make changes attempting to move current state towards desired state**) processes.
 - Logically, each [controller](#) is a separate process, but to reduce complexity, they are all compiled into a single binary and run in a single process.
- Cloud-controller-manager

Deployment Controller

- A *Deployment* provides declarative updates for [Pods](#) [ReplicaSets](#).
- You describe a *desired state* in a Deployment, and the Deployment [Controller](#) changes the actual state to the *desired state* at a controlled rate. You can define Deployments to create new ReplicaSets, or to remove existing Deployments and adopt all their resources with new Deployments.
- Deployment ensures that only a certain number of Pods are down while they are being updated. By default, it ensures that at least 75% of the desired number of Pods are up (25% max unavailable).
- Deployment also ensures that only a certain number of Pods are created above the desired number of Pods. By default, it ensures that at most 125% of the desired number of Pods are up (25% max surge).
- <https://kubernetes.io/docs/concepts/workloads/controllers/deployment/#creating-a-deployment>

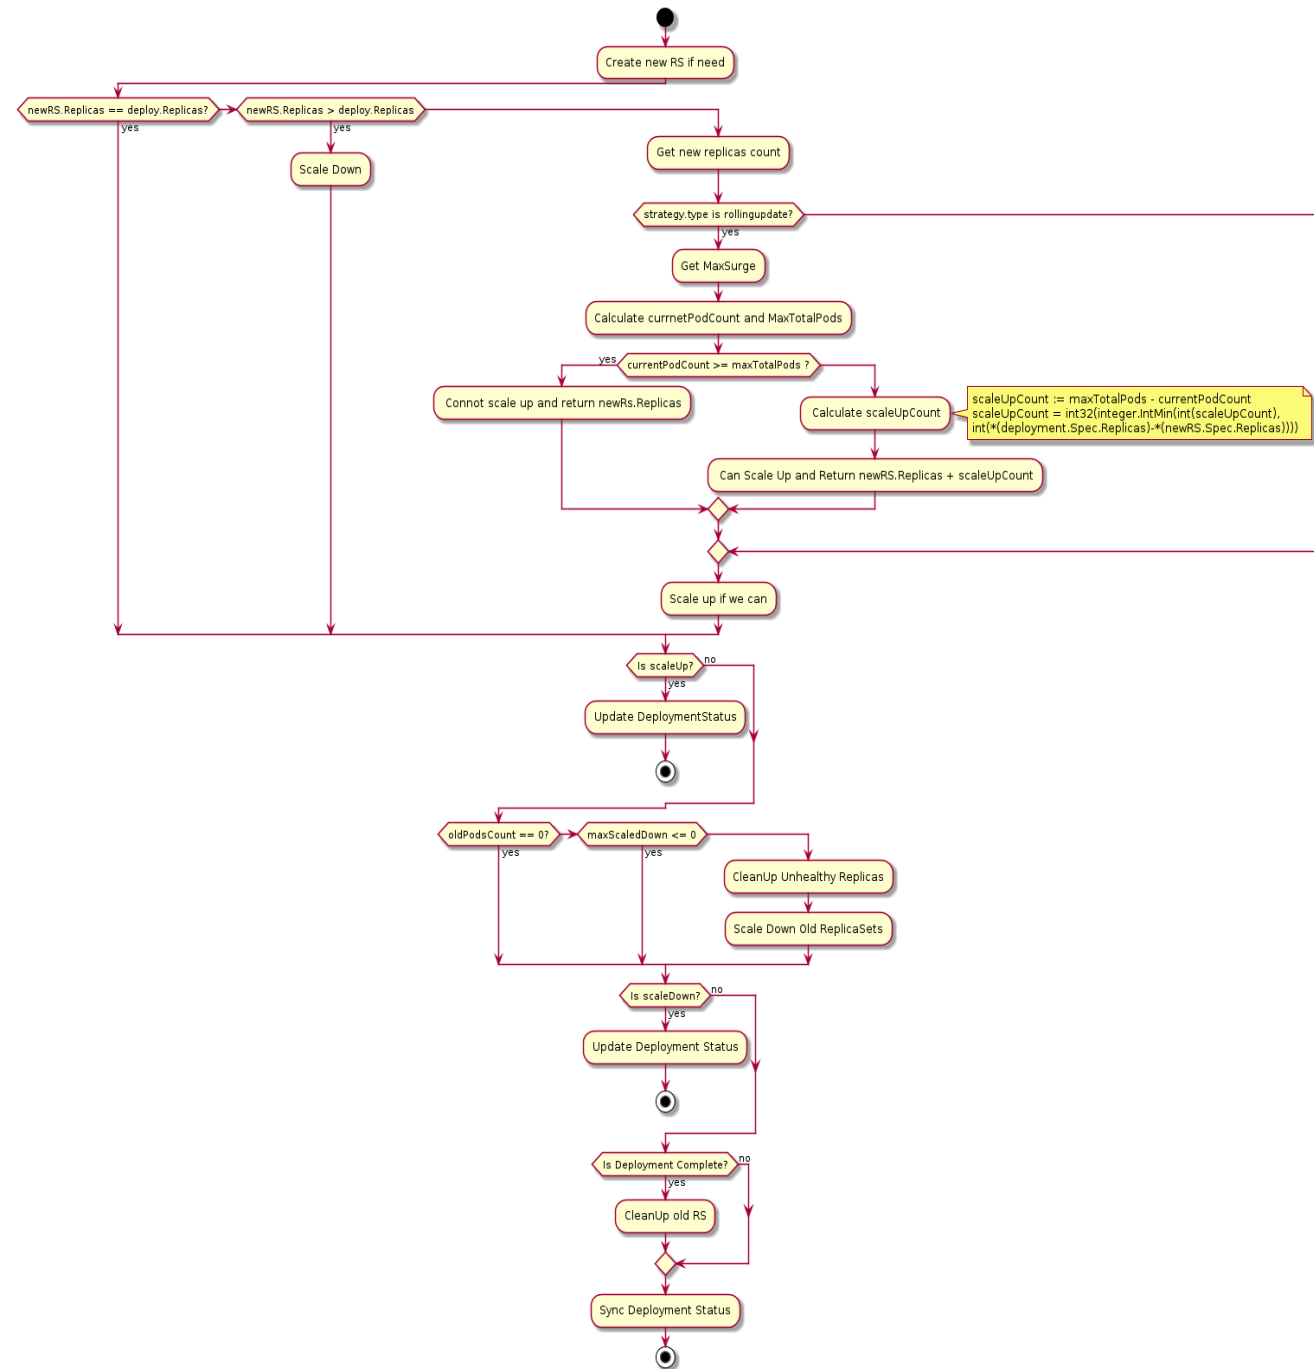


Deployment Controller

• Strategy Type

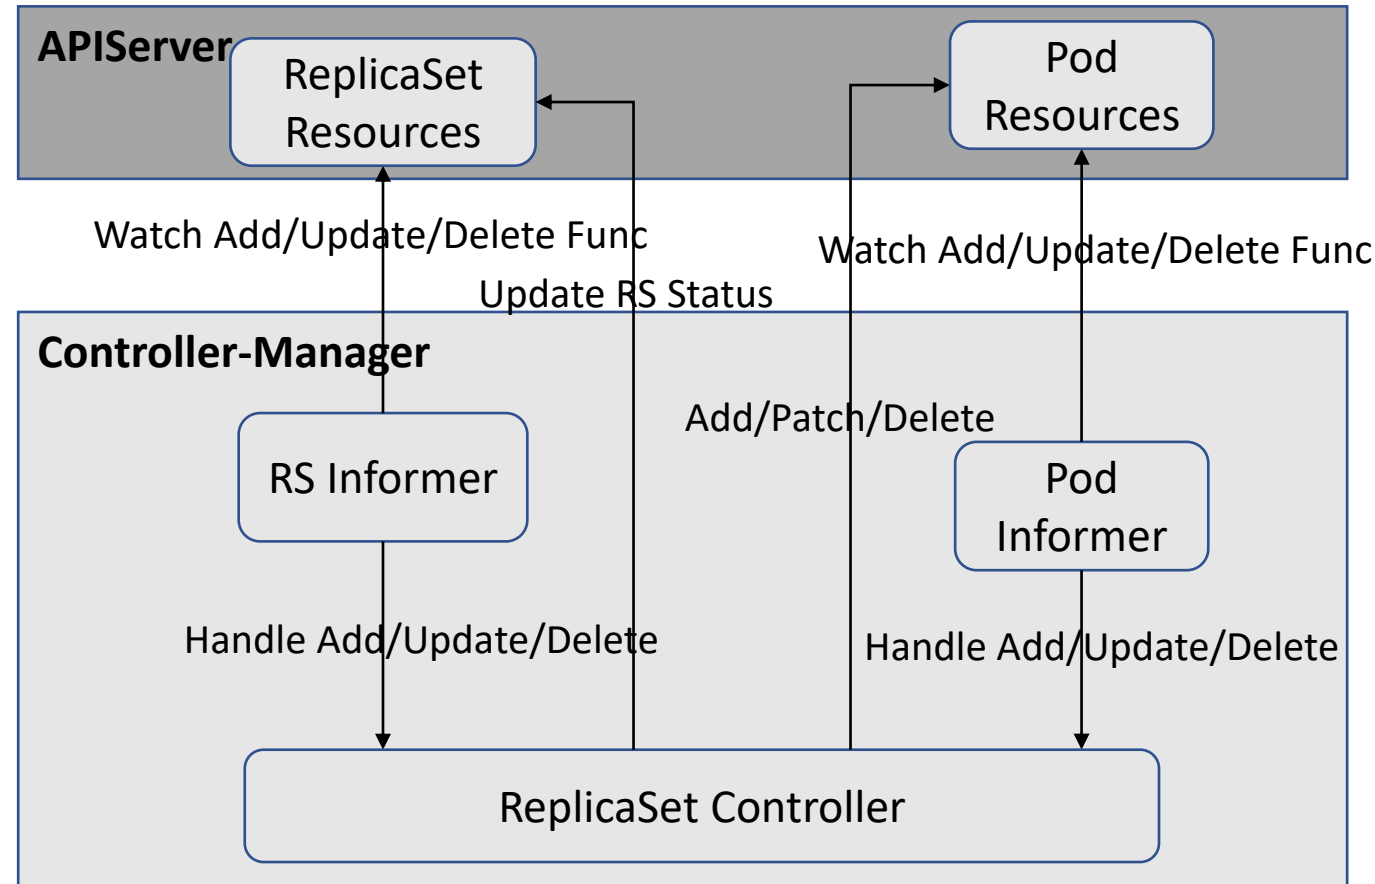
- Recreate – scale down old replica set first then scale up new replica set.
- RollingUpdate – scale up the new replica set if we can or scale down the old replica set if we can.

$\text{maxScaleDown} = \text{allPodsCount} -$
 $(*(\text{deployment.Spec.Replicas}) -$
 $\text{maxUnavailable}) -$
 $(*(\text{newRS.Spec.Replicas}) -$
 $\text{newRS.Status.AvailableReplicas})$



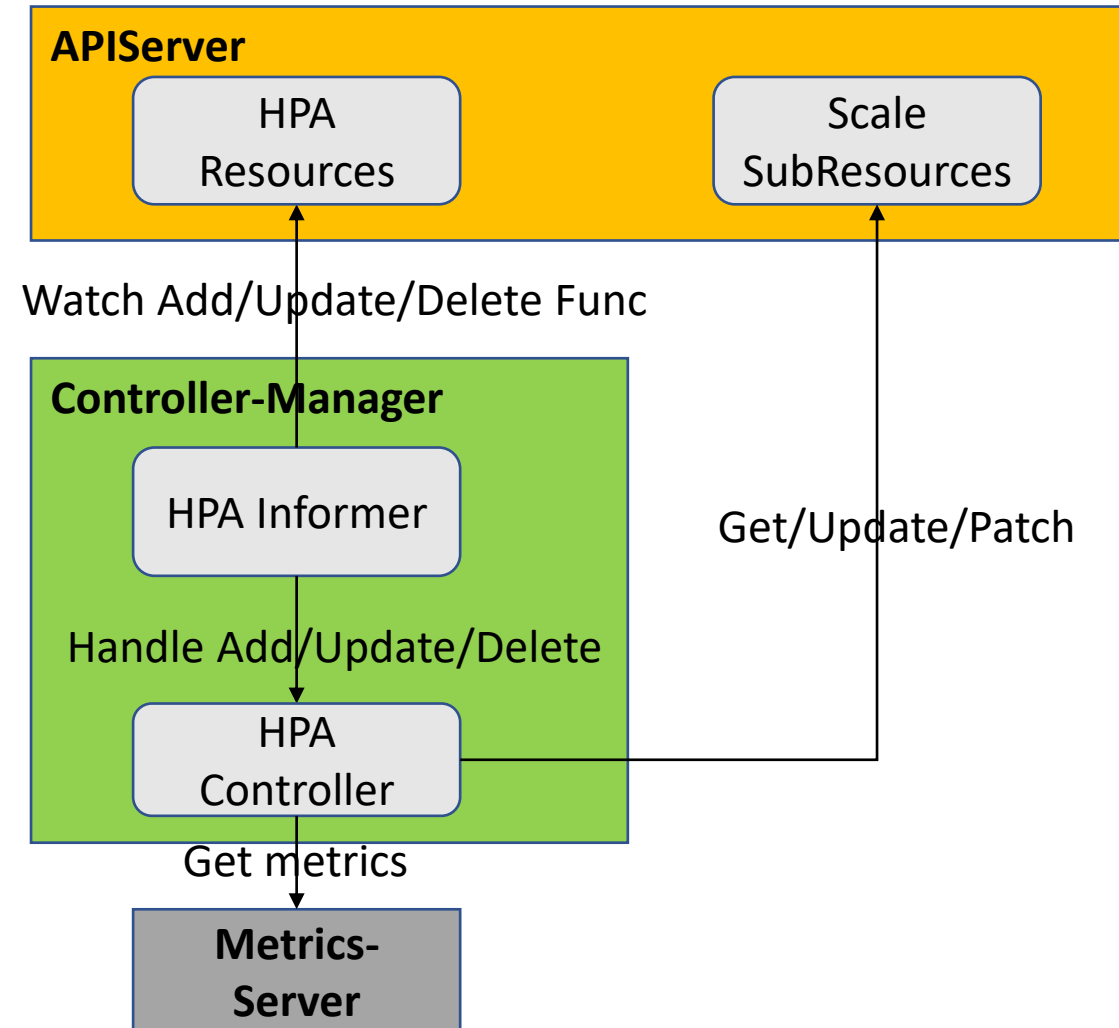
ReplicaSet Controller

- A ReplicaSet's purpose is to maintain a stable set of replica Pods running at any given time. As such, it is often used to guarantee the availability of a specified number of identical Pods.
- A ReplicaSet is defined with fields, including a selector that specifies how to identify Pods it can acquire, a number of replicas indicating how many Pods it should be maintaining, and a pod template specifying the data of new Pods it should create to meet the number of replicas criteria. A ReplicaSet then fulfills its purpose by creating and deleting Pods as needed to reach the desired number. When a ReplicaSet needs to create new Pods, it uses its Pod template.
- A ReplicaSet is linked to its Pods via the Pods' [metadata.ownerReferences](#) field, which specifies what resource the current object is owned by. All Pods acquired by a ReplicaSet have their owning ReplicaSet's identifying information within their ownerReferences field. It's through this link that the ReplicaSet knows of the state of the Pods it is maintaining and plans accordingly.
- A ReplicaSet identifies new Pods to acquire by using its selector. If there is a Pod that has no OwnerReference or the OwnerReference is not a [Controller](#) and it matches a ReplicaSet's selector, it will be immediately acquired by said ReplicaSet.
- <https://kubernetes.io/docs/concepts/workloads/controllers/replicaset/>



PodAutoScaler Controller

- The Horizontal Pod Autoscaler automatically scales the number of pods in a replication controller, deployment, replica set or stateful set based on observed CPU utilization (or, with [custom metrics](#) support, on some other application-provided metrics). Note that Horizontal Pod Autoscaling does not apply to objects that can't be scaled, for example, DaemonSets.
- The Horizontal Pod Autoscaler is implemented as a Kubernetes API resource and a controller. The resource determines the behavior of the controller. The controller periodically adjusts the number of replicas in a replication controller or deployment to match the observed average CPU utilization to the target specified by user.
- Algorithm details please reference this <https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale/#algorithm-details>.

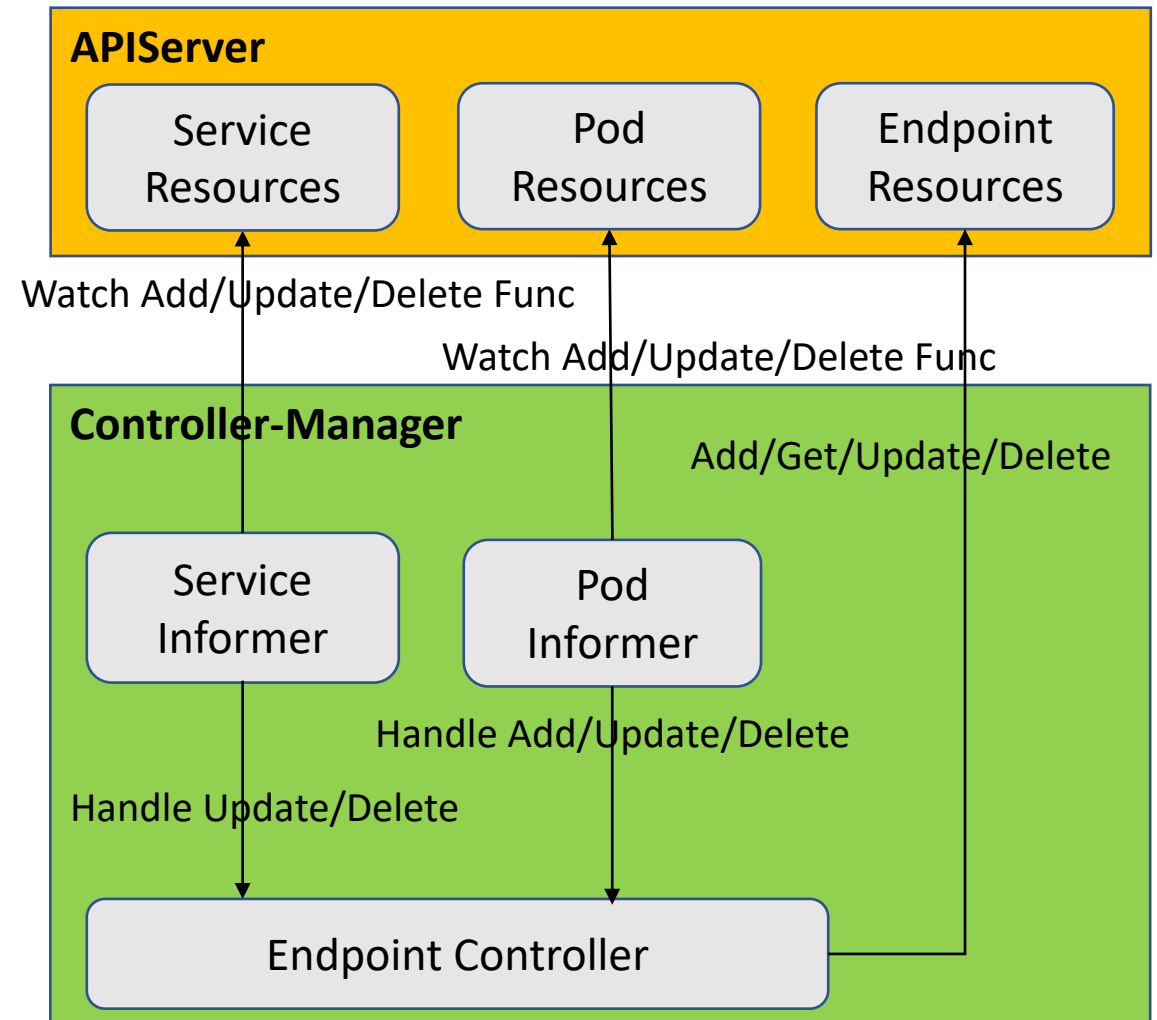


Service Controller

- TODO

Endpoint Controller

- Endpoint Controller is a mainly method to create and delete endpoint resource. It will obtain all of relevant pods according to Selector in the Service object. And mapping port in the between Service and Pod to generate a EndpointPort struct.
- Endpoint Controller will generate a new EndpointSubset for each Pod, it including Pod's IP and Port, and import Port and target Port in the Service. And at last it will according to EndpointSubset to repackage and to create a new endpoint resource.



Garbage Collector Controller

- TODO

Custom controller

- Using **client-go** to write a controller directly.
 - This is an official example for using client-go to write a controller.
<https://github.com/kubernetes/sample-controller>.
- Using **controller-runtime** to write a controller.
 - TODO
- Using **kubebuilder** or **operator SDK** to generate a controller.
 - These two tools can generate a CRD and its controller. Hence if you want to write a CRD, you can choose one of these two tools. But if you only want to write a controller to control k8s resources. I recommend you directly use client-go or controller-runtime.

References

- <https://kubernetes.io/docs/concepts/architecture/controller/>
- <https://kubernetes.io/docs/concepts/architecture/cloud-controller/>
- <https://github.com/kubernetes/sample-controller>
- <https://kubernetes.io/docs/concepts/workloads/controllers/>
- <https://kubernetes.io/docs/tasks/run-application/horizontal-pod-autoscale/>

Scheduler

- In Kubernetes, *scheduling* refers to making sure that [Pods](#) are matched to [Nodes](#) so that [Kubelet](#) can run them.
- A scheduler watches for newly created Pods that have no Node assigned. For every Pod that the scheduler discovers, the scheduler becomes responsible for finding the best Node for that Pod to run on.
- [kube-scheduler](#) is the default scheduler for Kubernetes and runs as part of the [control plane](#). kube-scheduler is designed so that, if you want and need to, you can write your own scheduling component and use that instead.
- For every newly created pod or other unscheduled pods, kube-scheduler selects an optimal node for them to run on. However, every container in pods has different requirements for resources and every pod also has different requirements. Therefore, existing nodes need to be filtered according to the specific scheduling requirements.
- In a cluster, Nodes that meet the scheduling requirements for a Pod are called *feasible* nodes. If none of the nodes are suitable, the pod remains unscheduled until the scheduler is able to place it.
- The scheduler finds feasible Nodes for a Pod and then runs a set of functions to score the feasible Nodes and picks a Node with the highest score among the feasible ones to run the Pod. The scheduler then notifies the API server about this decision in a process called *binding*.
- Factors that need taken into account for scheduling decisions include individual and collective resource requirements, hardware / software / policy constraints, affinity and anti-affinity specifications, data locality, inter-workload interference, and so on.

Scheduler

- kube-scheduler selects a node for a pod in following 2-step operation:
 - Filtering
 - Scoring
- The **filtering** step finds the set of Nodes where it's feasible to schedule the Pod. For example, the PodFitsResources filter checks whether a candidate Node has enough available resource to meet a Pod's specific resource requests. After this step, the node list contains any suitable Nodes; often, there will be more than one. If the list is empty, that Pod isn't (yet) schedulable.
- The **scoring** step, the scheduler ranks the remaining nodes to choose the most suitable Pod placement. The scheduler assigns a score to each Node that survived filtering, basing this score on the active scoring rules.
- Finally, kube-scheduler assigns the Pod to the Node with the highest ranking. If there is more than one node with equal scores, kube-scheduler selects one of these at random.

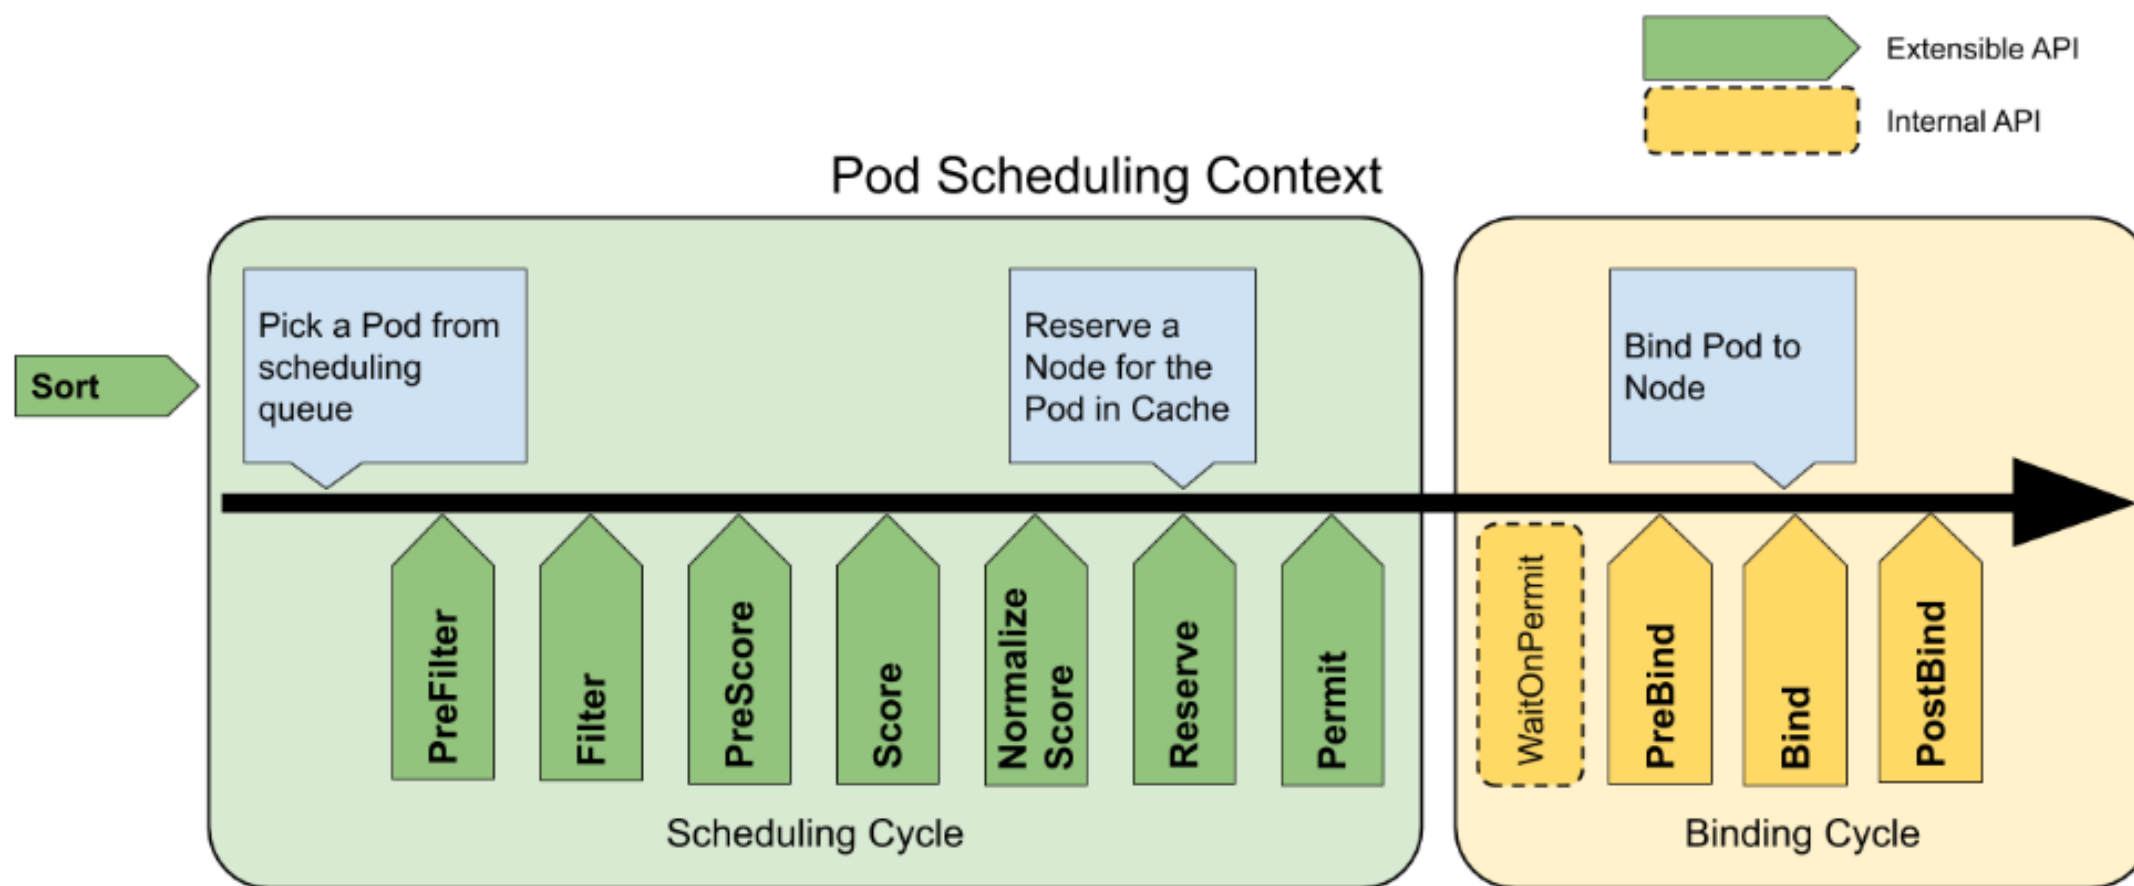
Scheduler

- There are two supported ways to configure the filtering and scoring behavior of the scheduler:
- [Scheduling Policies](#) allow you to configure *Predicates* for filtering and *Priorities* for scoring.
- [Scheduling Profiles](#) allow you to configure Plugins that implement different scheduling stages, including : QueueSort, Filter, Score, Bind, Reserve, Permit, and others. You can also configure the kube-scheduler to run different profiles.

Scheduling Framework

- The scheduling framework is a pluggable architecture for Kubernetes Scheduler that makes scheduler customizations easy. It adds a new set of "plugin" APIs to the existing scheduler. Plugins are compiled into the scheduler. The APIs allow most scheduling features to be implemented as plugins, while keeping the scheduling "core" simple and maintainable. Refer to the [design proposal of the scheduling framework](#) for more technical information on the design of the framework.
- The Scheduling Framework defines a few extension points. Scheduler plugins register to be invoked at one or more extension points. Some of these plugins can change the scheduling decisions and some are informational only.
- Each attempt to schedule one Pod is split into two phases, the **scheduling cycle** and the **binding cycle**.

Extension Points



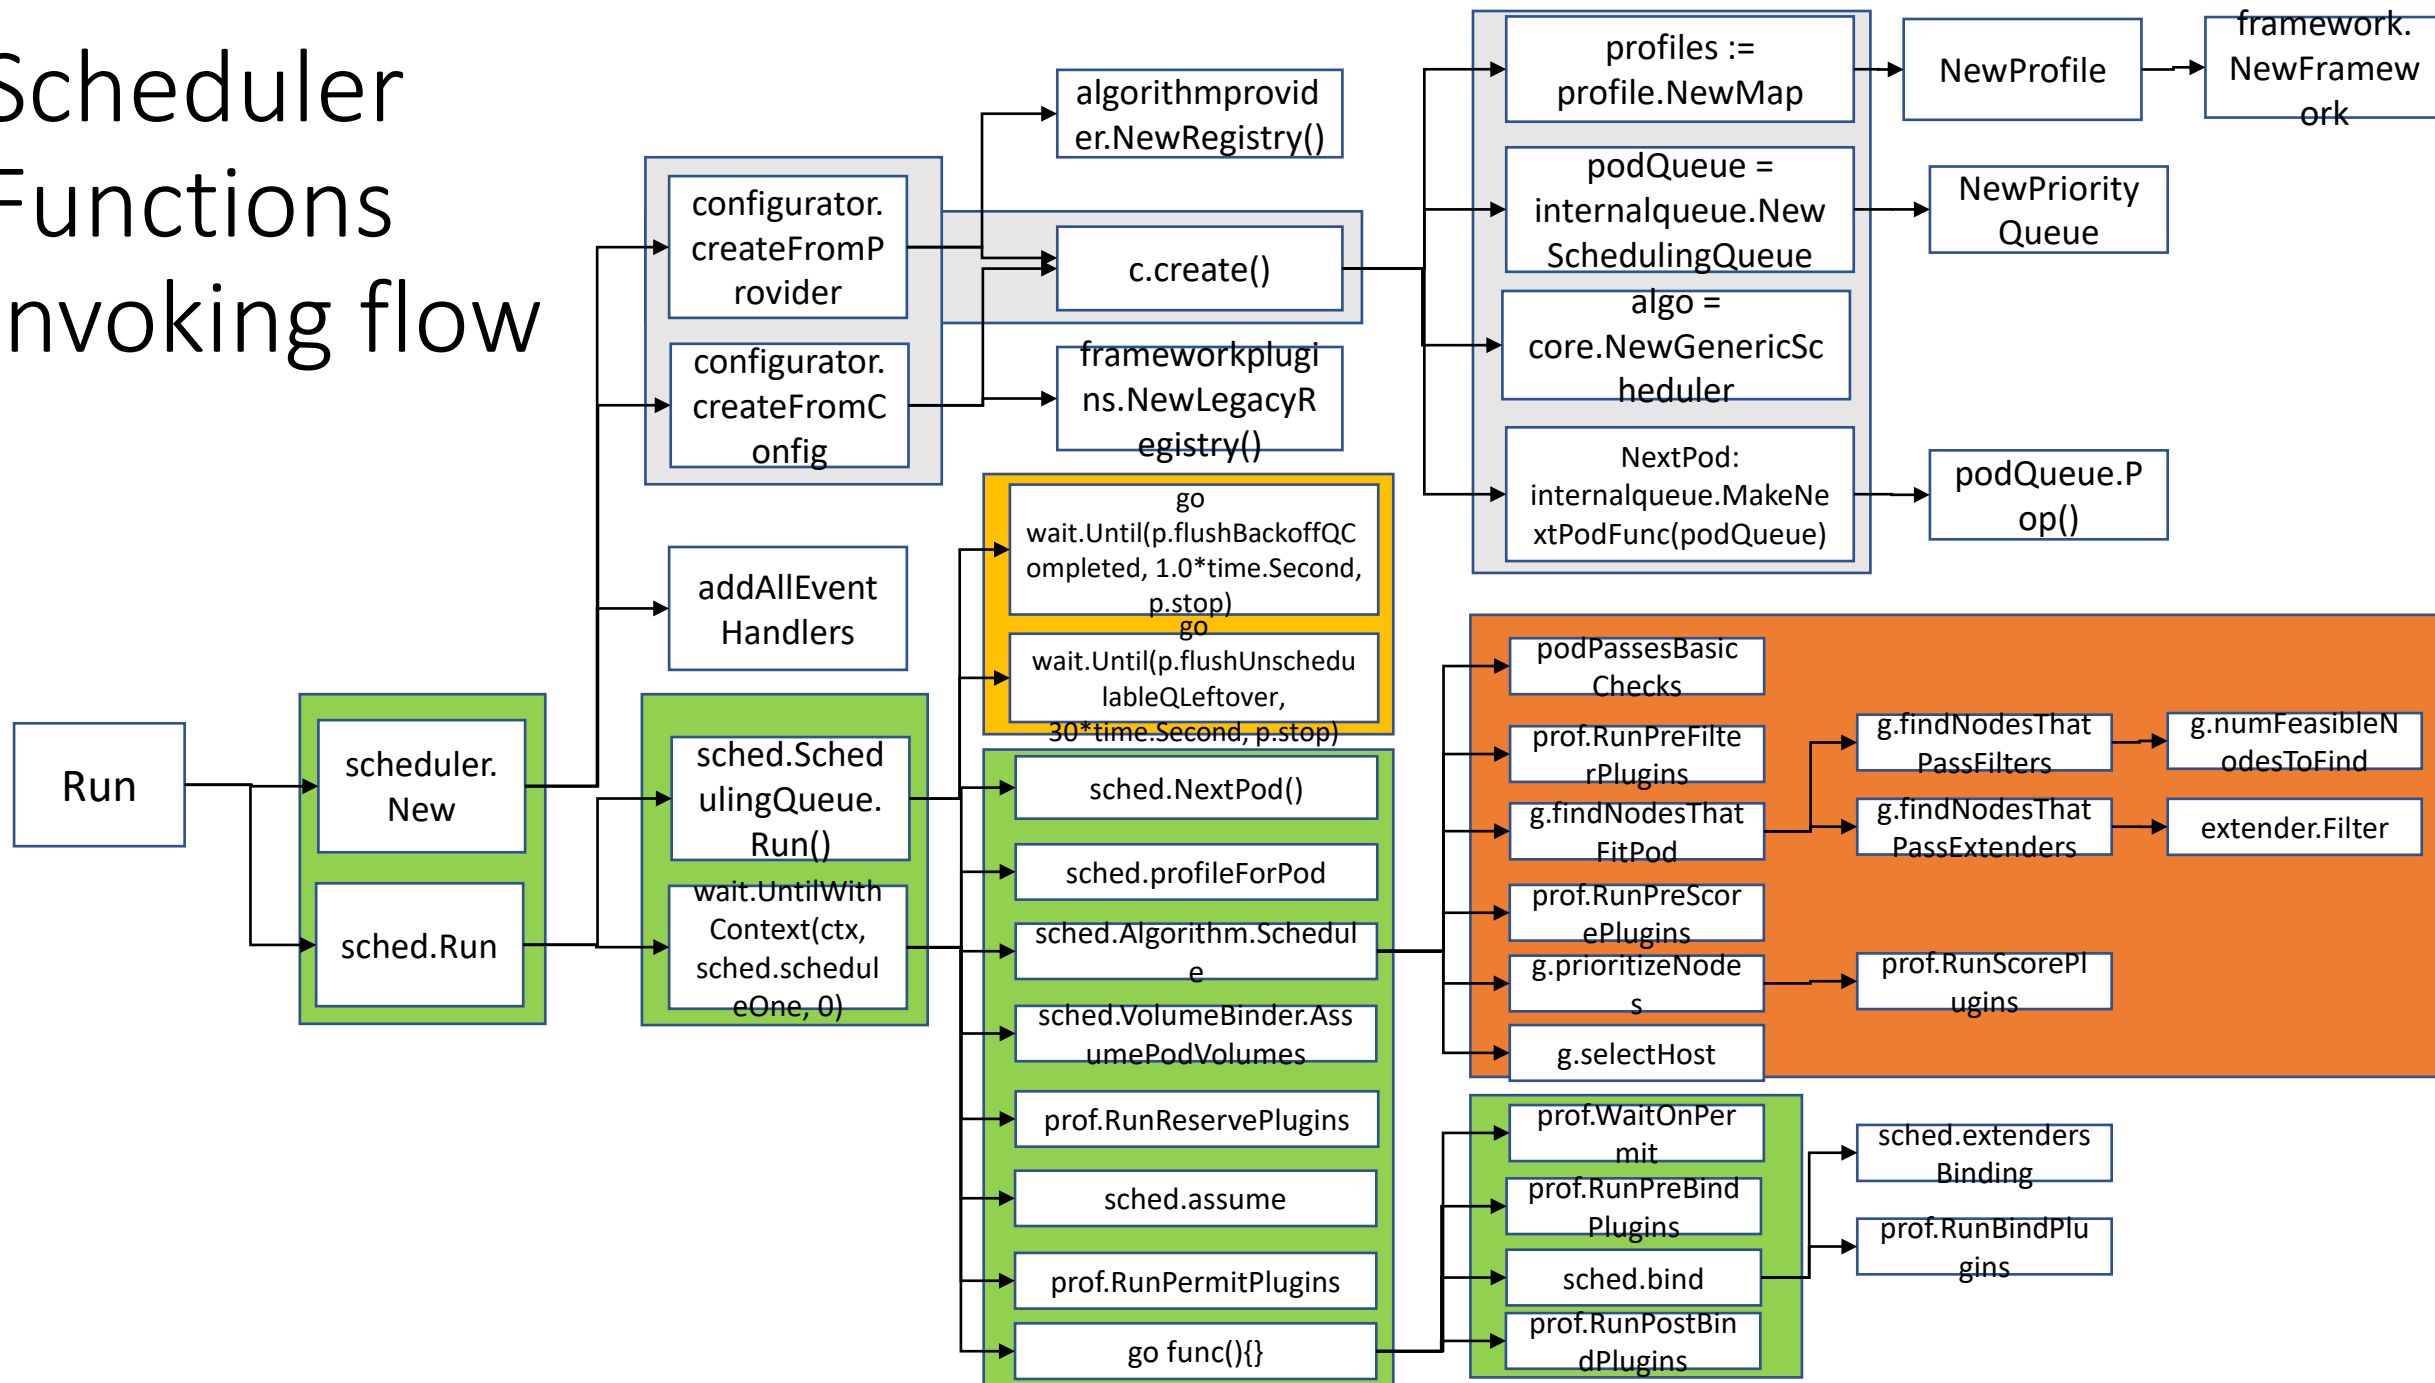
Pod Scheduling Context

- **Queue sort** - These plugins are used to sort Pods in the scheduling queue. A queue sort plugin essentially provides a `Less(Pod1, Pod2)` function. Only one queue sort plugin may be enabled at a time.
- **PreFilter** - These plugins are used to pre-process info about the Pod, or to check certain conditions that the cluster or the Pod must meet. If a PreFilter plugin returns an error, the scheduling cycle is aborted.
- **Filter** - These plugins are used to filter out nodes that cannot run the Pod. For each node, the scheduler will call filter plugins in their configured order. If any filter plugin marks the node as infeasible, the remaining plugins will not be called for that node. Nodes may be evaluated concurrently.
- **PreScore** - These plugins are used to perform "pre-scoring" work, which generates a sharable state for Score plugins to use. If a PreScore plugin returns an error, the scheduling cycle is aborted.
- **Score** - These plugins are used to rank nodes that have passed the filtering phase. The scheduler will call each scoring plugin for each node. There will be a well defined range of integers representing the minimum and maximum scores. After the [NormalizeScore](#) phase, the scheduler will combine node scores from all plugins according to the configured plugin weights.
- **Reserve** - This is an informational extension point. Plugins which maintain runtime state (aka "stateful plugins") should use this extension point to be notified by the scheduler when resources on a node are being reserved for a given Pod. This happens before the scheduler actually binds the pod to the Node, and it exists to prevent race conditions while the scheduler waits for the bind to succeed.
- **Permit** - *Permit* plugins are invoked at the end of the scheduling cycle for each Pod, to prevent or delay the binding to the candidate node.
- **PreBind** - These plugins are used to perform any work required before a Pod is bound. For example, a pre-bind plugin may provision a network volume and mount it on the target node before allowing the Pod to run there.
- **Bind** - These plugins are used to bind a Pod to a Node. Bind plugins will not be called until all PreBind plugins have completed. Each bind plugin is called in the configured order. A bind plugin may choose whether or not to handle the given Pod. If a bind plugin chooses to handle a Pod, **the remaining bind plugins are skipped**.
- **PostBind** - This is an informational extension point. Post-bind plugins are called after a Pod is successfully bound. This is the end of a binding cycle, and can be used to clean up associated resources.

Default Schedule Algorithms

- Defaultpodtopologyspread
- Imagelocality
- Tainttoleration
- Nodename
- Nodeports
- Nodepreferavoidpods
- Nodeaffinity
- Podtopologyspread
- Nodeunschedulable
- Noderesources
 - NodeResourcesFit
 - NodeResourcesBalancedAllocation
 - NodeResourcesLeastAllocated
 - NodeResourcesMostAllocated
 - RequestedToCapacityRatio
 - NodeResourceLimits
- Volumebinding
- Volumerestrictions
- Volumezone
- Nodevolumelimits
 - NodeVolumeLimits(CSI)
 - AzureDiskLimits
 - CinderLimits
 - EBSLimits
 - GCEPDLimits
- Interpodaffinity
- Nodelabel
- Serviceaffinity
- Queuesort
- Defaultbinder

Scheduler Functions invoking flow



Sample-Scheduler

- If you want to write a custom scheduler, you can implement some of following interfaces defined by Kubernetes.
<https://github.com/kubernetes/kubernetes/blob/master/pkg/scheduler/framework/v1alpha1/interface.go>. Such as you can reference these examples provided by k8s scheduler framework resource <https://github.com/kubernetes/kubernetes/blob/release-1.19/pkg/scheduler/framework/plugins/examples>. And also you can reference this example to implement a custom scheduler <https://github.com/Luffy110/k8s-studies/tree/master/playground/sample-scheduler>.

References

- <https://kubernetes.io/docs/reference/scheduling/policies/>
- <https://kubernetes.io/docs/concepts/scheduling-eviction/kube-scheduler/>
- <https://github.com/kubernetes/enhancements/blob/master/keps/sig-scheduling/20180409-scheduling-framework.md>
- <https://www.servicemeshher.com/blog/202003-k8s-scheduling-framework/>
- <https://www.qikqiak.com/post/custom-kube-scheduler/>
- <https://github.com/kubernetes/enhancements/blob/master/keps/sig-scheduling/624-scheduling-framework/README.md>

Kube-proxy

- Kube-proxy is a core component of the K8s, it be deployed in every Node to implement communication and load balancing of the between services.
- Kube-proxy is mainly responsible for creating proxy service for pod, and obtaining all of services information from APIServer, then according to these information to create relevant proxy service so that to implement request route and forward from Server to Pod.
- Now It supports **user space** proxy mode, **iptables** proxy mode and **ipvs** proxy mode.

What is Iptables?

- **iptables** is the user space command line program used to configure the Linux 2.4.x and later packet filtering ruleset. It is targeted towards system administrators.
- Since Network Address Translation(NAT) is also configured from the packet filter ruleset, **iptables** is used for this, too.
- Features:
 - listing the contents of the packet filter ruleset
 - adding/removing/modifying rules in the packet filter ruleset
 - listing/zeroing per-rule counters of the packet filter ruleset

What is IPVS?

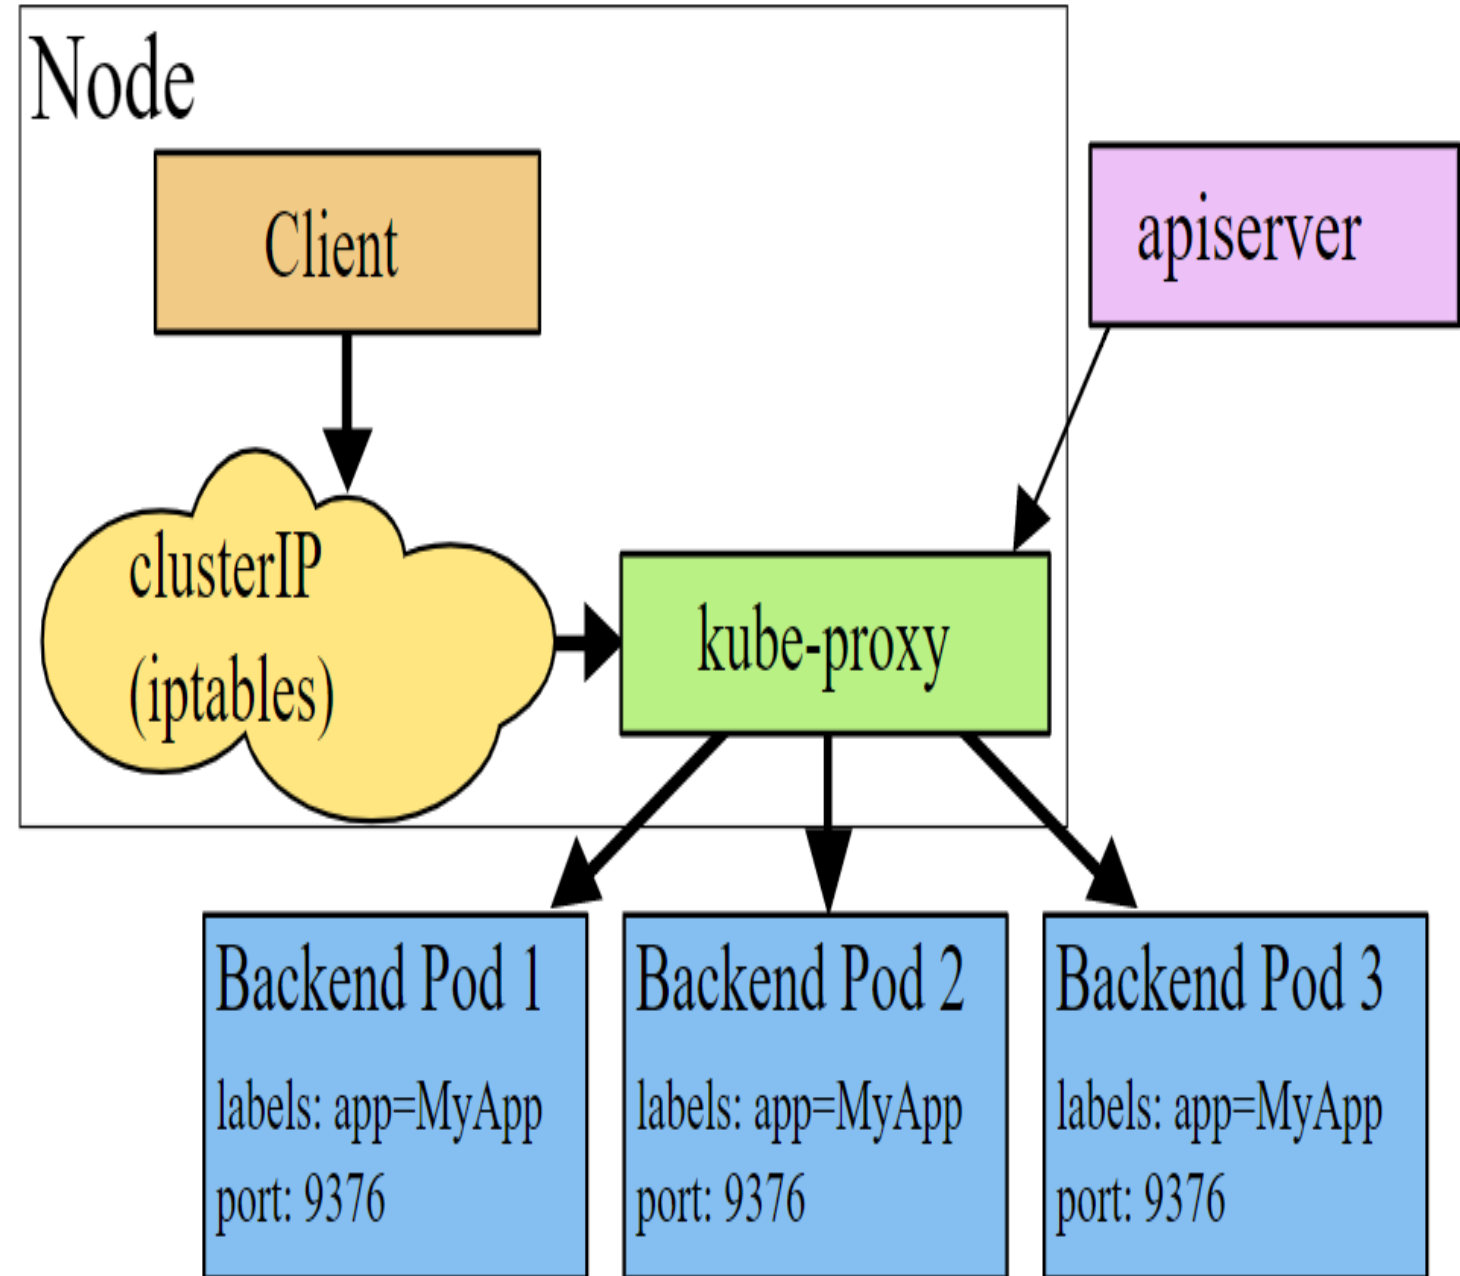
- **IPVS (IP Virtual Server)** implements transport-layer [load balancing](#), usually called [Layer 4 LAN switching](#), as part of the [Linux kernel](#). It's configured via the user-space utility `ipvsadm(8)` tool.
- IPVS is incorporated into the [Linux Virtual Server](#) (LVS), where it runs on a host and acts as a load balancer in front of a cluster of real servers. IPVS can directly request for TCP- and UDP-based services to the real servers, and making services of the real servers appear as virtual services on a single [IP address](#). IPVS is built on top of the [Netfilter](#).^[1]

What is Conntrack?

- Please reference this blog: <http://arthurchiao.art/blog/conntrack-design-and-implementation-zh/>.

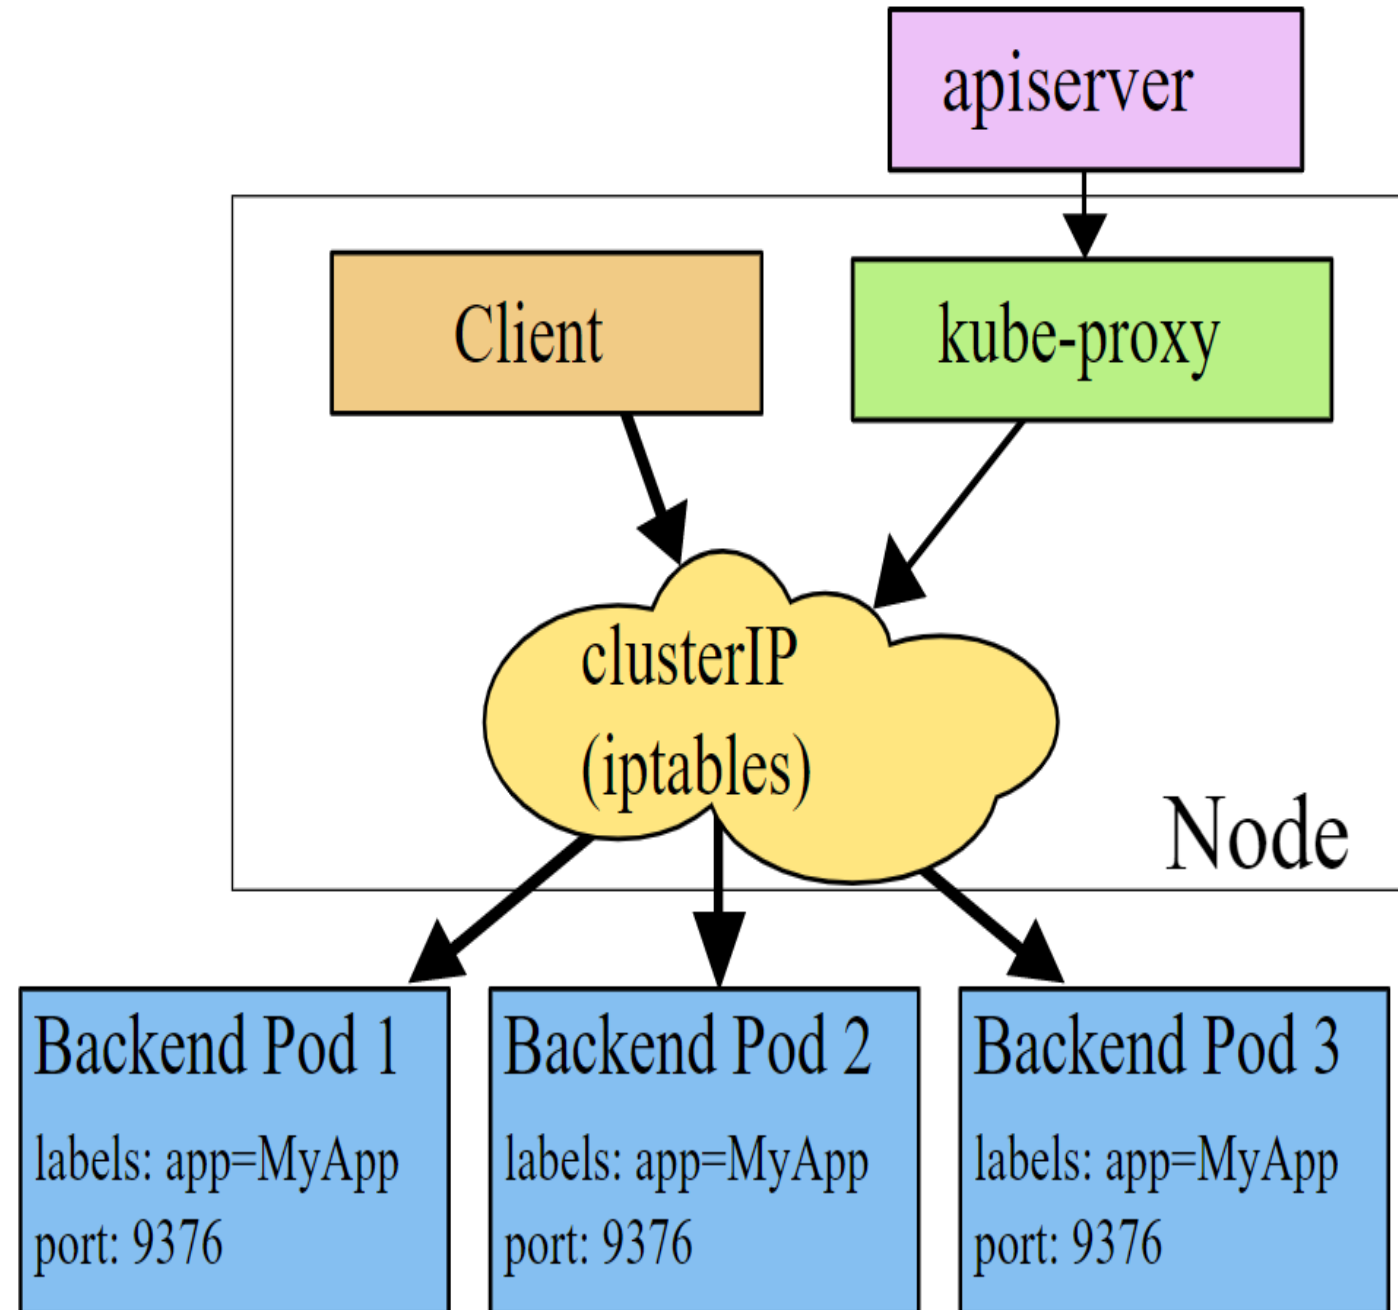
User Space Proxy Mode

- In this mode, kube-proxy watches the Kubernetes master for the addition and removal of Service and Endpoint objects. For each Service it opens a port (randomly chosen) on the local node. Any connections to this "proxy port" are proxied to one of the Service's backend Pods (as reported via Endpoints). kube-proxy takes the **SessionAffinity** setting of the Service into account when deciding which backend Pod to use.
- Lastly, the user-space proxy installs iptables rules which capture traffic to the Service's **clusterIP** (which is virtual) and **port**. The rules redirect that traffic to the proxy port which proxies the backend Pod.
- By default, kube-proxy in userspace mode chooses a backend via a round-robin algorithm.




Iptables Proxy Mode

- In this mode, kube-proxy watches the Kubernetes control plane for the addition and removal of Service and Endpoint objects. For each Service, it installs iptables rules, which capture traffic to the Service's **clusterIP** and **port**, and redirect that traffic to one of the Service's backend sets. For each Endpoint object, it installs iptables rules which select a backend Pod.
- Using iptables to handle traffic has a lower system overhead, because traffic is handled by Linux netfilter without the need to switch between userspace and the kernel space. This approach is also likely to be more reliable.
- The way kube-proxy programs the iptables rules means that it is nominally an **O(n)** style algorithm, where n grows roughly in proportion to your cluster size



IPVS Proxy Mode

- In **IPVS** mode, kube-proxy watches Kubernetes Services and Endpoints, calls **netlink** interface to create IPVS rules accordingly and synchronizes IPVS rules with Kubernetes Services and Endpoints periodically. This control loop ensures that IPVS status matches the desired state. When accessing a Service, IPVS directs traffic to one of the backend Pods.
- The IPVS proxy mode is based on netfilter hook function that is similar to iptables mode, but uses a hash table as the underlying data structure and works in the kernel space. That means kube-proxy in IPVS mode redirects traffic with lower latency than kube-proxy in iptables mode, with much better performance when synchronising proxy rules. Compared to the other proxy modes, IPVS mode also supports a higher throughput of network traffic.
- IPVS provides more options for balancing traffic to backend Pods; these are round-robin, least connection, destination hashing, source hashing etc.
- The result is that kube-proxy's connection processing in IPVS mode has a nominal computational complexity of **$O(1)$** . In other words, in most scenarios, its connection processing performance will stay constant independent of your cluster size.



Iptables vs IPVS


- IPVS supplied better extendibility and performance for large cluster.
- IPVS supports more complex balancing algorithms than iptables.
- IPVS supports service healthy check and connection retry functionalities and so on.

Iptables vs IPVS

- Round-Trip Response Times

- The right chart shows two key things:
 - The difference in average round-trip response times between iptables and IPVS is trivially insignificant until you get beyond 1,000 services (10,000 backend pods).
 - The difference in average round-trip response times is only discernible when not using keepalive connections. i.e. when using a new connection for every request.

Round-Trip Response Time vs Number of Services

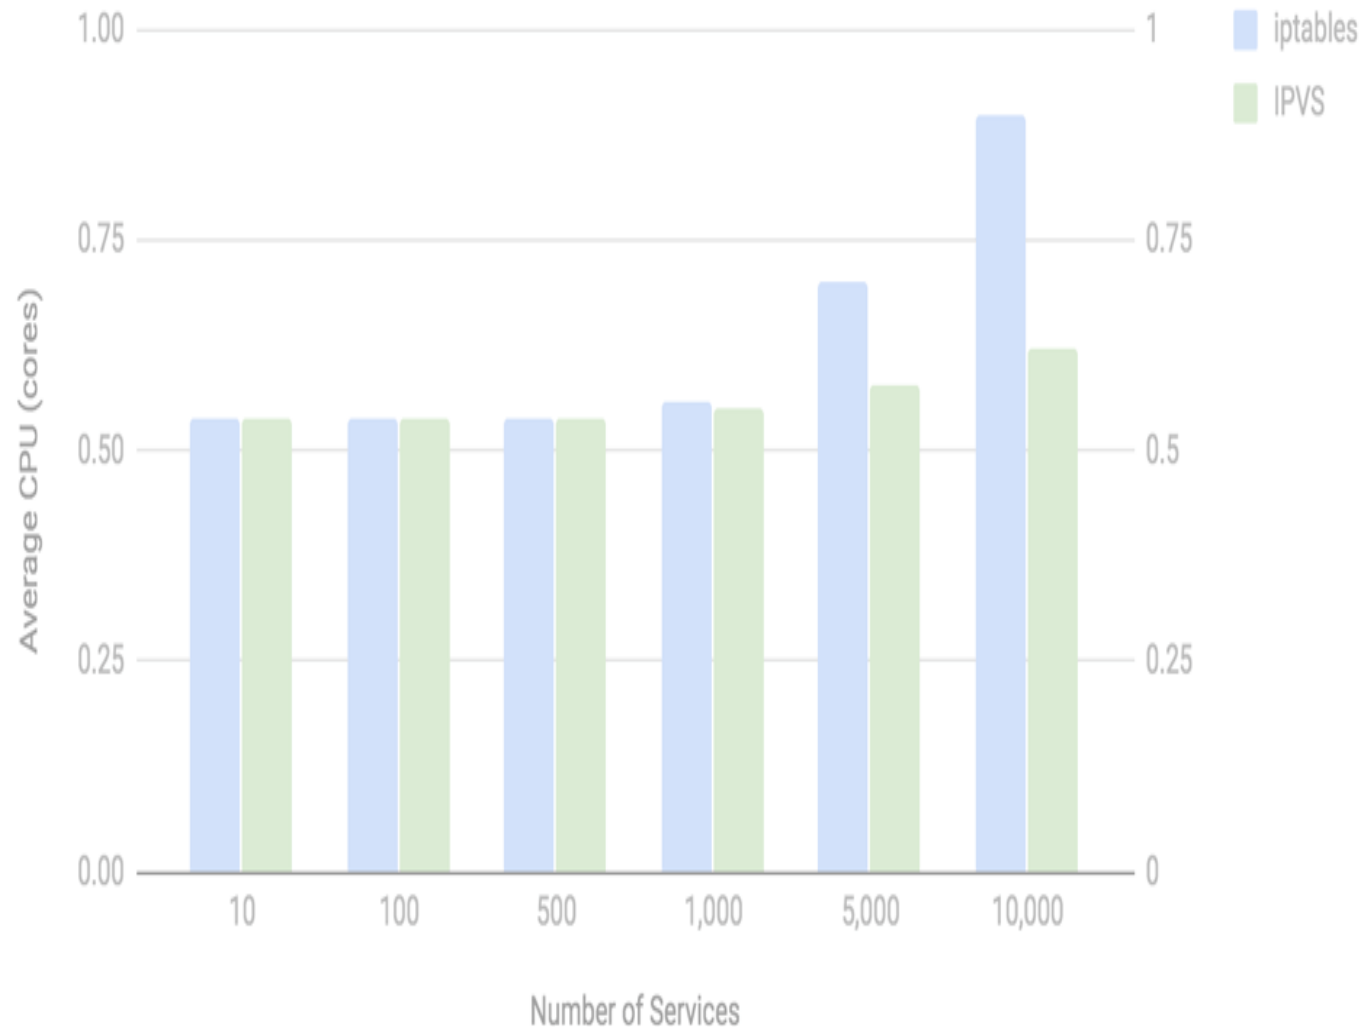


Iptables vs IPVS

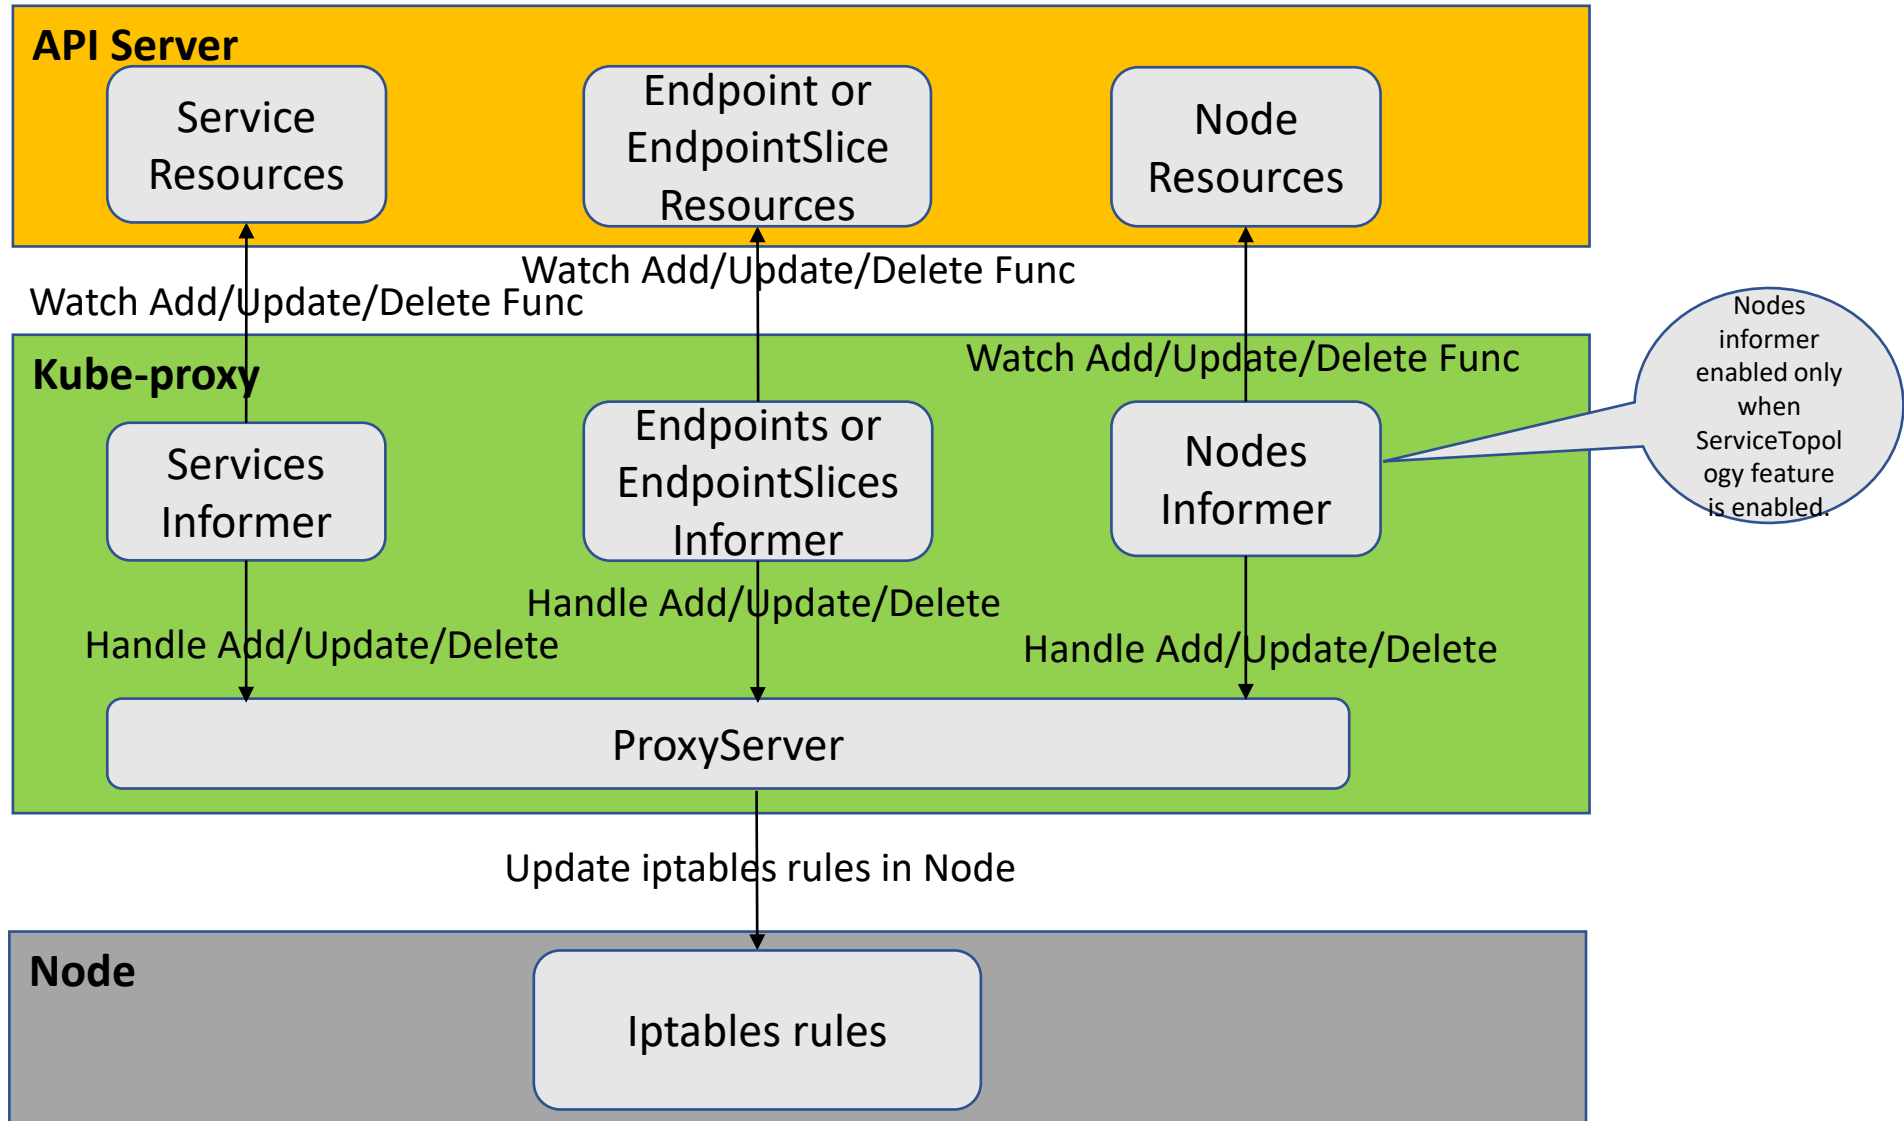
- Total CPU Usage

- The right chart shows two key things:
 - The difference in CPU usage between iptables and IPVS is relatively insignificant until you get beyond 1,000 services (with 10,000 backend pods).
 - At 10,000 services (with 100,000 backend pods), the increase in CPU with iptables is ~35% of a core, and with IPVS is ~8% of a core.

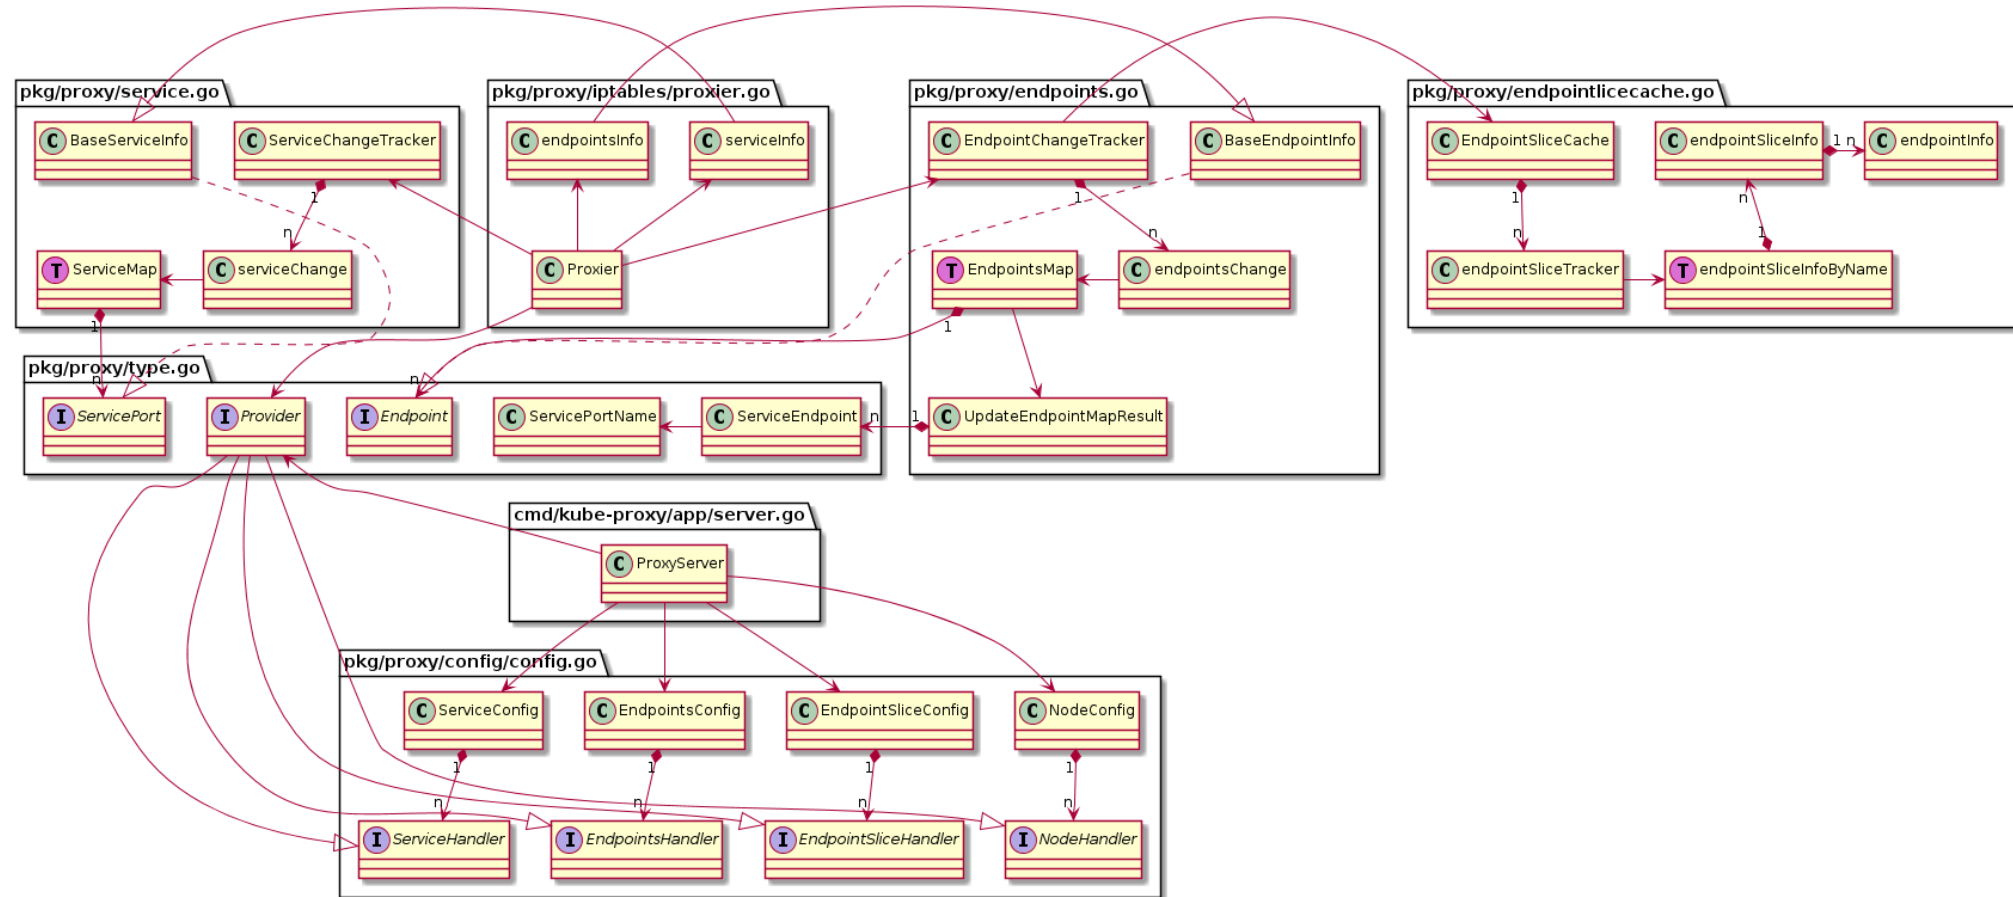
CPU vs Number of Services



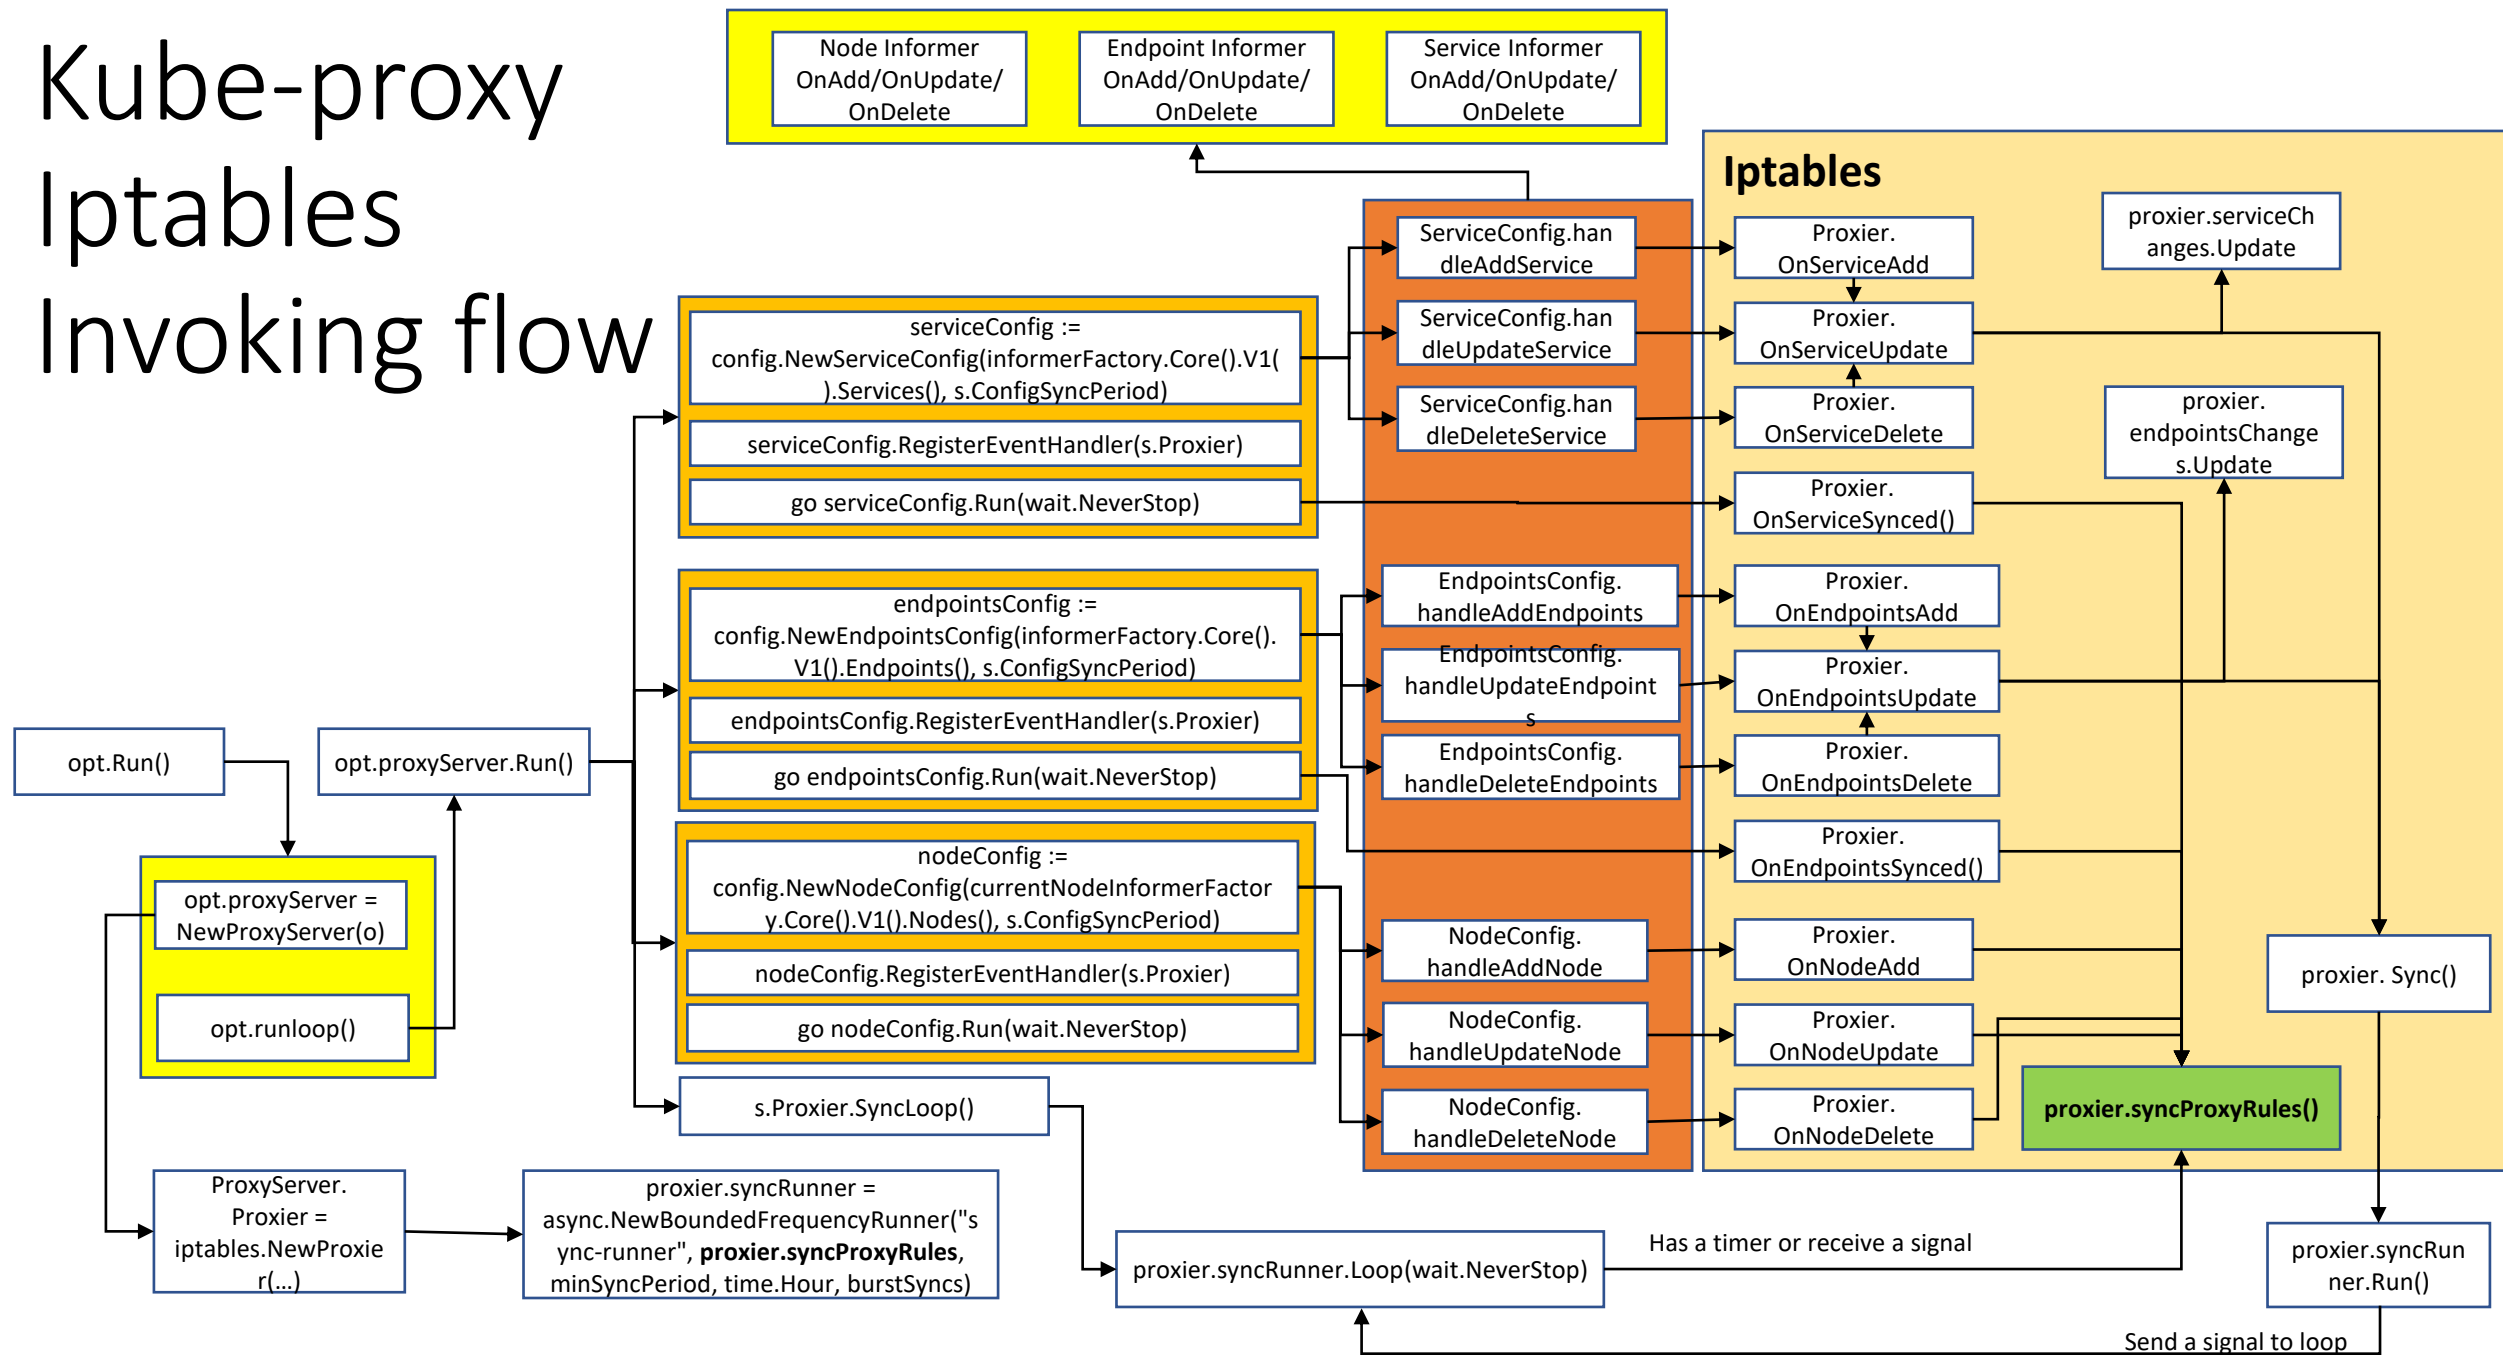
Kube-proxy informer flow



Kube-proxy class diagram



Kube-proxy Iptables Invoking flow



References

- <https://kubernetes.io/docs/reference/command-line-tools-reference/kube-proxy/>
- <https://kubernetes.io/docs/concepts/services-networking/service/>
- <https://cloud.tencent.com/developer/article/1097449>
- <https://www.tigera.io/blog/comparing-kube-proxy-modes-iptables-or-ipvs/>
- <https://netfilter.org/projects/iptables/index.html>
- https://en.wikipedia.org/wiki/IP_Virtual_Server
- <https://www.gikqiak.com/post/how-to-use-ipvs-in-kubernetes/>
- <https://opengrs.github.io/openstack/openstack-base-netfilter-framework-overview/>
- <https://www.frozentux.net/iptables-tutorial/chunkyhtml/x1309.html>
- <http://arthurchiao.art/blog/contrack-design-and-implementation-zh/>
- <http://arthurchiao.art/blog/deep-dive-into-iptables-and-netfilter-arch-zh/>
- <https://contrack-tools.netfilter.org/manual.html>
- <http://arthurchiao.art/blog/contrack-design-and-implementation-zh/>

kubelet

- The kubelet is the primary "node agent" that runs on each node. It makes sure that [containers](#) are running in a [Pod](#). It can register the node with the apiserver using one of: the hostname; a flag to override the hostname; or specific logic for a cloud provider.
- The kubelet works in terms of a PodSpec. A PodSpec is a YAML or JSON object that describes a pod. The kubelet takes a set of PodSpecs that are provided through various mechanisms (primarily through the apiserver) and ensures that the containers described in those PodSpecs are running and healthy. The kubelet doesn't manage containers which were not created by Kubernetes.
- Other than from a PodSpec from the apiserver, there are three ways that a container manifest can be provided to the Kubelet.
 - **File:** Path passed as a flag on the command line. Files under this path will be monitored periodically for updates. The monitoring period is 20s by default and is configurable via a flag.
 - **HTTP endpoint:** HTTP endpoint passed as a parameter on the command line. This endpoint is checked every 20 seconds (also configurable with a flag).
 - **HTTP server:** The kubelet can also listen for HTTP and respond to a simple API (underspec'd currently) to submit a new manifest.

Kubelet - containerManager

- Source code in `/pkg/kubelet/cm/`

Kubelet - containerGC

- containerGC provides two methods GarbageCollect() and DeleteAllUnusedContainers() to collect containers and deletes all unused containers respectively. And the real implementation is the implementor of Runtime interfaces. The implementor is KubeRuntimeManager in kubelet source code.
- GC has three types Policy:
 - **MinAge** - Minimum age at which a container can be garbage collected
 - **MaxPerPodContainer** - Max number of dead containers any single pod (UID, container name) pair is allowed to have
 - **MaxContainers** - Max number of total dead containers
- Source code in /pkg/kubelet/container/container_gc.go

Kubelet - imageManager

- imageManager is for pulling images, it has two methods to pull images(**serial** and **parallel**).
- imageManager mainly encapsulated **Image Service** interface of CRI, and also it provides a way of throttle image pulling, and it used a **TokenBucketRateLimiter** to control flow if provided qps and burst parameters when invoking **NewImageManager** function.
- Source code in /pkg/kubelet/images/image_manager.go
- /pkg/kubelet/images/image_gc_manager.go

Kubelet – imageGCManager

- ImageGCManager is mainly for removing unused images if over the max threshold for freeing enough place, also it uses some interfaces of CRI.
- GC Policies:
 - **HighThresholdPercent** - Any usage above this threshold will always trigger garbage collection(This is the highest usage we will allow).
 - **LowThresholdPercent** - Any usage below this threshold will never trigger garbage collection(This is the lowest threshold we will try to garbage collect to).
 - **MinAge** - Minimum age at which an image can be garbage collected.
- Source code in `/pkg/kubelet/images/image_gc_manager.go`

Kubelet - statusManager

- The statusManager is mainly for updating pod statuses in APIServer, and it writes only when the new status has changed.
- Source code in `/pkg/kubelet/status/`

Kubelet - probeManager

- Manager manages pod probing. It creates a probe "worker" for every container that specifies a probe (AddPod). The worker periodically probes its assigned container and caches the results. The manager use the cached probe results to set the appropriate Ready state in the PodStatus when requested (UpdatePodStatus). Updating probe parameters is not currently supported.
- In implementation there are a concept of **worker**, it mainly for handling he periodic probing of its assigned container. Each worker has a go-routine associated with it which runs the probe loop until the container permanently terminates, or the stop channel is closed. The worker uses the probe Manager's statusManager to get up-to-date container IDs
- It has readinessManager, livenessManager and startupManager. Respectively for managing the results of readiness probes, the results of liveness probes and the results of startup probes.
- The prober provided following three types prober. it will select one that you provided in v1.Probe to send probe.
 - Exec
 - HTTP
 - tcp
- Source code in /pkg/kubelet/prober/

Kubelet - evictionManager

- Memory pressure
- Disk pressure
- PID pressure
- Local storage eviction (if enabled LocalStorageCapacityIsolation feature)
 - containerEphemeralStorageLimitEviction
 - podEphemeralStorageLimitEviction
 - emptyDirLimitEviction
- Source code in `/pkg/kubelet/eviction/`

Kubelet - volumeManager

- <https://draveness.me/kubernetes-volume/>

Kubelet - runtimeManager

- runtimeManager provides a KubeGenericRuntimeManager. It mainly implemented all of container runtime interfaces(detailed interface, you can look at `/pkg/kubelet/container/runtime.go`).
- The runtimeManager is a real implementor of containerGC mechanism.
- Source code in `/pkg/kubelet/kuberuntime/`

Kubelet - podManager

- Manager stores and manages access to pods, maintaining the mappings between static pods and mirror pods. The kubelet discovers pod updates from 3 sources: file, http, and apiserver. Pods from non-apiserver sources are called static pods, and API server is not aware of the existence of static pods. In order to monitor the status of such pods, the kubelet creates a mirror pod for each static pod via the API server.
- A mirror pod has the same pod full name (name and namespace) as its static counterpart (albeit different metadata such as UID, etc). By leveraging the fact that the kubelet reports the pod status using the pod full name, the status of the mirror pod always reflects the actual status of the static pod. When a static pod gets deleted, the associated orphaned mirror pod will also be removed.

Kubelet - Pod Lifecycle Event Generator(PLEG)

- Getting pod current status from runtime, generating a pod lifecycle event and syncing it to kubelet pods cache.

Kubelet - cAdvisor

- Implemented all of interfaces of cAdvisor, so the real implementation, you should look at cAdvisor lib.

Kubelet - OOMWatcher

- Only start cAdvisor oomparser to do OOM, and this watcher only record events.
- So detailed OOM handling, please look at `StreamOoms(chan<-*oomparser.OomInstance)` method of github.com/google/cadvisor/utils/oomparser

References

- https://www.infoq.cn/article/odSLCISjvO8BNx*MbRbK
- <https://feisky.xyz/posts/kubernetes-container-runtime/>
- <https://github.com/opencontainers/runtime-spec/blob/master/config-linux.md>
- <https://kubernetes.io/docs/reference/command-line-tools-reference/kubelet/>
- <https://www.bookstack.cn/read/source-code-reading-notes/kubernetes-kubelet-modules.md>

Container-Runtime

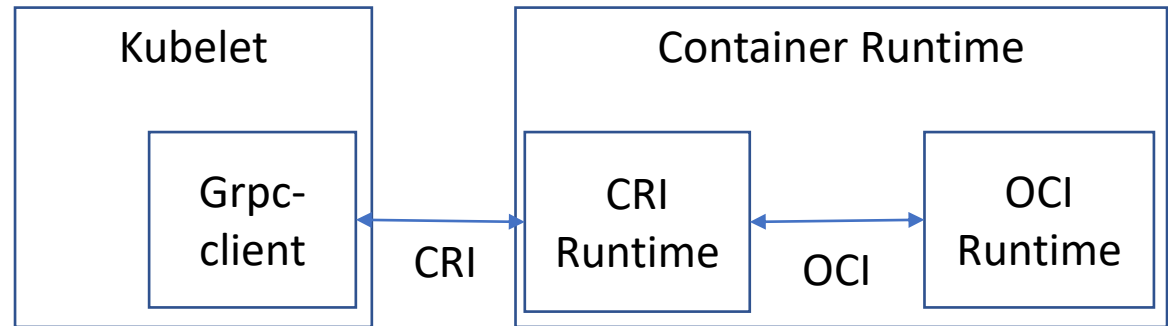
- Low-level container runtimes are responsible for setting up namespaces and cgroups for containers, and then running commands inside those namespaces and cgroups. Low-level runtimes support using these operating system features.
- High-level runtimes are responsible for transport and management of container images, unpacking the image, and passing off to the low-level runtime to run the container. Typically, high-level runtimes provide a daemon application and an API that remote applications can use to logically run containers and monitor but they sit on top of and delegate to low-level runtimes or other high-level runtimes for the actual work.

Container-Runtime

- Kubernetes runtimes are high-level container runtimes that support the [Container Runtime Interface](#) (CRI). and acts as a bridge between the [kubelet](#) and the container runtime. High-level container runtimes that want to integrate with Kubernetes are expected to implement CRI. The runtime is expected to handle the management of images and to support [Kubernetes pods](#), as well as manage the individual containers.

Container-Runtime

- The kubelet uses CRI to communicate with the container runtime running on that same node. In this way CRI is simply an abstraction layer or API that allows you to switch out container runtime implementations instead of having them built into the kubelet.



Container Runtime Interface(CRI)

- CRI (*Container Runtime Interface*) consists of a [protobuf API](https://github.com/kubernetes/kubernetes/tree/master/pkg/kubelt/server/streaming), specifications/requirements (to-be-added), and [libraries] (<https://github.com/kubernetes/kubernetes/tree/master/pkg/kubelt/server/streaming>) for container runtimes to integrate with kubelet on a node.
- <https://github.com/kubernetes/cri-api/blob/master/pkg/apis/runtime/v1alpha2/api.proto>

CRI Implementations

- Dockershim
 - Embedded into kubelet. And Dockershim communicates with docker.
- CRI-O
- Containerd
- Frakti
- rktlet

Open Container Initiative(OCI)

- The Open Container Initiative (OCI) is a lightweight, open governance structure (project), formed under the auspices of the Linux Foundation, for the express purpose of creating open industry standards around container formats and runtime.
- The OCI currently contains two specifications: **the Runtime Specification (runtime-spec)** and **the Image Specification (image-spec)**. The Runtime Specification outlines how to run a “filesystem bundle” that is unpacked on disk. At a high-level an OCI implementation would download an OCI Image then unpack that image into an OCI Runtime filesystem bundle. At this point the OCI Runtime Bundle would be run by an OCI Runtime.

OCI Runtime Specification

- The Open Container Initiative Runtime Specification aims to specify the configuration, execution environment, and lifecycle of a container.
- A container's configuration is specified as the config.json for the supported platforms and details the fields that enable the creation of a container. The execution environment is specified to ensure that applications running inside a container have a consistent environment between runtimes along with common actions defined for the container's lifecycle.

OCI Image Format Specification

- This specification defines an OCI Image, consisting of a [manifest](#), an [image index](#) (optional), a set of [filesystem layers](#), and a [configuration](#).
- The goal of this specification is to enable the creation of interoperable tools for building, transporting, and preparing a container image to run.

OCI Runtime Implementations

- runc
- Kata containers

Container

- Containers are implemented using [Linux namespaces](#) and [cgroups](#).
 - Namespaces let you virtualize system resources, like the file system or networking, for each container.
 - Cgroups provide a way to limit the amount of resources like CPU and memory that each container can use.

Kernel Namespaces

- **Namespaces** are a feature of the [Linux kernel](#) that partitions kernel resources such that one set of [processes](#) sees one set of resources while another set of processes sees a different set of resources. The feature works by having the same namespace for a set of resources and processes, but those namespaces refer to distinct resources. Resources may exist in multiple spaces. Examples of such resources are process IDs, hostnames, user IDs, file names, and some names associated with network access, and [interprocess communication](#).

Kernel Namespaces

- Process ID(PID) namespace
 - The [PID](#) namespace provides processes with an independent set of process IDs (PIDs) from other namespaces. PID namespaces are nested, meaning when a new process is created it will have a PID for each namespace from its current namespace up to the initial PID namespace. Hence the initial PID namespace is able to see all processes, albeit with different PIDs than other namespaces will see processes with.
 - The first process created in a PID namespace is assigned the process id number 1 and receives most of the same special treatment as the normal [init](#) process, most notably that [orphaned processes](#) within the namespace are attached to it. This also means that the termination of this PID 1 process will immediately terminate all processes in its PID namespace and any descendants.[\[5\]](#)
- Mount namespace
 - Mount namespaces control mount points. Upon creation the mounts from the current mount namespace are copied to the new namespace, but mount points created afterwards do not propagate between namespaces (using shared subtrees, it is possible to propagate mount points between namespaces[\[4\]](#)).
 - The clone flag used to create a new namespace of this type is CLONE_NEWNS - short for "NEW NameSpace". This term is not descriptive (as it doesn't tell which kind of namespace is to be created) because mount namespaces were the first kind of namespace and designers did not anticipate there being any others.
- Network namespace
 - Network namespaces virtualize the [network stack](#). On creation a network namespace contains only a [loopback](#) interface.
 - Each [network interface](#) (physical or virtual) is present in exactly 1 namespace and can be moved between namespaces.
 - Each namespace will have a private set of [IP addresses](#), its own [routing table](#), [socket](#) listing, connection tracking table, [firewall](#), and other network-related resources.
 - Destroying a network namespace destroys any virtual interfaces within it and moves any physical interfaces within it back to the initial network namespace.
- Inter Process Communication(IPC) namespace
- User ID namespace
- Control group(cgroup) namespace
- Time namespace
- UTS(UNIX Time Sharing) namespace

Control Groups (cgroups)

- **cgroups** (abbreviated from **control groups**) is a [Linux kernel](#) feature that limits, accounts for, and isolates the [resource usage](#) (CPU, memory, disk I/O, network, etc.) of a collection of [processes](#).
- One of the design goals of cgroups is to provide a unified interface to many different [use cases](#), from controlling single processes (by using [nice](#), for example) to full [operating system-level virtualization](#) (as provided by [OpenVZ](#), [Linux-VServer](#) or [LXC](#), for example).
- Cgroups provides:
 - **Resource limiting**
 - groups can be set to not exceed a configured [memory](#) limit, which also includes the [file system cache](#)
 - **Prioritization**
 - some groups may get a larger share of CPU utilization^[11] or disk I/O throughput^[12]
 - **Accounting**
 - measures a group's resource usage, which may be used, for example, for billing purposes^[13]
 - **Control**
 - freezing groups of processes, their [checkpointing](#) and restarting^[13]

Container Image

- A **Docker image** is an immutable (unchangeable) file that contains the source code, libraries, dependencies, tools, and other files needed for an application to run.

UnionFS

- **Unionfs** is a filesystem service for [Linux](#), [FreeBSD](#) and [NetBSD](#) which implements a [union mount](#) for other [file systems](#). It allows files and directories of separate file systems, known as branches, to be transparently overlaid, forming a single coherent file system. Contents of directories which have the same path within the merged branches will be seen together in a single merged directory, within the new, virtual filesystem.
- When mounting branches, the priority of one branch over the other is specified. So when both branches contain a file with the same name, one gets priority over the other.
- The different branches may be either *read-only* or *read/write* file systems, so that writes to the virtual, merged copy are directed to a specific real file system. This allows a file system to appear as writable, but without actually allowing writes to change the file system, also known as [copy-on-write](#). This may be desirable when the media is physically read-only, such as in the case of [Live CDs](#).

References

- <https://www.ianlewis.org/en/container-runtimes-part-1-introduction-container-r>
- <https://jvns.ca/blog/2016/10/10/what-even-is-a-container/>
- https://en.wikipedia.org/wiki/Linux_namespaces
- <https://en.wikipedia.org/wiki/Cgroups>
- <https://developers.redhat.com/blog/2018/02/22/container-terminology-practical-introduction/>
- <https://github.com/opencontainers/runtime-spec/blob/master/spec.md>
- <https://github.com/opencontainers/image-spec/blob/master/spec.md>
- <https://kubernetes.io/blog/2016/12/container-runtime-interface-cri-in-kubernetes/>
- <https://github.com/kubernetes/kubernetes/blob/242a97307b34076d5d8f5bb154fa4d97c9ef1d/docs/devel/container-runtime-interface.md>
- <https://www.jianshu.com/p/3ba255463047>
- <https://en.wikipedia.org/wiki/UnionFS>
- <https://unionfs.filesystems.org/>
- <https://phoenixnap.com/kb/docker-image-vs-container>
- <https://xuanwo.io/2019/08/06/oci-intro/>

Client-go

- Go clients for talking to a [kubernetes](#) cluster.
- What's included
 - The `kubernetes` package contains the clientset to access Kubernetes API.
 - The `discovery` package is used to discover APIs supported by a Kubernetes API server.
 - The `dynamic` package contains a dynamic client that can perform generic operations on arbitrary Kubernetes API objects.
 - The `plugin/pkg/client/auth` packages contain optional authentication plugins for obtaining credentials from external sources.
 - The `transport` package is used to set up auth and start a connection.
 - The `tools/cache` package is useful for writing controllers.

REST Client

- This is a RESTful Client which supports json and protobuf and native resources and CRD resources.

Dynamic Client

- dynamicClient is a dynamic client, in fact, it encapsulated rest client. It supports dynamic configuration resource Group, version and which resource used.
- Example: <https://github.com/kubernetes/client-go/tree/master/examples/dynamic-create-update-delete-deployment>
- NOTE: This example is based on out of cluster env. If you want to use it in cluster env(such as you want to use it in pod container). You need use “rest.InClusterConfig()” to generate a configuration. More detailed example please reference <https://github.com/kubernetes/client-go/tree/master/examples/in-cluster-client-configuration>.

Typed ClientSet

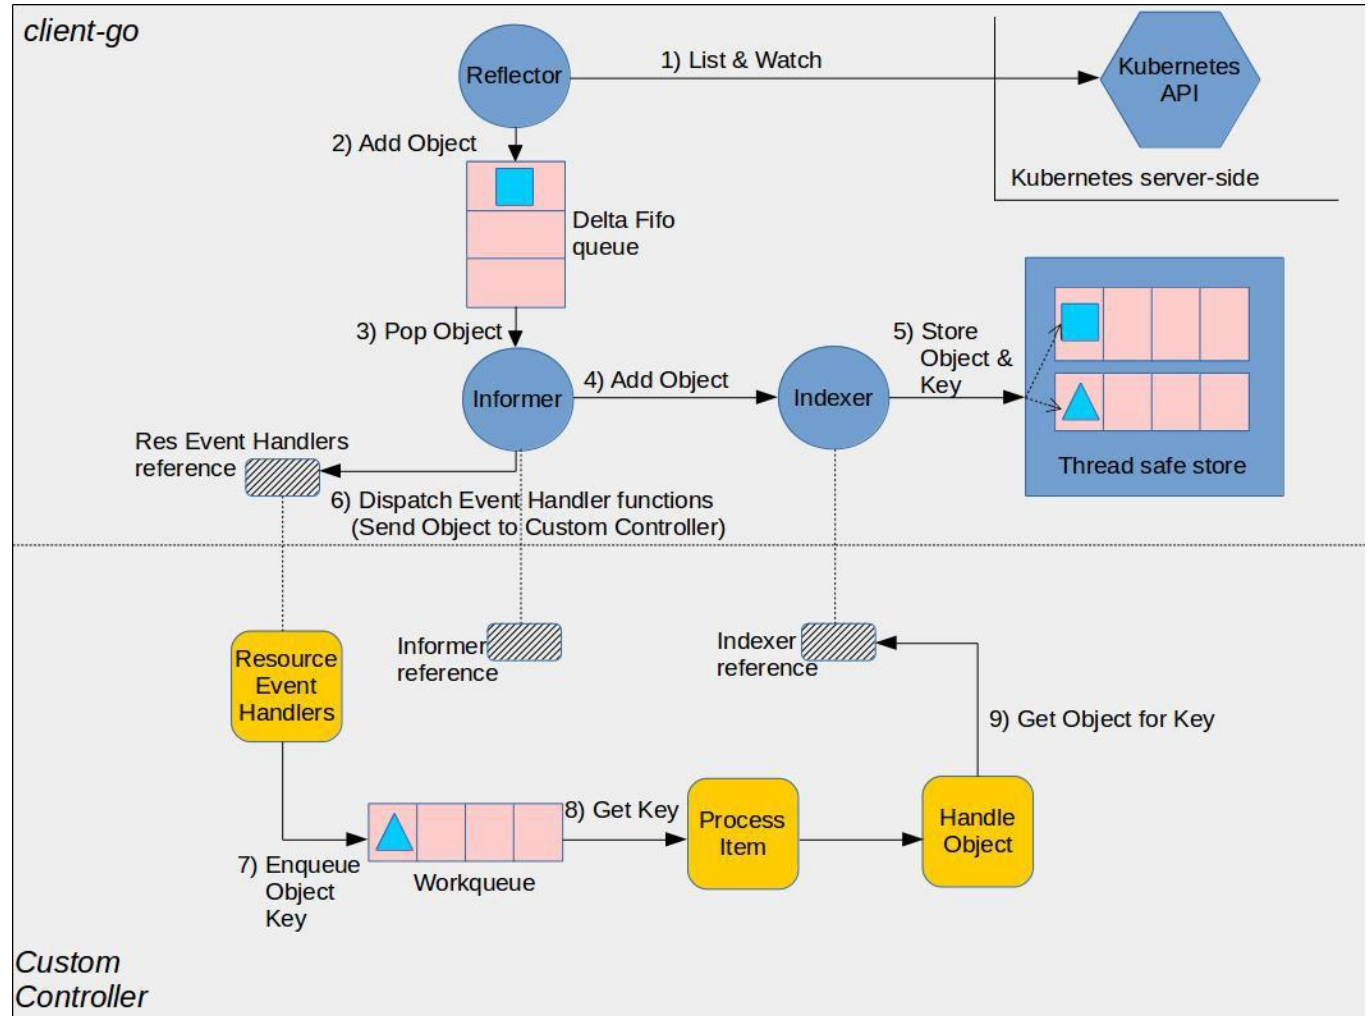
- Typed clientSet is a kubernetes client. It defined all resources API operations. So you can directly use it to operate resources of the all versions.
- Example: <https://github.com/kubernetes/client-go/tree/master/examples/create-update-delete-deployment>

Typed Vs Dynamic

- The typed client sets make it simple to communicate with the API server using pre-generated local API objects to achieve an RPC-like programming experience. Typed clients use program compilations to enforce data safety and some validation. However, when using typed clients, programs are forced to be tightly coupled with the version and the types used.
- The **dynamic** package on the other hand, uses a simple type, **unstructured.Unstructured**, to represent all object values from the API server. Type **Unstructured** uses a collection of nested **map[string]interface{}** values to create an internal structure that closely resembles the REST payload from the server.
- The dynamic package defers all data bindings until runtime. This means programs that use the dynamic client will not get any of the benefits of type validation until the program is running. This may be a problem for certain types of applications that require strong data type check and validation.
- Being loosely coupled, however, means that programs that use the **dynamic** package do not require recompilation when the client API changes. The client program has more flexibility in handling updates to the API surface without knowing ahead of time what those changes are.

Informer

This is a pictorial representation showing how the various components in the client-go library work and their interaction points with the custom controller code that you will write.



Informer Mechanism

- **client-go components**

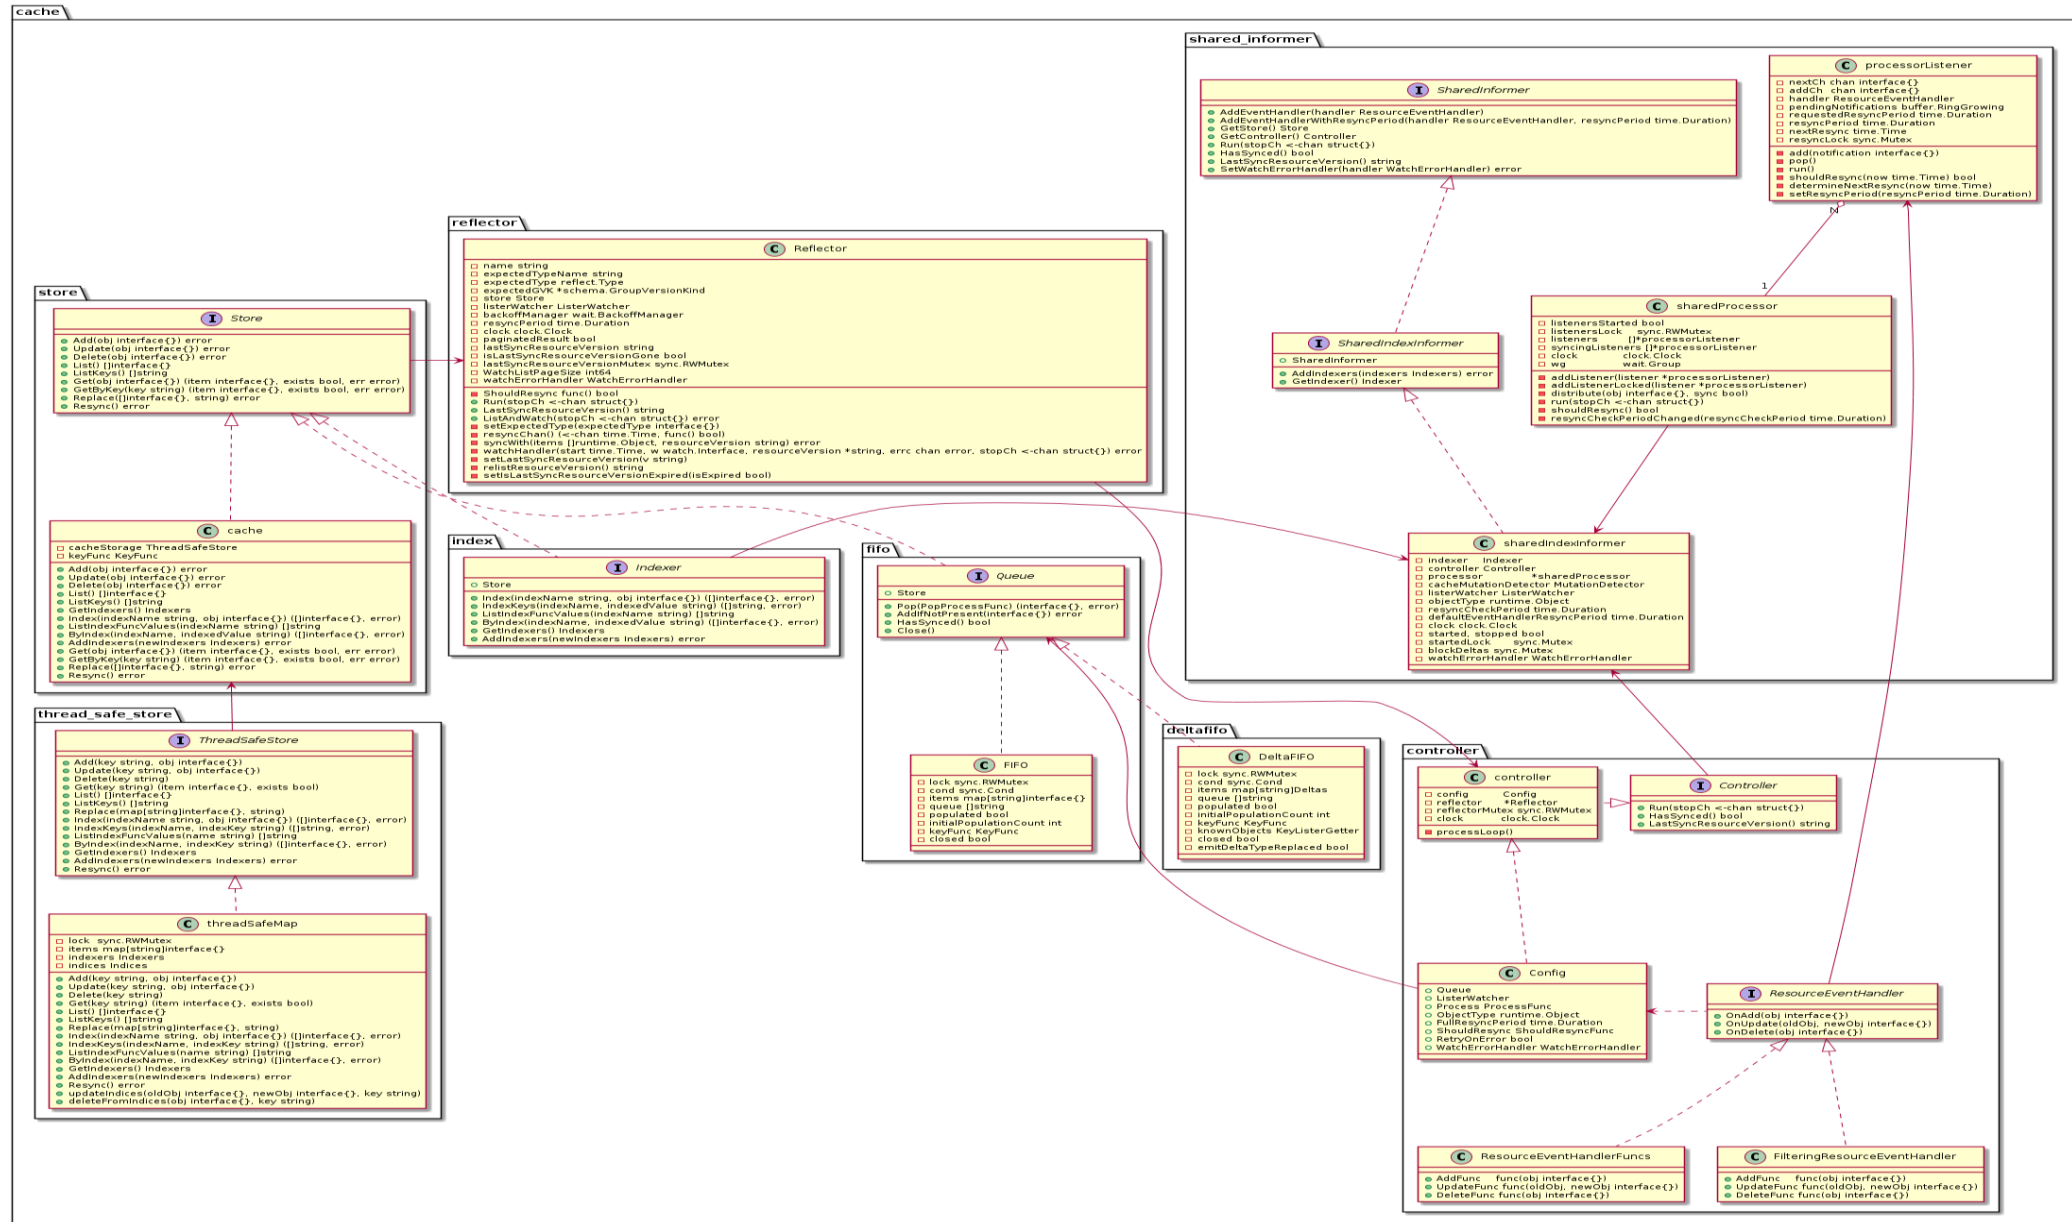
- **Reflector:** A reflector, which is defined in [type *Reflector* inside package *cache*](#), watches the Kubernetes API for the specified resource type (kind). The function in which this is done is *ListAndWatch*. The watch could be for an in-built resource or it could be for a custom resource. When the reflector receives notification about existence of new resource instance through the watch API, it gets the newly created object using the corresponding listing API and puts it in the Delta Fifo queue inside the *watchHandler* function.
- **Informer:** An informer defined in the [base controller inside package *cache*](#) pops objects from the Delta Fifo queue. The function in which this is done is *processLoop*. The job of this base controller is to save the object for later retrieval, and to invoke our controller passing it the object.
- **Indexer:** An indexer provides indexing functionality over objects. It is defined in [type *Indexer* inside package *cache*](#). A typical indexing use-case is to create an index based on object labels. Indexer can maintain indexes based on several indexing functions. Indexer uses a thread-safe data store to store objects and their keys. There is a default function named *MetaNamespaceKeyFunc* defined in [type *Store* inside package *cache*](#) that generates an object's key as <namespace>/<name> combination for that object.

Informer Mechanism

- **Custom Controller components**

- **Informer reference:** This is the reference to the Informer instance that knows how to work with your custom resource objects. Your custom controller code needs to create the appropriate Informer.
- **Indexer reference:** This is the reference to the Indexer instance that knows how to work with your custom resource objects. Your custom controller code needs to create this. You will be using this reference for retrieving objects for later processing.
- **Resource Event Handlers:** These are the callback functions which will be called by the Informer when it wants to deliver an object to your controller. The typical pattern to write these functions is to obtain the dispatched object's key and enqueue that key in a work queue for further processing.
- **Work queue:** This is the queue that you create in your controller code to decouple delivery of an object from its processing. Resource event handler functions are written to extract the delivered object's key and add that to the work queue.
- **Process Item:** This is the function that you create in your code which processes items from the work queue. There can be one or more other functions that do the actual processing. These functions will typically use the [Indexer reference](#), or a Listing wrapper to retrieve the object corresponding to the key.

Informer class diagram

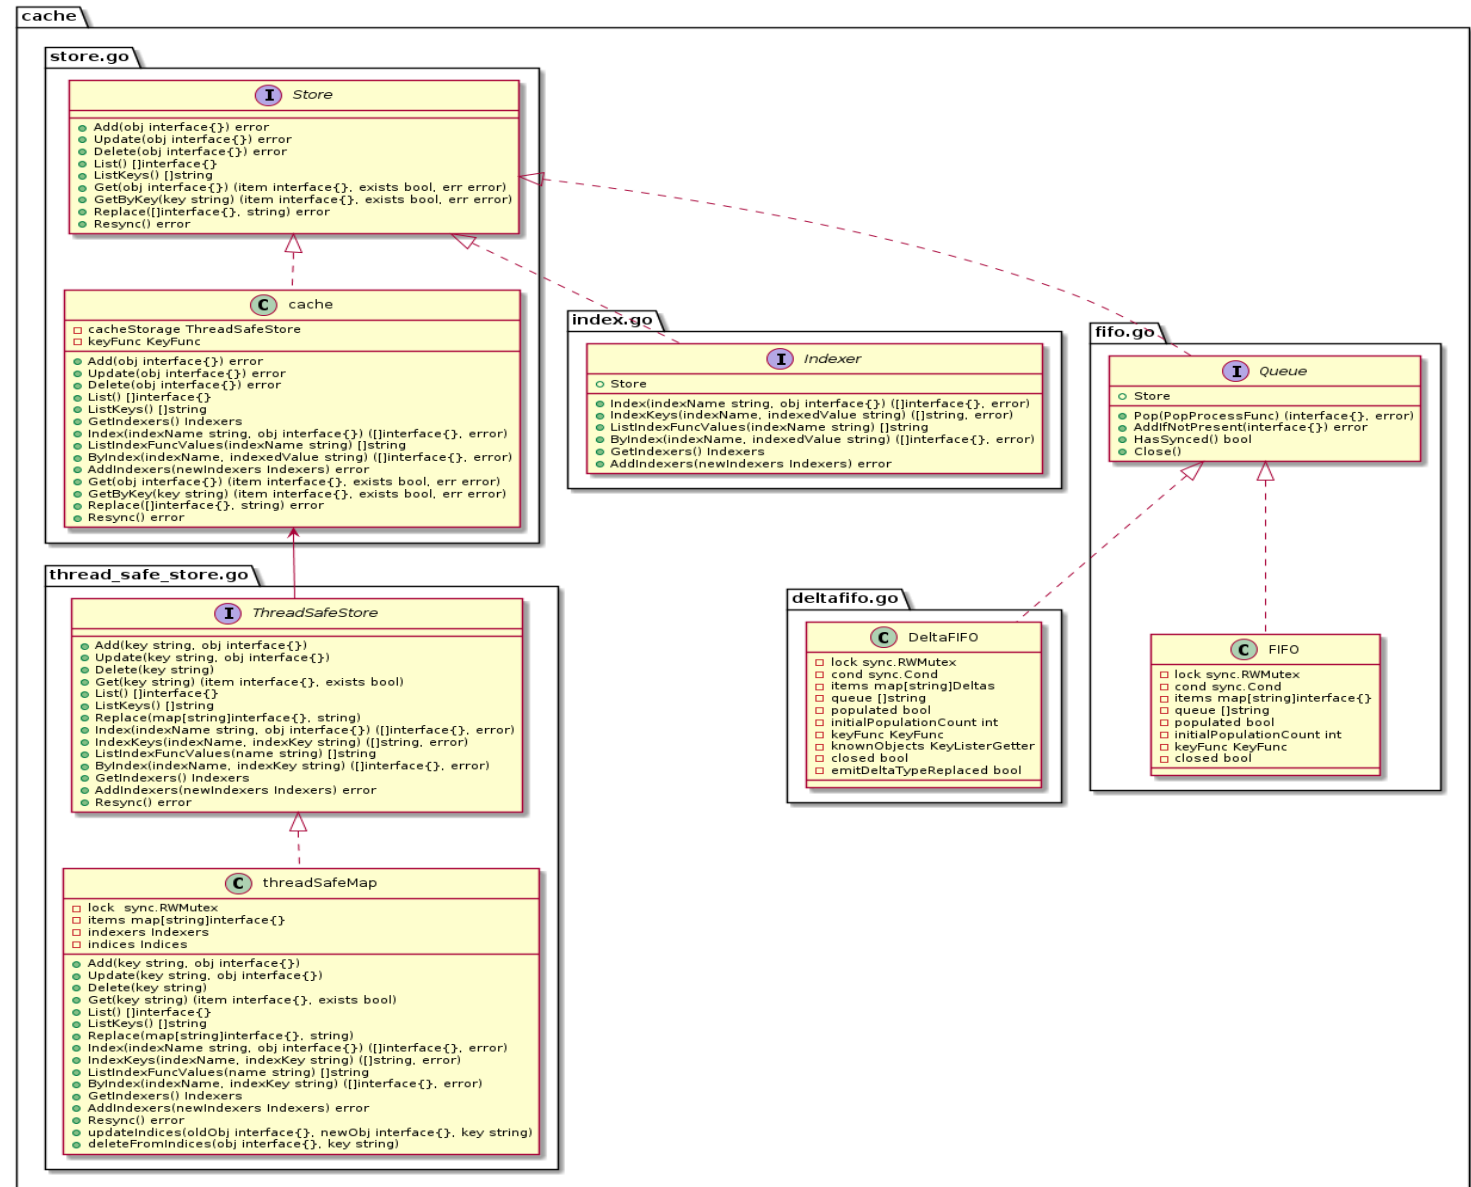


DeltaFIFO ReSync

- ReSync Indexer data into DeltaFIFO.
 - This mechanism can ensure events will not be lost, but it will trigger some redundancy events.
- Why:
 - Q1: controller has already implemented work queue mechanism, why need this resync operation?
 - A1: not all of controllers have implemented this work queue. And work queue has a limitation which is that it has a size restriction(default value is 100). So if controller didn't implement work queue or work queue overloaded due to too many failures. We need this mechanism to avoid event lost.
 - Q2: how to ensure resync will not overwrite event that it had already existed in DeltaFIFO?
 - A2: If the event has already existed in DeltaFIFO, then ignore this resync operation. If not, add this event into DeltaFIFO. Read source code https://github.com/kubernetes/client-go/blob/release-1.18/tools/cache/delta_fifo.go#L367
 - Q3: why resync operation will trigger an onUpdate event?
 - A3: read source code https://github.com/kubernetes/client-go/blob/release-1.18/tools/cache/shared_informer.go#L494 and https://github.com/kubernetes/client-go/blob/release-1.18/tools/cache/delta_fifo.go#L595
 - Q4: if a listener has not event in a long time, then this resync mechanism will always trigger some redundancy events?
 - A4: Yes.

Indexer

- **Store** is a generic object storage and processing interface
- **Indexer** extends Store with multiple indices and restricts each accumulator to simply hold the current object (and be empty after Delete).
- **cache** implements Indexer in terms of a ThreadSafeStore and an associated KeyFunc.
- **ThreadSafeStore** is an interface that allows concurrent indexed access to a storage backend. It is like Indexer but does not (necessarily) know how to extract the Store key from a given object.
- **threadSafeMap** implements ThreadSafeStore



threadSafeMap

- **threadSafeMap** includes following attributes:
 1. **items** which is saved all of data
 2. **indexers** which is mapping a name to an IndexFunc
 3. **indices** which is mapping a name to an Index.

```
// threadSafeMap implements ThreadSafeStore
type threadSafeMap struct {
    lock sync.RWMutex
    items map[string]interface{}

    // indexers maps a name to an IndexFunc
    indexers Indexers

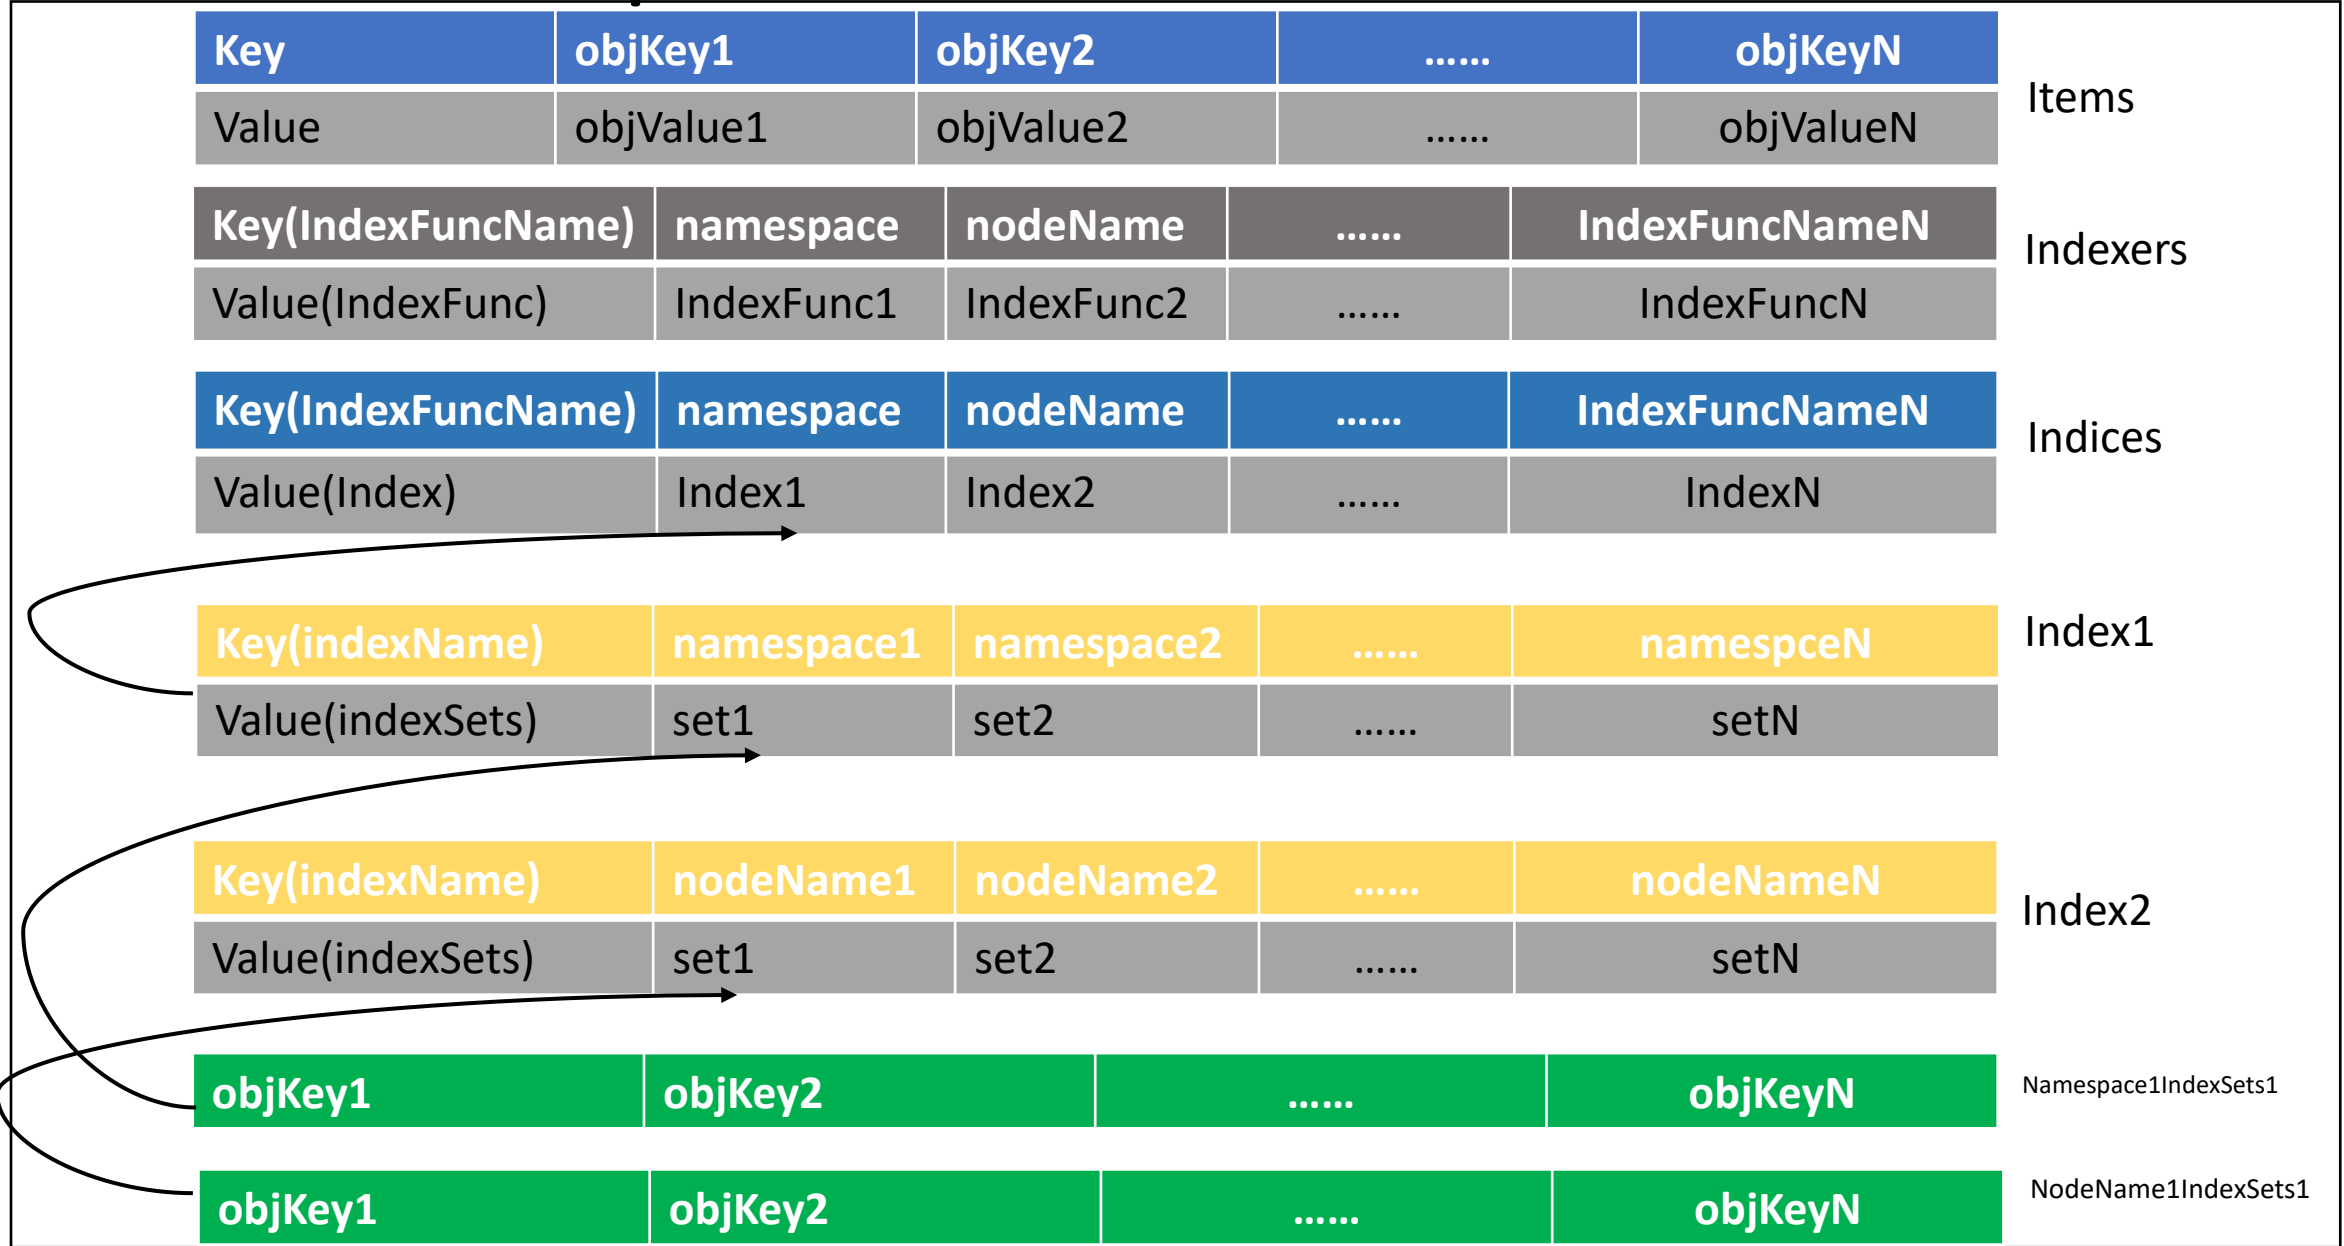
    // indices maps a name to an Index
    indices Indices
}
```

```
// Index maps the indexed value to a set of keys in the store that match on that value
type Index map[string]sets.String

// Indexers maps a name to a IndexFunc
type Indexers map[string]IndexFunc

// Indices maps a name to an Index
type Indices map[string]Index
```

threadSafeMap



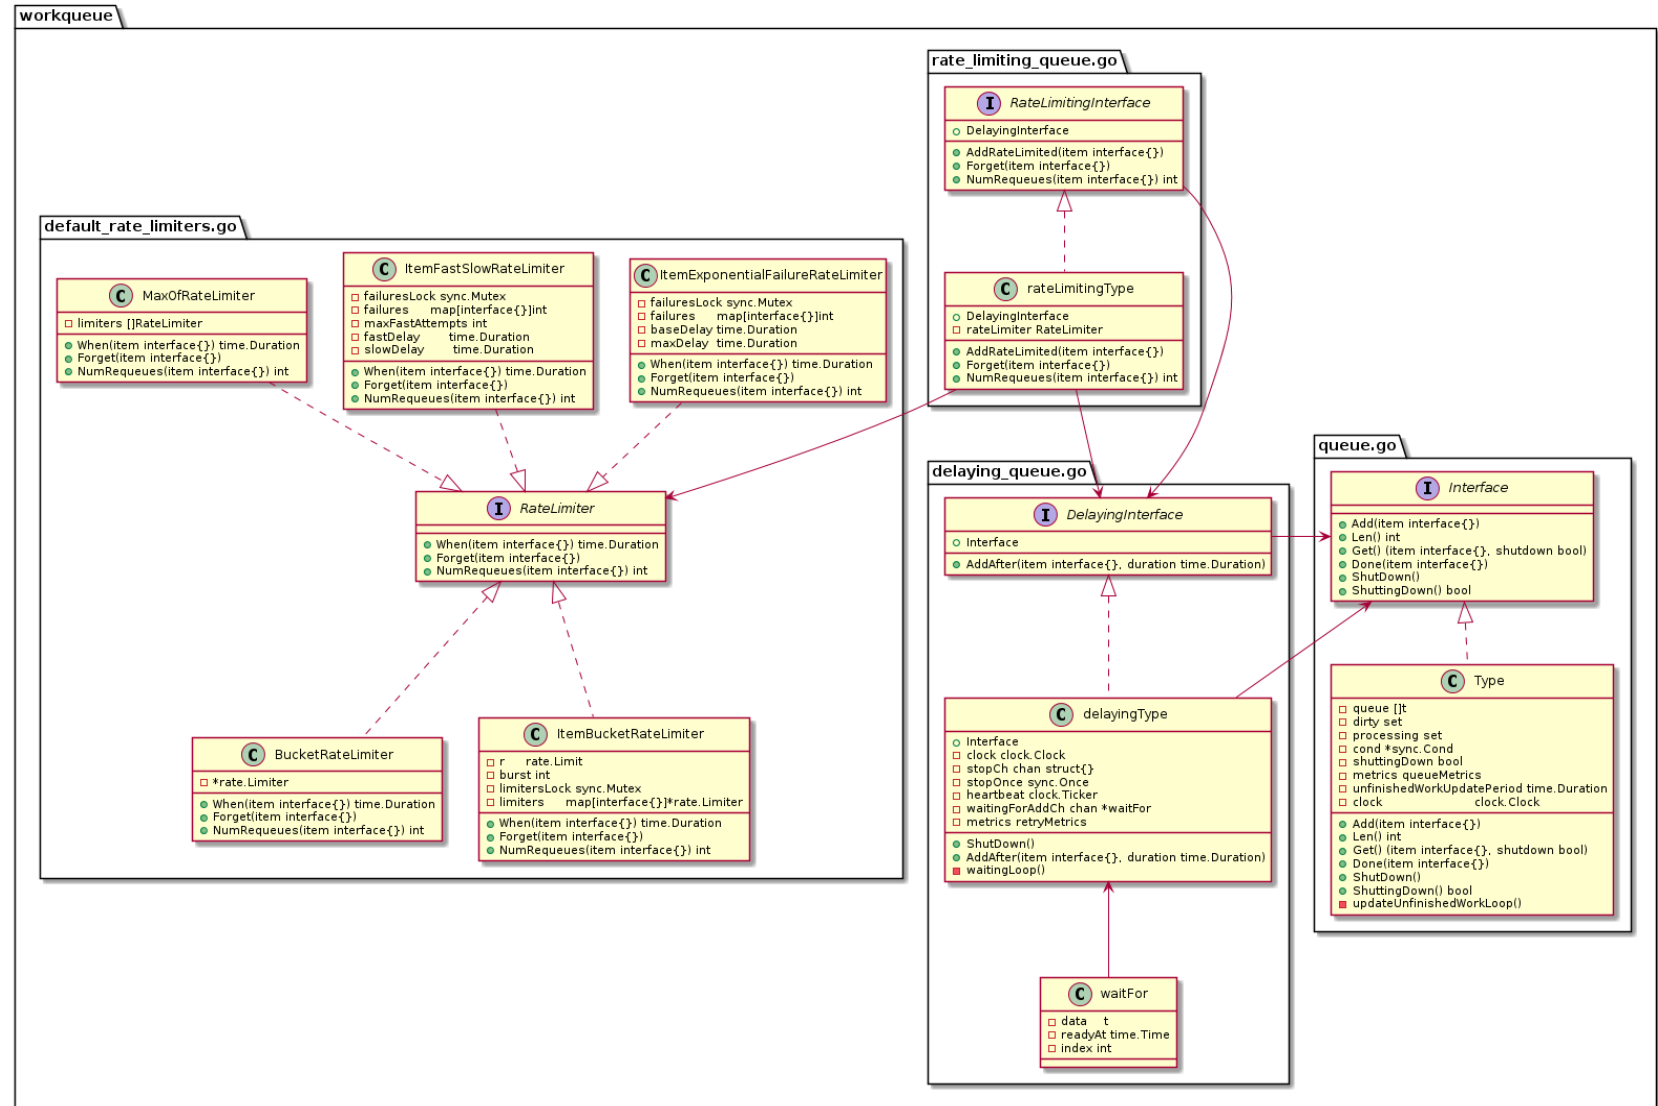
threadSafeMap

Example: Add a newObj into map

1. Add a newObj into items that specified key
2. Update this new key into indices
 1. Delete oldObj from indices if it exists.
 2. Iterate indexers to find indexValues of the newObj.
 3. Use indexer name to find index from indices.
 4. Create a new index if index is not existed. Otherwise next step.
 5. Iterate indexValues to find set from index.
 6. Create a new set if set is not existed.
 7. Add the new specified key into set.

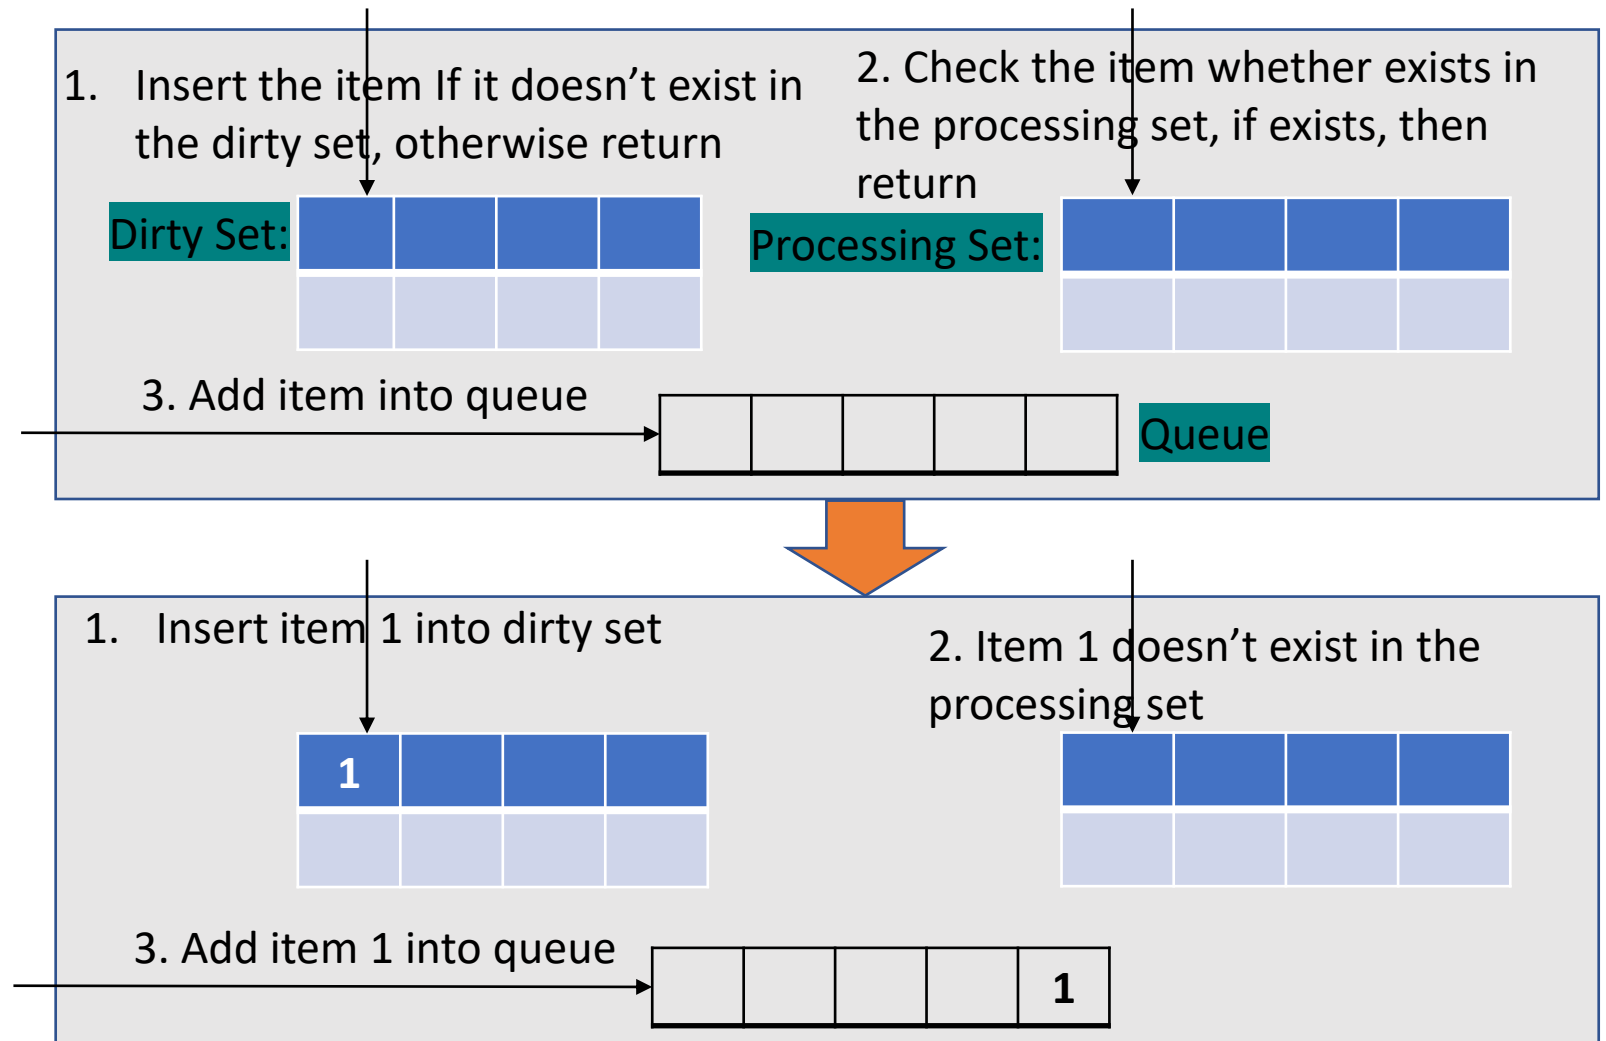
WorkQueue

- **Interface** . A FIFO queue interface , supports deduplication.
- **DelayingInterface**. A delay queue interface, encapsulated based on **Interface** interface.
- **RateLimitingInterface**. A rate limiting interface, encapsulated based on **DelayingInterface** .



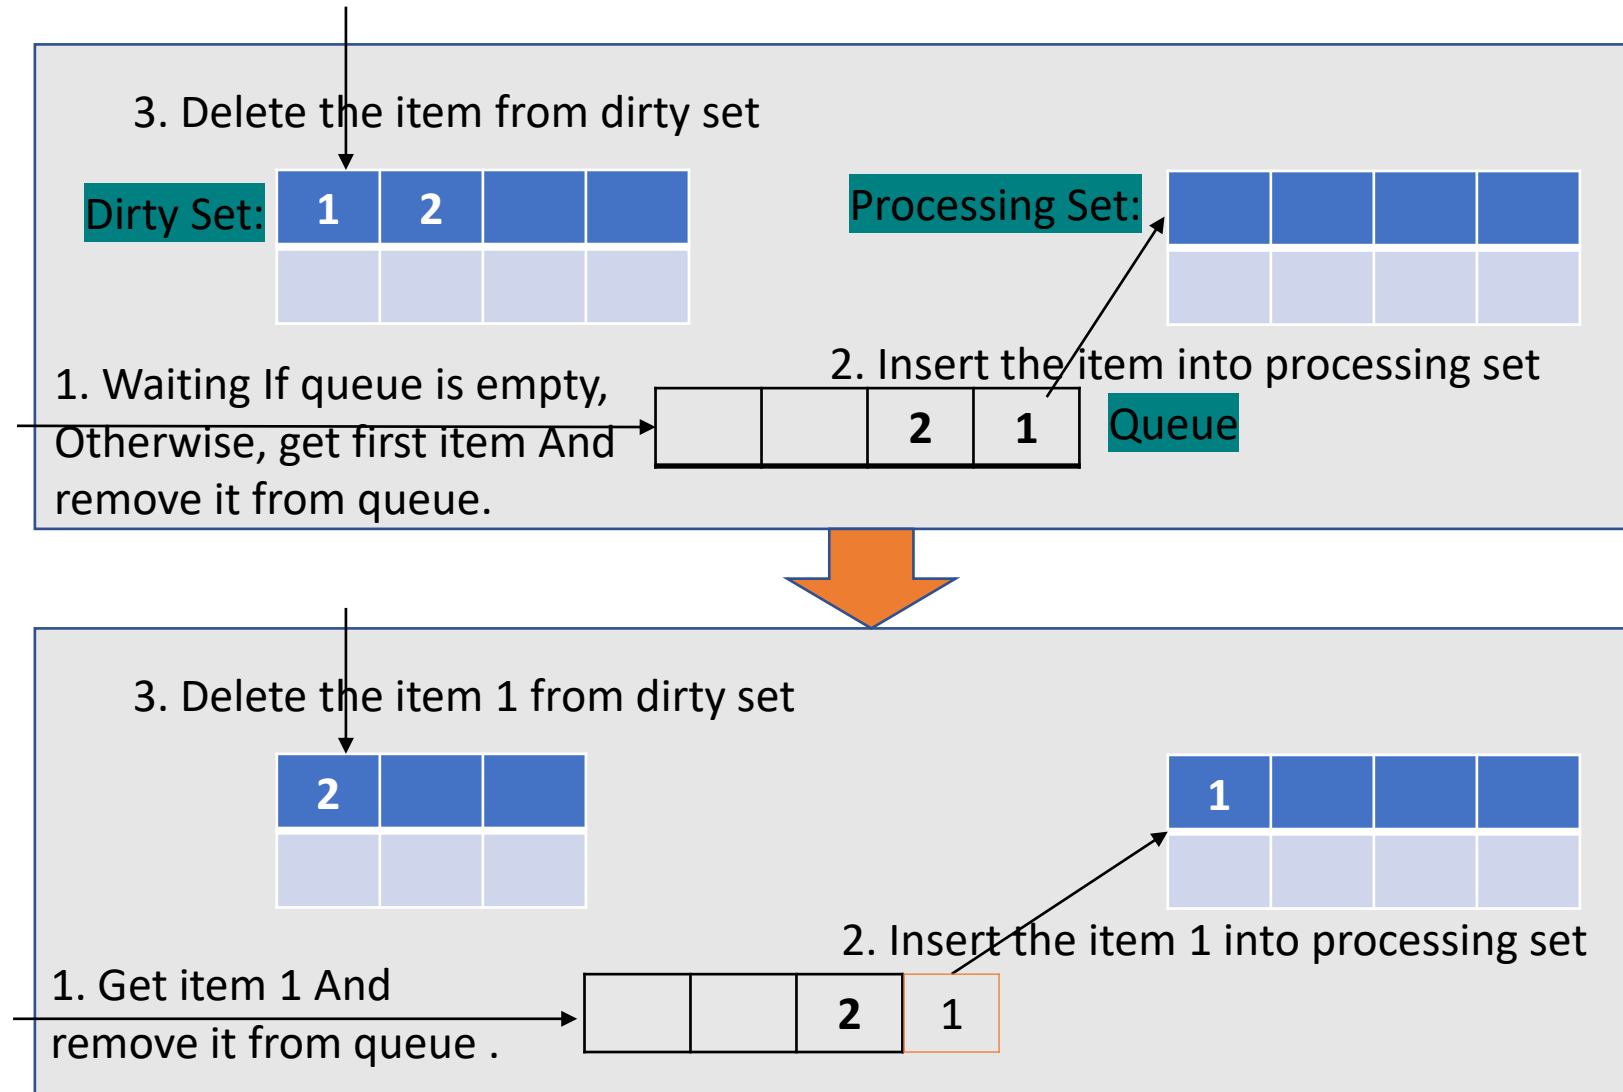
Basic Queue

- **Add** an item handling.
- **Add** marks item as needing processing.



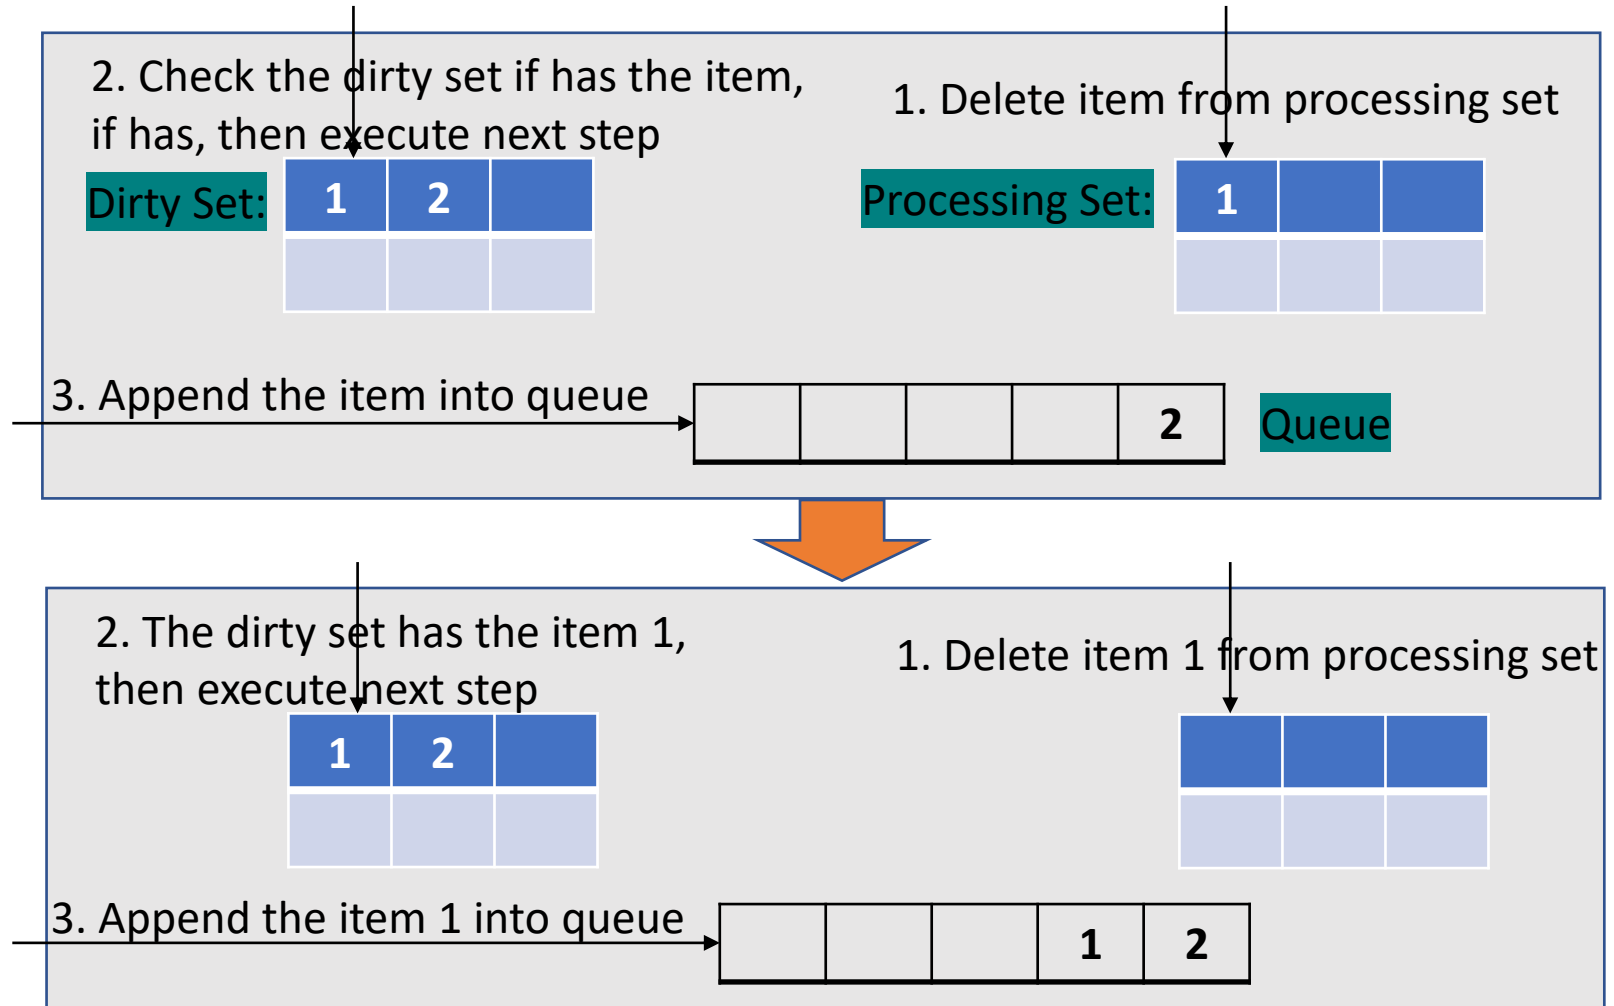
Basic Queue

- **Get** an item handling.
- **Get** blocks until it can return an item to be processed. If shutdown = true, the caller should end their goroutine. You must call Done with item when you have finished processing it.



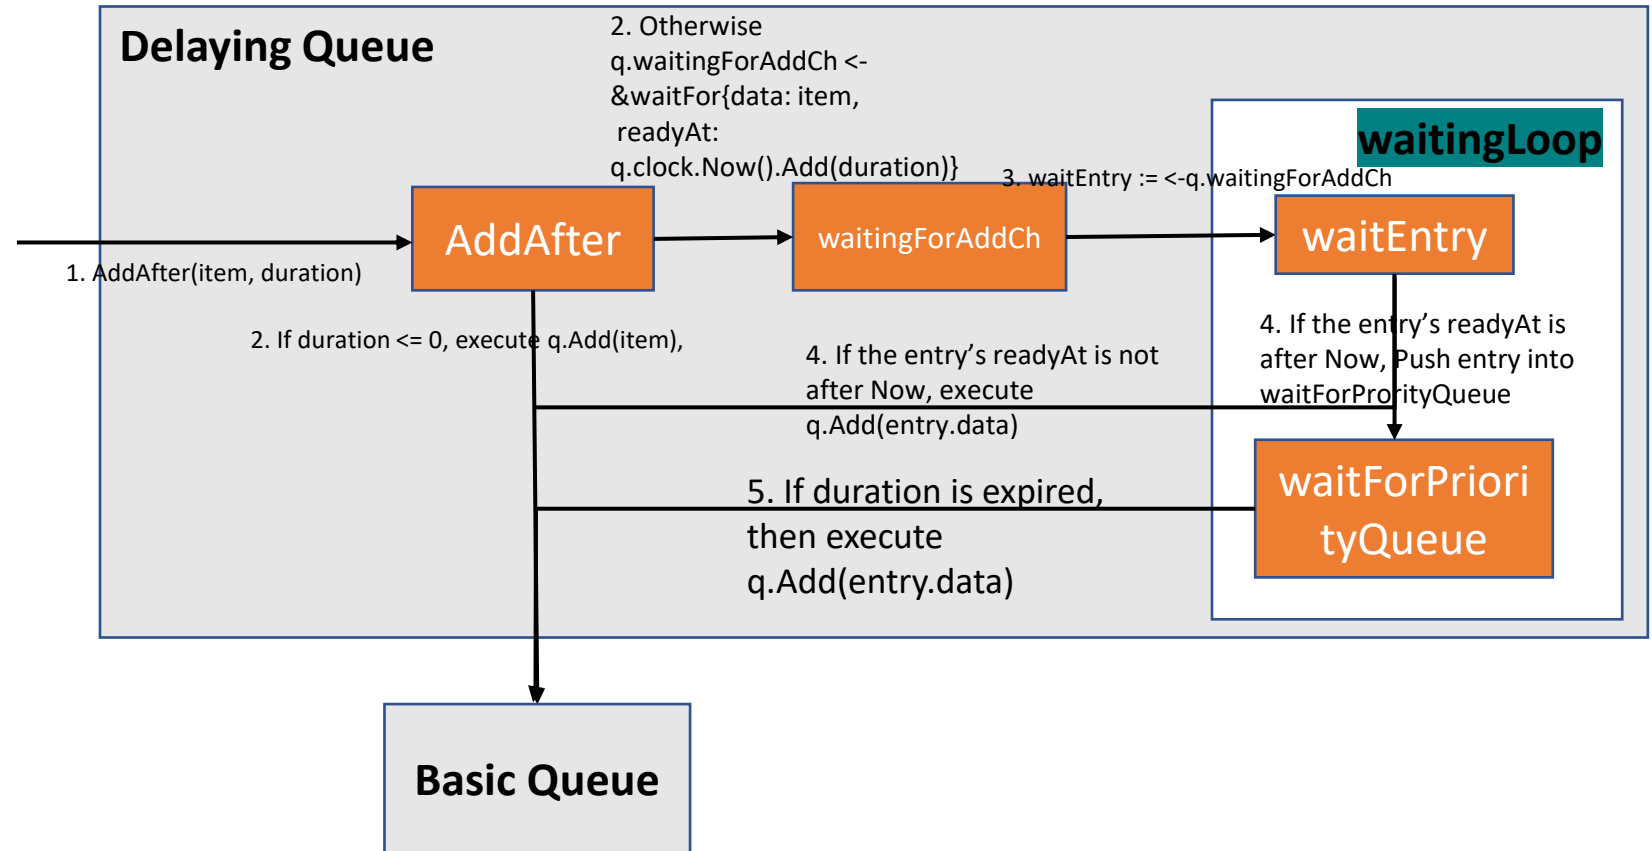
Basic Queue

- **Done** an item handling.
- **Done** marks item as done processing, and if it has been marked as dirty again while it was being processed, it will be re-added to the queue for re-processing.



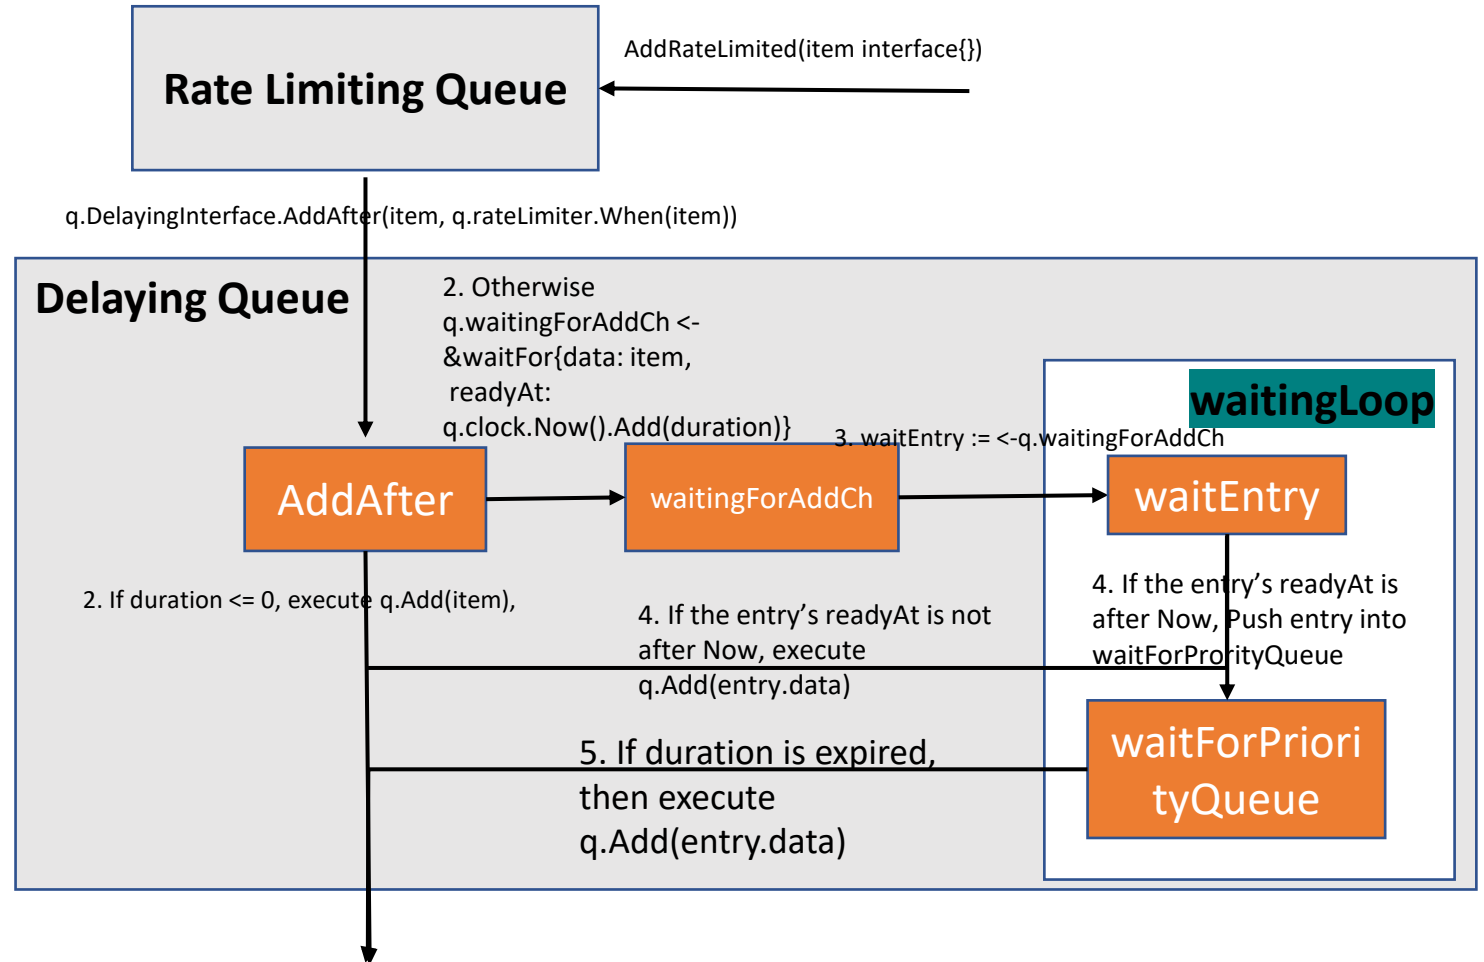
Delaying Queue

- **DelayingInterface** is an Interface that can Add an item at a later time. This makes it easier to requeue items after failures without ending up in a hot-loop.



Rate limiting Queue

- **RateLimitingInterface** is an interface that rate limits items being added to the queue.
- Rate limiting queue is an interface, so if you don't want using default queue, then you can define a custom rate limiting queue.
- Default rate limiting queue:
 - BucketRateLimiter
 - ItemBucketRateLimiter
 - ItemExponentialFailureRateLimiter
 - ItemFastSlowRateLimiter
 - MaxOfRateLimiter

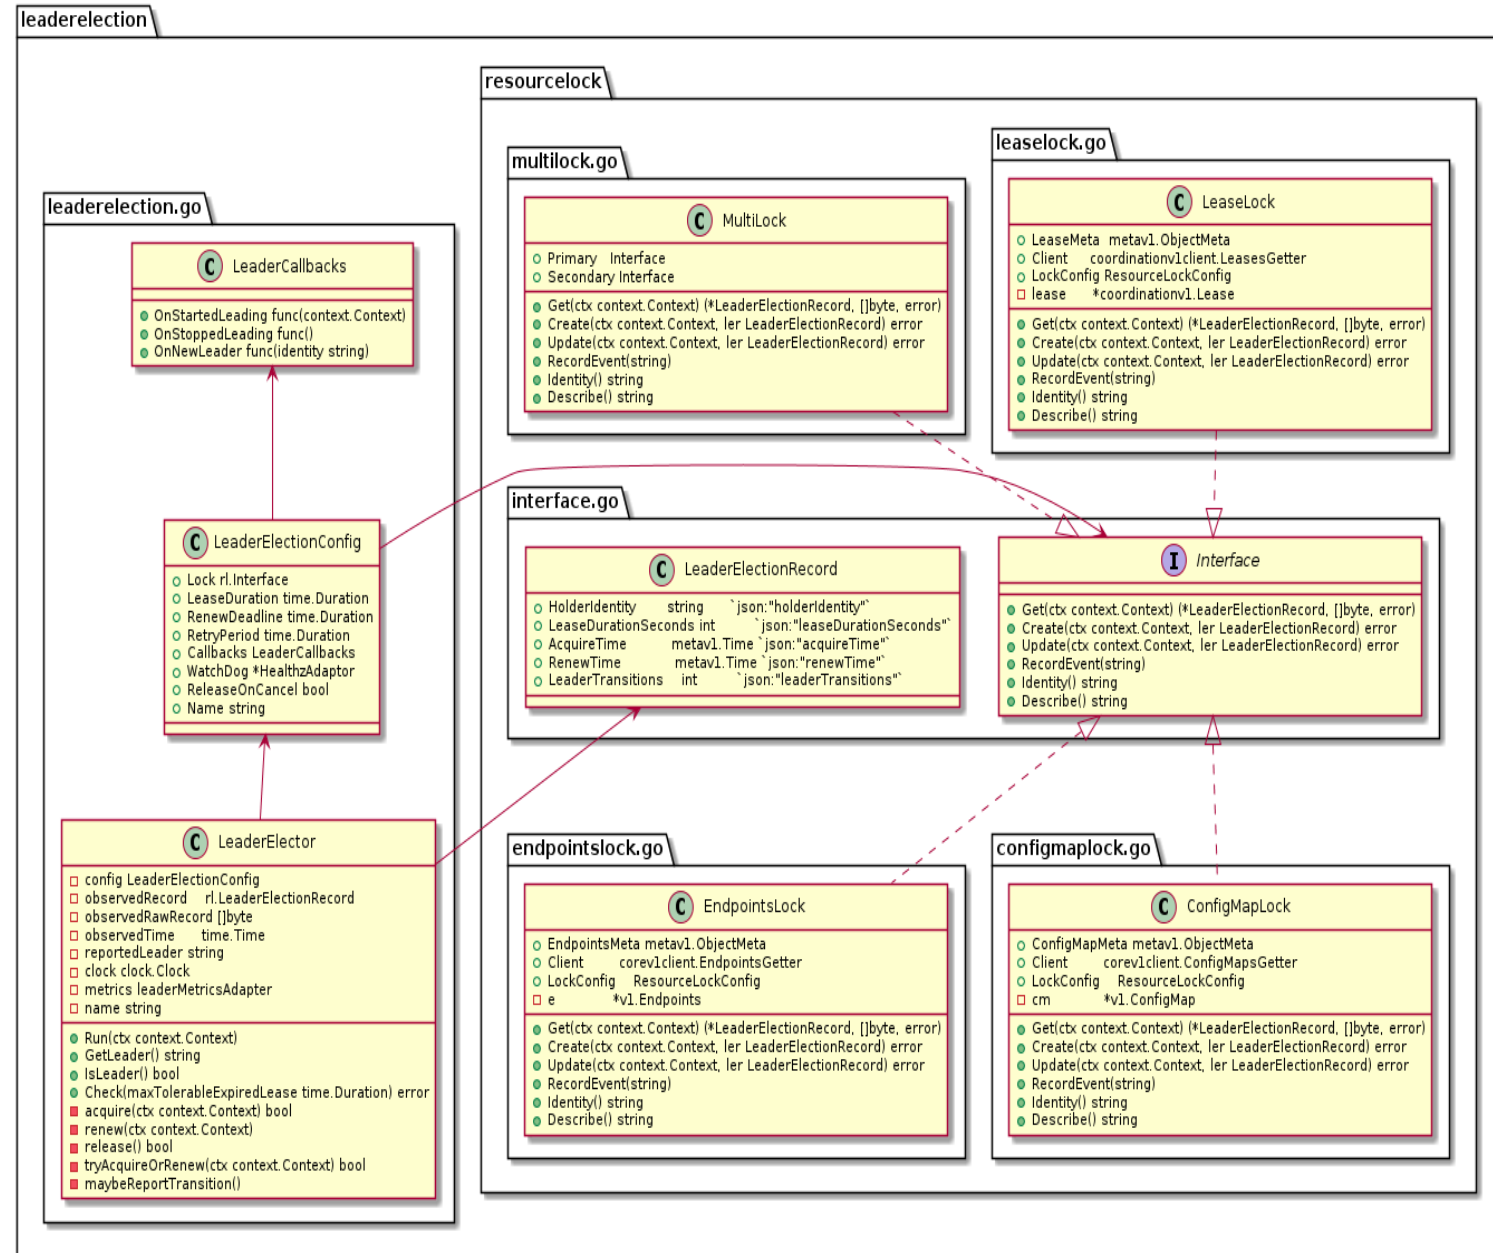


Informer example

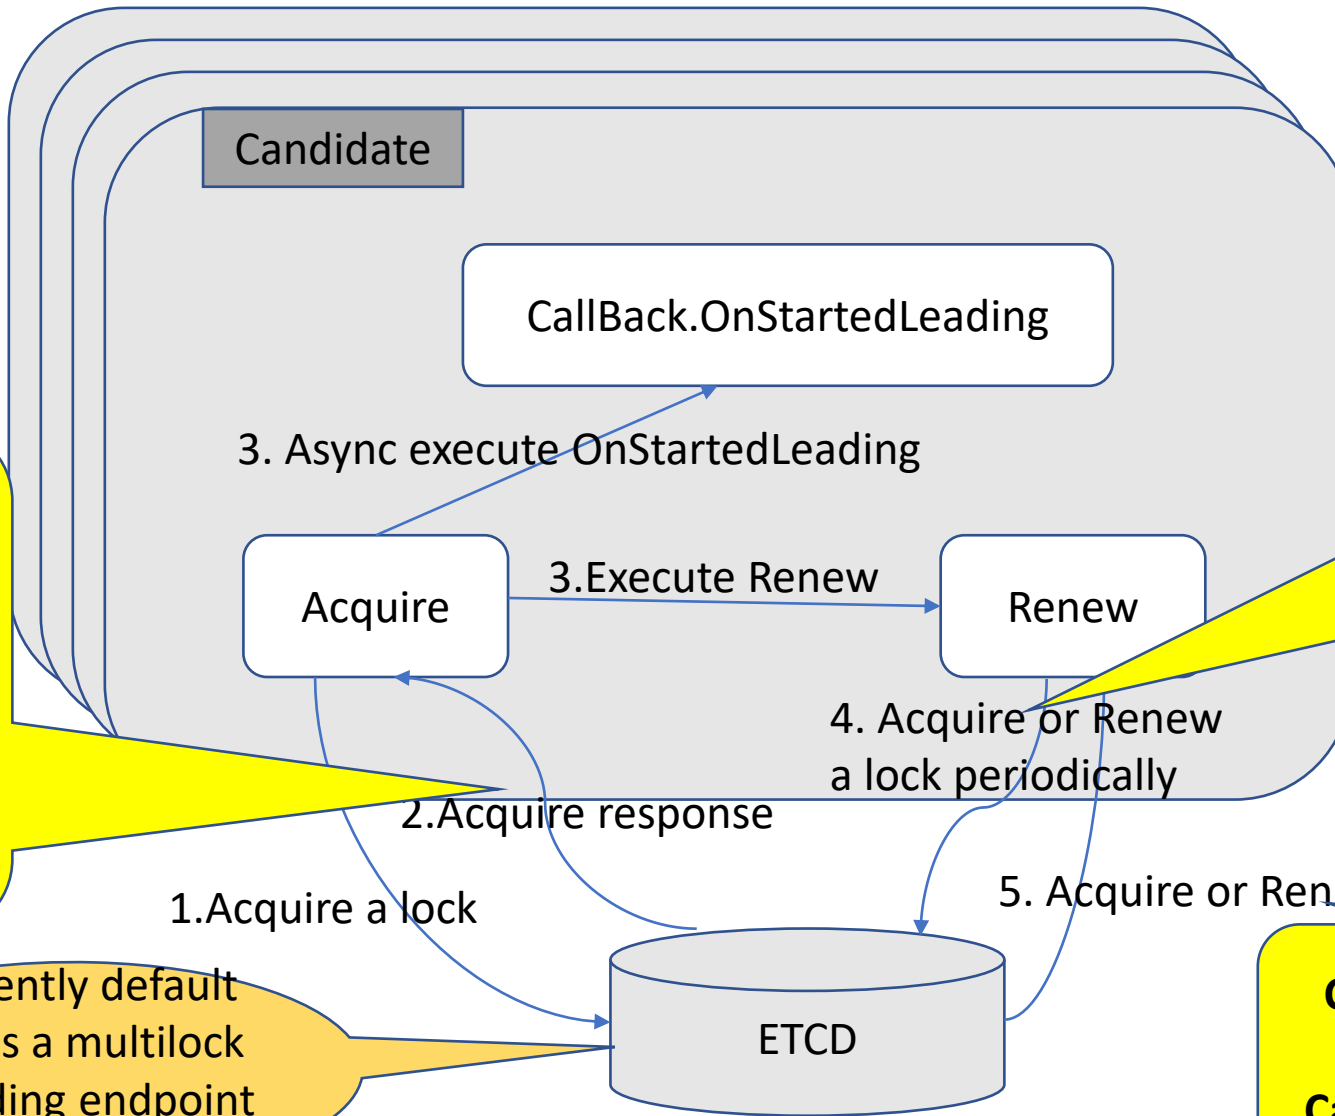
- You can reference k8s official guide(Kubernetes/sample-controller <https://github.com/kubernetes/sample-controller>)
- You can also reference client-go official example(workqueue <https://github.com/kubernetes/client-go/tree/master/examples/workqueue>). This is a typically example use client-go informer, it is fully based on this client-go informer architecture. And almost all of K8S controllers are implemented based on this architecture.

Leader Election

- leaderelection implements leader election of a set of endpoints. It uses an annotation in the endpoints object to store the record of the election state. This implementation does not guarantee that only one client is acting as a leader (a.k.a. fencing).



Leader Election



Case1: if succeeded, candidate will become a leader, and execute Step3.
Case2: if failed, execute Step1 with some interval.

Currently default lock is a multilock including endpoint and lease

Case1: creating a lock if lock don't exist.
Case2: updating lock record information if lock exists.

Case1: if succeeded, execute step4 with some interval.
Case2: if failed, exit and release lock.

Leader Election

- Currently kube-controller-manager is using “endpointlease” resource lock as a default value. So it will have two resource locks, the primary one is endpoint lock, the secondary one is lease lock.
- Q1: why introduce multilock?
- Q2: how to ensure one client is acting as a leader?

References

- Source code: <https://github.com/kubernetes/client-go>
- Article :
 - <https://github.com/opsnull/kubernetes-dev-docs/tree/master/client-go>
 - <https://tangxusc.github.io/blog/2019/05/client-go-informer%E6%9C%BA%E5%88%B6/>
 - <https://www.jianshu.com/p/76e7b1a57d2c>
- Informer analysis:
 - <https://xigang.github.io/2019/09/21/client-go/>.
 - <http://www.programmersought.com/article/6135240470/>.
 - <https://www.huweihuang.com/kubernetes-notes/code-analysis/kube-controller-manager/sharedIndexInformer.html>
- Sample-controller : <https://github.com/kubernetes/sample-controller/blob/master/docs/controller-client-go.md>
- Indexer <https://studygolang.com/articles/20402>
- Leader election:
 - <https://medium.com/michaelbi-22303/deep-dive-into-kubernetes-simple-leader-election-3712a8be3a99>
 - <https://kubernetes.io/blog/2016/01/simple-leader-election-with-kubernetes/>

Controller-runtime

- The Kubernetes controller-runtime Project is a set of go libraries for building Controllers. It is leveraged by [Kubebuilder](#) and [Operator SDK](#).
- Controller
- webhook

References

- <https://github.com/kubernetes-sigs/controller-runtime>
- <https://godoc.org/github.com/kubernetes-sigs/controller-runtime/pkg>

Admission webhook

- Admission webhooks are HTTP callbacks that receive admission requests and do something with them. You can define two types of admission webhooks, [validating admission webhook](#) and [mutating admission webhook](#). Mutating admission webhooks are invoked first, and can modify objects sent to the API server to enforce custom defaults. After all object modifications are complete, and after the incoming object is validated by the API server, validating admission webhooks are invoked and can reject requests to enforce custom policies.

References

- <https://www.qikqiak.com/post/k8s-admission-webhook/>
- <https://kubernetes.io/blog/2019/03/21/a-guide-to-kubernetes-admission-controllers/>

Thanks!